

NATURAL LANGUAGE PROCESSING

Foundations, Web Scraping, Sentiment Analysis & Applications

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2026-I



UNIVERSIDAD DISTRITAL
FRANCISCO JOSÉ DE CALDAS

Outline

- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection
- 3 Sentiment Analysis
- 4 Applications and Integration



Outline


- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection
- 3 Sentiment Analysis
- 4 Applications and Integration



What is Natural Language Processing?

Definition

NLP is the intersection of **linguistics**, **computer science**, and **machine learning** enabling computers to **understand**, **process**, and **generate** human language.



Why NLP is Hard:


- *Ambiguity: "I saw the man with the telescope" — who has it?*
- *Context: meaning depends on co-text and world knowledge*
- *Learning from examples*



What is Natural Language Processing?

Definition

NLP is the intersection of **linguistics**, **computer science**, and **machine learning** enabling computers to **understand**, **process**, and **generate** human language.



Why NLP is Hard:


- **Ambiguity:** *"I saw the man with the telescope"* — who has it?
- **Context:** meaning depends on *co-text* and *world knowledge*
- **Variability:** spelling, slang, typos, *dialects*



What is Natural Language Processing?

Definition

NLP is the intersection of **linguistics**, **computer science**, and **machine learning** enabling computers to **understand**, **process**, and **generate** human language.



Why NLP is Hard:


- **Ambiguity:** *"I saw the man with the telescope"* — who has it?
- **Context:** meaning depends on **co-text** and *world knowledge*
- **Variability:** spelling, slang, typos, dialects



What is Natural Language Processing?

Definition

NLP is the intersection of **linguistics**, **computer science**, and **machine learning** enabling computers to **understand**, **process**, and **generate** human language.




Why NLP is Hard:

- **Ambiguity:** *"I saw the man with the telescope"* — who has it?
- **Context:** meaning depends on **co-text** and *world knowledge*
- **Variability:** spelling, slang, typos, **dialects**



The NLP Processing Pipeline



- **Tokenization:** split into **words**, subwords (BPE), or characters
- **Normalization:** lowercase, accent removal, punctuation stripping
- **Stemming:** *running* → *run*; **Lemmatization:** *better* → *good* (grammatically correct root)




Text Representation: From Words to Vectors

Bag-of-Words & TF-IDF

- Doc $\rightarrow \mathbf{v} \in \mathbb{R}^{|V|}$; **order ignored**
- $\text{tfidf}(t, d) = \text{tf} \cdot \log \frac{|D|}{df_t}$
- Downweights **common**, upweights **rare** terms

Word Embeddings

- Dense $\mathbb{R}^d$; semantic similarity $\approx$ distance
- king - man + woman $\approx$ queen



Word2Vec 2-D projection




Text Representation: From Words to Vectors

Bag-of-Words & TF-IDF

- Doc $\rightarrow \mathbf{v} \in \mathbb{R}^{|V|}$; **order ignored**
- $\text{tfidf}(t, d) = \text{tf} \cdot \log \frac{|D|}{df_t}$
- Downweights **common**, upweights **rare** terms

Word Embeddings

- Dense $\mathbb{R}^d$; semantic similarity $\approx$ distance
- king - man + woman $\approx$ queen



Word2Vec 2-D projection




Text Representation: From Words to Vectors

Bag-of-Words & TF-IDF

- Doc $\rightarrow \mathbf{v} \in \mathbb{R}^{|V|}$; **order ignored**
- $\text{tfidf}(t, d) = \text{tf} \cdot \log \frac{|D|}{df_t}$
- Downweights **common**, upweights **rare** terms

Word Embeddings


- Dense $\mathbb{R}^d$; **semantic similarity** $\approx$ distance
- king-man+woman $\approx$ queen



Word2Vec 2-D projection



Key NLP Tasks Overview



BERT, GPT, T5: one model, all tasks via fine-tuning



Outline


- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection**
- 3 Sentiment Analysis
- 4 Applications and Integration



The HTTP Request — Response Cycle



Parsed HTML Tree



• Status: 200 OK | 403 Forbidden | 429 Rate Limited

• Dynamic pages: JS renders after HTTP response


→ need Selenium



The HTTP Request — Response Cycle



Parsed HTML Tree




- Status: 200 OK | 403 Forbidden | 429 Rate Limited
- Dynamic pages: JS renders *after* HTTP response
→ need Selenium



The HTTP Request — Response Cycle




Parsed HTML Tree



- Status: 200 OK | 403 Forbidden | 429 Rate Limited
 - Dynamic pages: JS renders *after* HTTP response
- need Selenium




Scraping Workflow with requests + BeautifulSoup



- Add `time.sleep(1--3)` between requests (**politeness** delay)
- Always respect `robots.txt` and Terms of Service




Scraping Workflow with requests + BeautifulSoup



- Add `time.sleep(1--3)` between requests (**politeness** delay)
- Always respect `robots.txt` and Terms of Service



Scraping vs. APIs vs. Ready Datasets




Legal & Ethical Bounds

Check `robots.txt`, Terms of Service, and GDPR constraints. Do **not** collect personal data without consent.



Scraping vs. APIs vs. Ready Datasets



Legal & Ethical Bounds

Check `robots.txt`, Terms of Service, and **GDPR** constraints. Do **not** collect personal data without consent.



Outline


- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection
- 3 Sentiment Analysis**
- 4 Applications and Integration



What is Sentiment Analysis?

Definition

Automatically identify the **emotional tone** of text: **positive**, **negative**, **neutral** — at multiple granularity levels.




• Output: binary (pos/neg), ternary (+/0/-), or 5-star regression



What is Sentiment Analysis?

Definition


Automatically identify the **emotional tone** of text: **positive**, **negative**, **neutral** — at multiple granularity levels.



- Output: **binary** (pos/neg), **ternary** (+/0/-), or **5-star regression**



Three Approaches: Lexicon, ML, Deep Learning



accuracy & data requirements increase →



Evaluation: *Metrics and Confusion Matrix*

Beyond Accuracy

Accuracy is misleading on **imbalanced** data. Always report per-class metrics.

- Precision = $\frac{TP}{TP + FP}$
- Recall = $\frac{TP}{TP + FN}$
- F1 = $\frac{2 \cdot P \cdot R}{P + R}$
- F1-macro: average over all classes

Confusion Matrix

		high prec.	
		Pos (pred)	Neg (pred)
high rec.	Pos (actual)	TP	FN
	Neg (actual)	FP	TN



Evaluation: *Metrics and Confusion Matrix*

Beyond Accuracy

Accuracy is misleading on **imbalanced** data. Always report per-class metrics.

- Precision = $\frac{TP}{TP + FP}$
- Recall = $\frac{TP}{TP + FN}$
- F1 = $\frac{2 \cdot P \cdot R}{P + R}$
- F1-macro: average over all classes

Confusion Matrix

		high prec.	
		Pos (pred)	Neg (pred)
high rec.	Pos (actual)	TP	FN
	Neg (actual)	FP	TN



Evaluation: *Metrics* and *Confusion Matrix*

Beyond Accuracy

Accuracy is misleading on **imbalanced** data. Always report per-class metrics.

- Precision = $\frac{TP}{TP + FP}$
- Recall = $\frac{TP}{TP + FN}$
- $F1 = \frac{2 \cdot P \cdot R}{P + R}$
- **F1-macro**: average over *all classes*

Confusion Matrix


		high prec.	
		Pos (pred)	Neg (pred)
high rec.	Pos (actual)	TP	FN
	Neg (actual)	FP	TN




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural conventions differ
- **Subjectivity:** “*Film runs 120 min*” has no sentiment

Negation Scope



Context Window



Practical Tip


Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves polarity inversion for BoW models.




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural conventions differ
- **Subjectivity:** “*Film runs 120 min*” has no sentiment

Negation Scope



Context Window



Practical Tip


Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves polarity inversion for BoW models.




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural conventions differ
- **Subjectivity:** “*Film runs 120 min*” has **no sentiment**

Negation Scope



Context Window



Practical Tip


Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves polarity inversion for BoW models.




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural **conventions** differ
- **Subjectivity:** “*Film runs 120 min*” has **no sentiment**

Negation Scope



Context Window



Practical Tip


Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves **polarity inversion** for BoW models.




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural **conventions** differ
- **Subjectivity:** “*Film runs 120 min*” has **no sentiment**

Negation Scope



Context Window



Practical Tip


Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves **polarity inversion** for BoW models.




Sentiment Analysis: *Linguistic Challenges*

- **Negation:** “*Not bad*” → **positive**
- **Sarcasm:** “*Great, another Monday*”
- **Domain shift:** medical $\neq$ movie vocabulary
- **Multilingual:** cultural **conventions** differ
- **Subjectivity:** “*Film runs 120 min*” has **no sentiment**

Negation Scope



Context Window



Practical Tip

Mark negation scope during preprocessing: `not_bad`, `not_great` — preserves **polarity inversion** for BoW models.




Outline

- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection
- 3 Sentiment Analysis
- 4 Applications and Integration**



End-to-End NLP Pipeline




Real-World Applications:

- Chatbots & virtual assistants
- Machine translation (DeepL)
- Search engine ranking
- NER in clinical records
- Spam / phishing detection
- Code generation (GitHub Copilot)



End-to-End NLP Pipeline



Real-World Applications:

- **Chatbots** & virtual assistants
- **Machine translation** (DeepL)
- **Search engine** ranking
- **NER** in clinical records
- **Spam / phishing** detection
- **Code generation** (GitHub Copilot)




Transformer Models: Self-Attention


Key Innovation (Vaswani *et al.*, 2017)

Replace recurrence with **attention**: every token simultaneously queries **all** other tokens, capturing **long-range context** in parallel.

“bank” attends most to “river” → river bank, not financial bank



Transformer Families: Architecture Comparison



BERT pre-training

MLM: predict masked tokens

"The [MASK] by the river"

GPT pre-training

CLM: predict next token


"The bank by the _"

T5 pre-training

Span corruption: mask spans

"The <X> river"

Transformer Families: Architecture Comparison



BERT pre-training

MLM: predict masked tokens

"The [MASK] by the river"

GPT pre-training

CLM: predict next token


"The bank by the ..."

T5 pre-training

Span corruption: mask spans

"The <X> river"

Transformer Families: Architecture Comparison



BERT pre-training

MLM: predict masked tokens

"The [MASK] by the river"

GPT pre-training

CLM: predict next token


"The bank by the _"

T5 pre-training

Span corruption: mask spans

"The <X> river"

Transformer Families: Architecture Comparison



BERT pre-training

MLM: predict masked tokens

"The [MASK] by the river"

GPT pre-training

CLM: predict next token

"The bank by the _"

T5 pre-training

Span corruption: mask spans

"The $\langle X \rangle$ river"

Sentence Embeddings & Semantic Search

Encoding a Sentence

- Feed full sentence through BERT
- Pool [CLS] token or mean-pool hidden states
- Output: $\mathbf{v} \in \mathbb{R}^{768}$ — a semantic fingerprint

Cosine Similarity

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

1 $\approx$ same meaning

0 $\approx$ unrelated



Semantic search in embedding space

*This is the Embed + Vector DB step
in RAG*



Sentence Embeddings & Semantic Search

Encoding a Sentence


- Feed full sentence through BERT
- Pool [CLS] token or mean-pool hidden states
- Output: $\mathbf{v} \in \mathbb{R}^{768}$ — a **semantic fingerprint**

Cosine Similarity

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

1 $\approx$ same meaning

0 $\approx$ unrelated



Semantic search in embedding space

*This is the Embed + Vector DB step
in RAG*



Sentence Embeddings & Semantic Search

Encoding a Sentence


- Feed full sentence through BERT
- Pool [CLS] token or mean-pool hidden states
- Output: $\mathbf{v} \in \mathbb{R}^{768}$ — a **semantic fingerprint**

Cosine Similarity

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

1 $\approx$ same meaning

0 $\approx$ unrelated



Semantic search in embedding space

*This is the Embed + Vector DB step
in RAG*



Sentence Embeddings & Semantic Search

Encoding a Sentence


- Feed full sentence through BERT
- Pool [CLS] token or mean-pool hidden states
- Output: $\mathbf{v} \in \mathbb{R}^{768}$ — a **semantic fingerprint**

Cosine Similarity

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

1 $\approx$ same meaning

0 $\approx$ unrelated



Semantic search in embedding space


*This is the Embed + Vector DB step
in RAG*



Retrieval-Augmented Generation (RAG)

Motivation

LLMs hallucinate. **Grounding** responses in a **verified document store** dramatically reduces factual errors.




- Tools: LangChain, LlamaIndex, Chroma, FAISS
- Enables domain-specific assistants without full retraining



Retrieval-Augmented Generation (RAG)

Motivation

LLMs hallucinate. **Grounding** responses in a **verified document store** dramatically reduces factual errors.



- Tools: LangChain, LlamaIndex, Chroma, FAISS
- Enables **domain-specific** assistants without full retraining



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc explanations
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for efficient fine-tuning
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with live **inference**.



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc explanations
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for efficient fine-tuning
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with live **inference**.



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc explanations
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for efficient fine-tuning
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with live *inference*.



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc **explanations**
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for **efficient fine-tuning**
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: *scrape reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy as a Streamlit app with live inference.*



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc **explanations**
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for **efficient fine-tuning**
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with **live inference**.



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc **explanations**
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for **efficient fine-tuning**
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with **live inference**.



Current Challenges in NLP

- **Hallucination:** LLMs generate **plausible but false** content
- **Bias:** stereotypes inherited from **training corpora**
- **Low-resource languages:** $\sim 7,000$ languages; most resources are **English-centric**
- **Interpretability:** SHAP, LIME, Attention maps as post-hoc **explanations**
- **Computational cost:** GPT-4 $\approx$ \$200M to train; LoRA/PEFT for **efficient fine-tuning**
- **Privacy:** text often contains PII — **GDPR** compliance required

Workshop Connection

Build a complete pipeline: scrape *reviews* $\rightarrow$ *preprocess* $\rightarrow$ *train sentiment classifier* $\rightarrow$ *deploy* as a Streamlit app with **live inference**.



Outline

- 1 Foundations and Basic Concepts
- 2 Web Scraping and Data Collection
- 3 Sentiment Analysis
- 4 Applications and Integration



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/data-analysis-programming>

