

Normal

ADVANCED DATABASES

Databases II

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2025-III



Outline

1 Object-Oriented Databases (OODB)

2 NoSQL Databases

3 Parallel Databases

4 Distributed Databases



Outline

1 Object-Oriented Databases (OODB)

2 NoSQL Databases

3 Parallel Databases

4 Distributed Databases



What is an OODB?

~2000s

- Combines object-oriented programming and database principles.
- Stores **objects** with their methods and attributes.
- Supports complex data and inheritance.

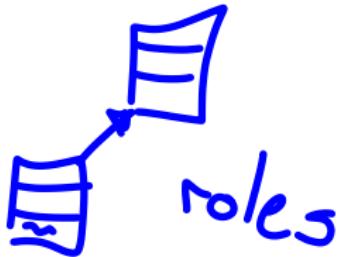
*calculations
queries*



What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.

trees
graphs



OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object storage and identity.
- Supports OQL (Object Query Language).



OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object *storage* and identity.
- Supports OQL (*Object Query Language*).

Demo time!



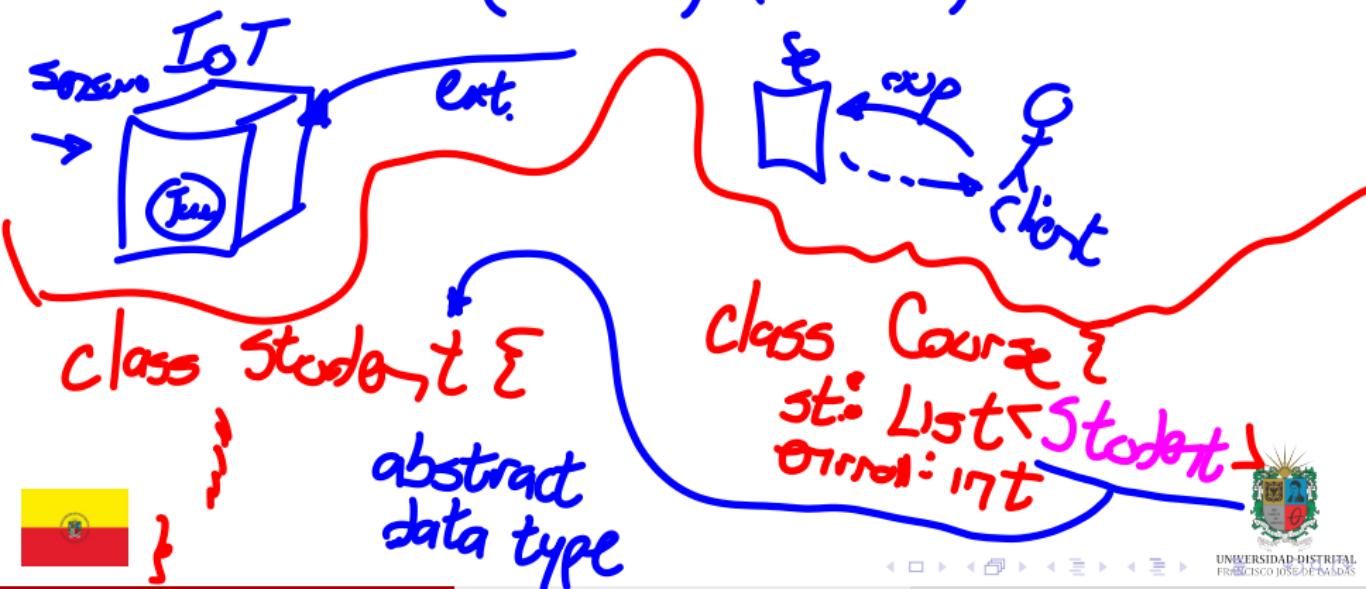
OODB Architecture

- Integrates:

JVM / MICSIL



- Common deployment (embedded) or client-server



Persistence Mechanisms

- By Reachability: Root object persistence.

Rule of Least Knowledge

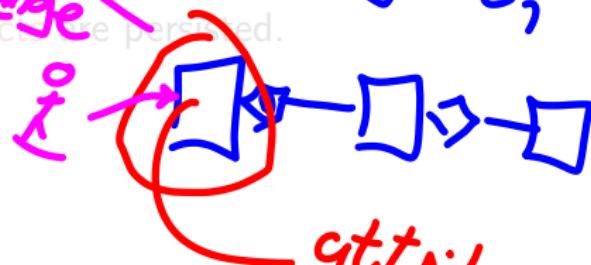
- By Depth Traversing: Annotate persistable objects.

- By Modification: Modified objects are persisted.

bad practice

`a.getB(); getc();`

Composite, aggregation,



attribute
abstract
type *data*



Persistence Mechanisms

- By Reachability: Root object persistence.
- By Explicit Marking: Annotate persistable objects.
- By Modification: Modified objects are persisted.

@Save

```
class A {  
    b:B; x  
    c:C; /  
}
```

```
class B {  
};  
{
```

@Save

```
class C {  
};  
{
```



Persistence Mechanisms

- **By Reachability:** Root object persistence.
- **By Explicit Marking:** Annotate persistable objects.
- **By Modification:** Modified objects are persisted.

att.: null
att2: "a"



OODB Query Example (OQL)

OOI

Student s = new Student()

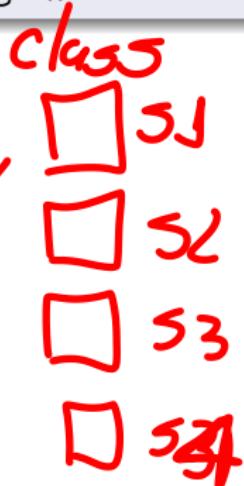
SELECT s.name FROM Student s WHERE s.average() > 4.0

- Queries can follow object references.
- Better for nested, recursive, or complex types.

Class Student {

```
public String name;
private double[] grades;
public double average()
    return Math.avg(grades);
}
```

Demo time!



OODB Query Example (OQL)

OQL

```
SELECT s.name FROM Student s WHERE s.average() > 4.0
```

- **Queries** can follow object references.
- Better for **nested**, **recursive**, or **complex types**.

Demo time!



Outline

1 Object-Oriented Databases (OODB)

2 NoSQL Databases

3 Parallel Databases

4 Distributed Databases

No mandatory



Types of NoSQL Databases

- **Key-Value** *Code*: Redis, DynamoDB (AWS)

- Document: MongoDB, CouchDB

Column-Family: Cassandra, Base

Graph: Neo4j, ArangoDB

Python Dictionary

- `.save(key, value)`
- `.get(key)`

{

"key1": ~ ,
"key2": ~ ,
:
"keyn": ~

}



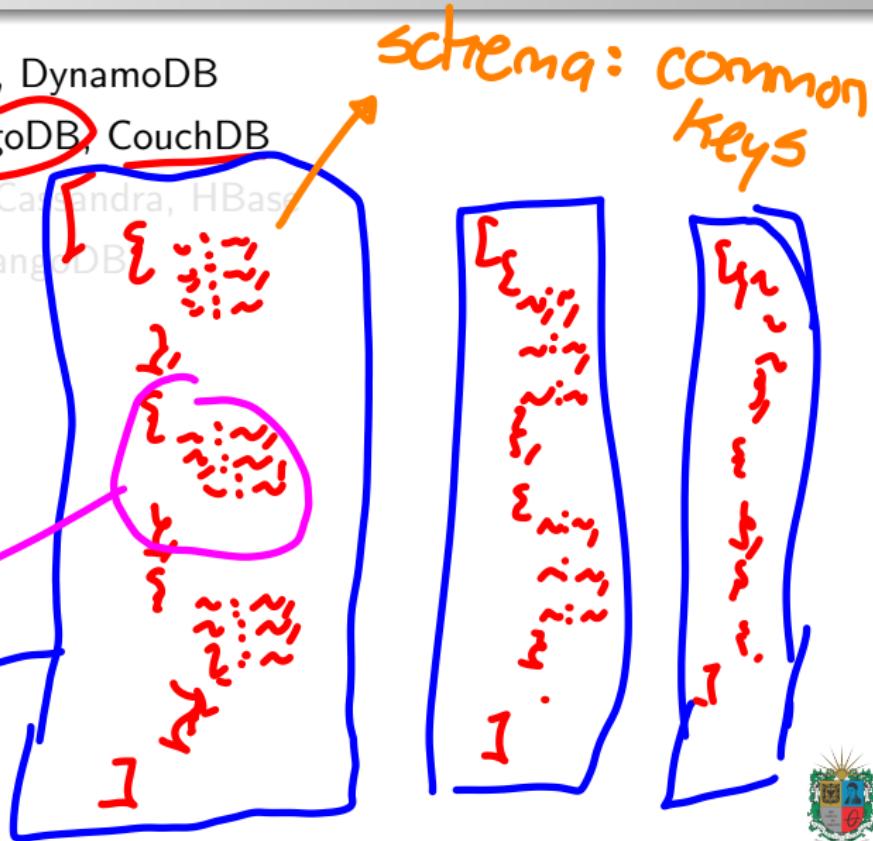
Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB

- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB

LIST OF Dictionaries
JSON

document
collection



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



CAP Theorem Deep Dive

Consistency (C)

Latest write is visible to all reads.

Availability (A)

System always responds.

Partition Tolerance (P)

Tolerates network splits.

Only two can be guaranteed at the same time.



CAP: Practical Examples

Database	CAP Preference	Comment
MongoDB	A + P	Eventual Consistency
Cassandra	A + P	Tunable consistency
Spanner	C + P	Sacrifices availability



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., `mongos` in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



Case Study: MongoDB vs PostgreSQL

- **MongoDB:** JSON-like schema, scalable, flexible.
- **PostgreSQL:** ACID-compliant, strong relational support.
- **Trade-offs** depend on data model **complexity** vs **transactional needs**.

Demo time!



Case Study: MongoDB vs PostgreSQL

- **MongoDB:** JSON-like schema, scalable, flexible.
- **PostgreSQL:** ACID-compliant, strong relational support.
- **Trade-offs** depend on data model **complexity** vs **transactional needs**.

Demo time!



Case Study: Redis with Web API

- In-memory key-value store.
- Excellent for caching, sessions, and real-time analytics.
- Integrates easily with web applications.

Demo time!



Case Study: Redis with Web API

- In-memory key-value store.
- Excellent for caching, sessions, and real-time analytics.
- Integrates easily with web applications.

Demo time!



Outline

1 Object-Oriented Databases (OODB)

2 NoSQL Databases

3 Parallel Databases

4 Distributed Databases



Parallel DB Concepts

- Uses multiple processors to speed up query execution.
- Exploits parallelism in data access and computation.
- Reduces query response time for big datasets.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- Vertical: Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Parallel Query Cost Model

- **Startup cost (S)**: Fixed overhead.
- **Communication cost (C)**: Inter-node data transfer.
- **Computation cost (T)**: Local processing time.
- Total: $T_{total} = S + C + T$



Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



Case Study: PostgreSQL + Citus

- Open-source extension for **PostgreSQL** enabling parallel, distributed queries.
- **MPP architecture** (massively parallel processing) for scaling out across nodes.
- Good for **real-time analytics**, multi-tenant SaaS, and large OLAP workloads.
- Supports **sharding**, **replication**, and **distributed transactions**.

Demo time!



Case Study: PostgreSQL + Citus

- Open-source extension for **PostgreSQL** enabling parallel, distributed queries.
- **MPP architecture** (massively parallel processing) for scaling out across nodes.
- Good for **real-time analytics**, multi-tenant SaaS, and large OLAP workloads.
- Supports **sharding**, **replication**, and **distributed transactions**.

Demo time!



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



What is a Distributed DB?

- **Data stored** across multiple **physical nodes**.
- *Appears* as a **single logical database**.
- Must ensure **consistency**, **availability**, and **partition tolerance**.



Transparency Goals

- **Location:** Hide physical location of data.
- **Replication:** Hide duplication.
- **Fragmentation:** Hide data partitioning.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Two-Phase Commit Protocol (2PC)

- ① **Prepare Phase:** Coordinator asks all participants to **prepare**.
- ② **Commit Phase:** If all vote yes, **coordinator commits**.

Issue: Blocking if coordinator crashes during commit.



Paxos (Simplified Consensus)

- Needed for agreement in **distributed systems**.
- **Roles:**
 - *Proposer*: Suggests a value.
 - *Acceptor*: Votes.
 - *Learner*: Learns chosen value.
- Ensures **consistency** even if nodes fail.



Case Study: Apache Cassandra

- Highly available and scalable NoSQL database.
- Designed for big data applications.
- Supports multi-data center replication.
- Offers a rich query language (CQL).

Demo time!



Case Study: Apache Cassandra

- Highly available and scalable NoSQL database.
- Designed for big data applications.
- Supports multi-data center replication.
- Offers a rich query language (CQL).

Demo time!



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Conclusion

- Data systems have evolved for **scalability** and **complexity**.
- Choosing the **right DB** model depends on the workload requirements.
- Understanding design **trade-offs** is key for architects.



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/databases-ii>

