

ADVANCED DATABASES

Databases II

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Lecturer
Computer Engineer
School of Engineering
Universidad Distrital Francisco José de Caldas

2025-I



- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.



What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.



OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object **storage** and **identity**.
- Supports **OQL** (*Object Query Language*).



Persistence Mechanisms

- **By Reachability:** Root object persistence.
- **By Explicit Marking:** Annotate persistable objects.
- **By Modification:** Modified objects are persisted.



Persistence Mechanisms

- **By Reachability:** **Root object** persistence.
- **By Explicit Marking:** **Annotate** persistable objects.
- **By Modification:** **Modified objects** are persisted.



Persistence Mechanisms

- **By Reachability:** **Root object** persistence.
- **By Explicit Marking:** **Annotate** persistable objects.
- **By Modification:** **Modified objects** are persisted.



OODB Query Example (OQL)

OQL

```
SELECT s.name FROM Student s WHERE s.average() > 4.0
```

- **Queries** can follow **object references**.
- Better for **nested**, **recursive**, or **complex types**.



OODB Architecture

- Integrates:
 - **Object-Oriented Language Runtime**
 - **Persistence Engine**
 - **Query Processor (OQL)**
- **Common deployment:** embedded or client-server.



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases**
- 3 Parallel Databases
- 4 Distributed Databases



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



CAP Theorem Deep Dive

Consistency (C)

Latest write is visible to all reads.

Availability (A)

System always responds.

Partition Tolerance (P)

Tolerates network splits.

Only two can be guaranteed at the same time.



CAP: Practical Examples

Database	CAP Preference	Comment
MongoDB	A + P	Eventual Consistency
Cassandra	A + P	Tunable consistency
Spanner	C + P	Sacrifices availability



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., **mongos** in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



Case Study: MongoDB vs PostgreSQL

- **MongoDB**: JSON-like schema, scalable, flexible.
- **PostgreSQL**: ACID-compliant, strong relational support.
- **Trade-offs** depend on data model **complexity** vs **transactional needs**.



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Parallel DB Concepts

- Uses **multiple processors** to **speed up query execution**.
- Exploits **parallelism** in data access and computation.
- **Reduces query response time** for big datasets.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



Parallel Query Cost Model

- **Startup cost (S):** Fixed overhead.
- **Communication cost (C):** Inter-node data transfer.
- **Computation cost (T):** Local processing time.
- Total: $T_{total} = S + C + T$



Parallel DB Architectures

- **Shared Memory:** Easy communication, poor scalability.
- **Shared Disk:** Easier fault-tolerance, contention risk.
- **Shared Nothing:** Best scalability, harder coordination.



Parallel DB Architectures

- **Shared Memory:** Easy communication, poor scalability.
- **Shared Disk:** Easier fault-tolerance, contention risk.
- **Shared Nothing:** Best scalability, harder coordination.



Parallel DB Architectures

- **Shared Memory:** Easy communication, poor scalability.
- **Shared Disk:** Easier fault-tolerance, contention risk.
- **Shared Nothing:** Best scalability, harder coordination.



Case Study: Greenplum

- Open-source **parallel DB** based on PostgreSQL.
- **MPP architecture** (massively parallel processing).
- Good for **OLAP** and analytics workloads.



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases**



What is a Distributed DB?

- Data stored across multiple physical nodes.
- Appears as a single logical database.
- Must ensure consistency, availability, and partition tolerance.



Transparency Goals

- **Location:** Hide physical location of data.
- **Replication:** Hide duplication.
- **Fragmentation:** Hide data partitioning.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Two-Phase Commit Protocol (2PC)

- 1 **Prepare Phase:** Coordinator asks all participants to prepare.
- 2 **Commit Phase:** If all vote yes, coordinator commits.

Issue: Blocking if coordinator crashes during commit.



Paxos (Simplified Consensus)

- Needed for agreement in **distributed systems**.
- **Roles:**
 - *Proposer*: Suggests a value.
 - *Acceptor*: Votes.
 - *Learner*: Learns chosen value.
- Ensures consistency despite node failures.



Case Study: Google Spanner

- Globally distributed **RDBMS**.
- Uses *TrueTime API* (atomic clocks + GPS).
- Offers external consistency and global transactions.



Conclusion

- **Data systems** have evolved for **scalability** and **complexity**.
- Choosing the **right DB** model depends on workload needs.
- Understanding design **trade-offs** is key for architects.



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/databases-ii>

