

# ADVANCED DATABASES

## Databases II

Author: Eng. Carlos Andrés Sierra, M.Sc.  
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor  
Computer Engineering Program  
School of Engineering  
Universidad Distrital Francisco José de Caldas

2025-III



UNIVERSIDAD DISTRITAL  
FRANCISCO JOSÉ DE CALDAS

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



# Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



# What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.



# What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.



# OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object **storage** and **identity**.
- Supports **OQL** (*Object Query Language*).

Demo time!



# OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object **storage** and **identity**.
- Supports **OQL** (*Object Query Language*).

## Demo time!



# OODB Architecture

- Integrates:
  - **Object-Oriented Language Runtime**
  - **Persistence Engine**
  - **Query Processor (OQL)**
- **Common deployment:** embedded or client-server.





# Persistence Mechanisms

- **By Reachability:** Root object persistence.
- **By Explicit Marking:** Annotate persistable objects.
- **By Modification:** Modified objects are persisted.



# Persistence Mechanisms

- **By Reachability:** Root object persistence.
- **By Explicit Marking:** Annotate persistable objects.
- **By Modification:** Modified objects are persisted.



# Persistence Mechanisms

- **By Reachability:** **Root object** persistence.
- **By Explicit Marking:** **Annotate** persistable objects.
- **By Modification:** **Modified objects** are persisted.



# OODB Query Example (OQL)

## OQL

```
SELECT s.name FROM Student s WHERE s.average() > 4.0
```

- **Queries** can follow **object references**.
- Better for **nested**, **recursive**, or **complex types**.

Demo time!



# OODB Query Example (OQL)

## OQL

```
SELECT s.name FROM Student s WHERE s.average() > 4.0
```

- **Queries** can follow **object references**.
- Better for **nested**, **recursive**, or **complex types**.

# Demo time!



# Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



# Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



# Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB





# Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



# Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



# CAP Theorem Deep Dive

## Consistency (C)

Latest write is visible to all reads.

## Availability (A)

System always responds.

## Partition Tolerance (P)

Tolerates network splits.

**Only two can be guaranteed at the same time.**



# CAP: Practical Examples

Database	CAP Preference	Comment
MongoDB	A + P	Eventual Consistency
Cassandra	A + P	Tunable consistency
Spanner	C + P	Sacrifices availability



# NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., **mongos** in MongoDB).



# NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



# NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).



# NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across **partitions**.
- **Replication**: Ensures **fault-tolerance**.
- **Routers** coordinate requests across **shards** (e.g., mongos in MongoDB).





# Case Study: MongoDB vs PostgreSQL

- **MongoDB**: JSON-like schema, scalable, flexible.
- **PostgreSQL**: ACID-compliant, strong relational support.
- **Trade-offs** depend on data model **complexity** vs **transactional needs**.

Demo time!



# Case Study: MongoDB vs PostgreSQL

- **MongoDB**: JSON-like schema, scalable, flexible.
- **PostgreSQL**: ACID-compliant, strong relational support.
- **Trade-offs** depend on data model **complexity** vs **transactional needs**.

## Demo time!



# Case Study: Redis with Web API

- **In-memory** key-value store.
- Excellent for **caching**, **sessions**, and **real-time analytics**.
- Integrates easily with **web applications**.

Demo time!



# Case Study: Redis with Web API

- **In-memory** key-value store.
- Excellent for **caching**, **sessions**, and **real-time analytics**.
- Integrates easily with **web applications**.

## Demo time!



# Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



# Parallel DB Concepts

- Uses **multiple processors** to **speed up query execution**.
- Exploits **parallelism** in data access and computation.
- **Reduces query response time** for big datasets.



# Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



# Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.





# Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



# Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.



# Parallel Query Cost Model

- **Startup cost (S):** Fixed overhead.
- **Communication cost (C):** Inter-node data transfer.
- **Computation cost (T):** Local processing time.
- Total:  $T_{total} = S + C + T$



# Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



# Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



# Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability**.
- **Shared Disk:** Easier **fault-tolerance**, contention risk.
- **Shared Nothing:** Best scalability, **harder coordination**.



# Case Study: PostgreSQL + Citus

- **Open-source** extension for **PostgreSQL** enabling parallel, distributed queries.
- **MPP architecture** (massively parallel processing) for scaling out across nodes.
- Good for **real-time analytics**, multi-tenant SaaS, and large OLAP workloads.
- Supports **sharding**, **replication**, and **distributed transactions**.

Demo time!



# Case Study: PostgreSQL + Citus

- **Open-source** extension for **PostgreSQL** enabling parallel, distributed queries.
- **MPP architecture** (massively parallel processing) for scaling out across nodes.
- Good for **real-time analytics**, multi-tenant SaaS, and large OLAP workloads.
- Supports **sharding**, **replication**, and **distributed transactions**.

## Demo time!





# Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases**



# What is a Distributed DB?

- **Data stored** across multiple **physical nodes**.
- *Appears* as a **single logical database**.
- Must ensure **consistency**, **availability**, and **partition tolerance**.



# Transparency Goals

- **Location:** Hide physical location of data.
- **Replication:** Hide duplication.
- **Fragmentation:** Hide data partitioning.



# Distributed DBMS Architecture

- **Client-Server Model:** Clients query; **servers** store data.
- **Federated Model:** **Semi-autonomous** DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



# Distributed DBMS Architecture

- **Client-Server Model:** Clients query; **servers** store data.
- **Federated Model:** **Semi-autonomous** DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



# Distributed DBMS Architecture

- **Client-Server Model:** Clients query; **servers** store data.
- **Federated Model:** **Semi-autonomous** DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



# Distributed DBMS Architecture

- **Client-Server Model:** Clients query; **servers** store data.
- **Federated Model:** **Semi-autonomous** DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



# Two-Phase Commit Protocol (2PC)

- 1 **Prepare Phase:** **Coordinator** asks all participants to **prepare**.
- 2 **Commit Phase:** If all vote yes, **coordinator commits**.

**Issue:** Blocking if coordinator crashes during commit.





# Paxos (Simplified Consensus)

- Needed for agreement in **distributed systems**.
- **Roles:**
  - *Proposer*: Suggests a value.
  - *Acceptor*: Votes.
  - *Learner*: Learns chosen value.
- Ensures **consistency** even if nodes fail.



# Case Study: Apache Cassandra

- **Highly available** and **scalable** NoSQL database.
- Designed for **big data** applications.
- Supports **multi-data center** replication.
- **Offers** a rich query language (CQL).

Demo time!



# Case Study: Apache Cassandra

- **Highly available** and **scalable** NoSQL database.
- Designed for **big data** applications.
- Supports **multi-data center** replication.
- **Offers** a rich query language (CQL).

## Demo time!



# Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



# Conclusion

- **Data systems** have evolved for **scalability** and **complexity**.
- Choosing the **right DB** model depends on the workload requirements.
- Understanding design **trade-offs** is key for architects.



# Thanks!

## Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/databases-ii>

