

ADVANCED DATABASES

Databases II

Author: Eng. Carlos Andrés Sierra, M.Sc.

cavirguezs@udistrital.edu.co


Lecturer
Computer Engineer
School of Engineering
Universidad Distrital Francisco José de Caldas

2025-I

Data Engineer

trade-off

Puede



index date



- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



What is an OODB?

(2000) → Java → 295%


- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data and inheritance**.

h



What is an OODB?

- Combines **object-oriented programming** and **database principles**.
- Stores **objects** with their **methods** and **attributes**.
- Supports **complex data** and **inheritance**.



OODB Features

- Encapsulation, Inheritance, Polymorphism.
- Direct object storage and identity.
- Supports OQL (Object Query Language).



public

is Scholarship : bool {
return grade > 4.5 ? true : false;
}

get
grant

set

public

is Approve : bool {

return grade >= 3.0 ? true : false;

Validation

grade > 50 ? true : false;

Serialization



Persistence Mechanisms

- **By Reachability:** Root object persistence.
- By Explicit Marking: Annotate persistable objects.
- By Modification: Modified objects are persisted.



Persistence Mechanisms

- **By Reachability:** Root object persistence.
- **By Explicit Marking:** Annotate persistable objects.
- **By Modification:** Modified objects are persisted.

```


@time
public int getAge() {
    ...
}

@persist
public class Student {
    ...
}
  
```



Persistence Mechanisms

- **By Reachability:** **Root object** persistence.
- **By Explicit Marking:** **Annotate** persistable objects.
- **By Modification:** **Modified objects** are persisted.



OODB Query Example (OQL)

OQL

```
SELECT s.name FROM Student s WHERE s.average() > 4.0
```

- **Queries** can follow object references.
- Better for nested, recursive, or complex types.

SELECT s.name FROM Student s
WHERE s.isScholarship();
(Legacy)



OODB Architecture

- Integrates:

- Object-Oriented Language Runtime
- Persistence Engine
- Query Processor (OQL)

- Common deployment: embedded or client-server.

Tree Optimization Execution

Virtual Memory

Object → attributes
 → methods



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Types of NoSQL Databases

- **Key-Value:** ^{Apartir de FOSS} Redis, DynamoDB → AWS

- Document: MongoDB, CouchDB

- Column-Family: Cassandra, HBase


- Graph: Neo4j, ArangoDB

\downarrow Dict $\rightarrow \{ \text{"key": value} \}$
 \downarrow "key": value
 \downarrow "key2": value
 \downarrow }
 redis['key'] = []; ↗ overwrite



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



write-update


No table

list of columns



Types of NoSQL Databases

- **Key-Value:** Redis, DynamoDB
- **Document:** MongoDB, CouchDB
- **Column-Family:** Cassandra, HBase
- **Graph:** Neo4j, ArangoDB



CAP Theorem Deep Dive

Consistency (C)

Latest write is visible to all reads.


Availability (A)

System always responds

Partition Tolerance (P)

Tolerates network splits.


Only two can be guaranteed at the same time.



CAP: Practical Examples

Database	CAP Preference	Comment
MongoDB	$A + P$	Eventual Consistency
Cassandra	$A + P$	Tunable consistency
Spanner	$C + P$	Sacrifices availability

↓
Latency




NoSQL Architecture Patterns

- **Shared-nothing architecture** enables scalability.

• **Sharding**: Distributes data across partitions.

• **Replication**: Ensures fault-tolerance.

• **Router**: Coordinates requests across shards (e.g., mongos in MongoDB)




NoSQL Architecture Patterns

- **Shared-nothing architecture** enables **scalability**.
- **Sharding**: Distributes data across partitions.

- **Replication**: Ensures fault-tolerance.

- **Routers** coordinate requests across shards (e.g. mongos in MongoDB)



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables scalability.
- **Sharding**: Distributes data across partitions.
- **Replication**: Ensures fault-tolerance.
- **Routers** coordinate requests across shards (e.g., mongos in MongoDB).



NoSQL Architecture Patterns

- **Shared-nothing architecture** enables scalability.
- **Sharding**: Distributes data across partitions.
- **Replication**: Ensures fault-tolerance.
- **Routers** coordinate requests across shards (e.g., mongos in MongoDB).

middleware



Case Study: MongoDB vs PostgreSQL

- **MongoDB**: JSON-like schema, scalable, flexible. *BSON → extending*
- **PostgreSQL**: ACID-compliant, strong relational support. *SQL*
- **Trade-offs** depend on data model complexity vs transactional needs.

Hybrid → no/bi-model system



Outline

1 Object-Oriented Databases (OODB)

•

2 NoSQL Databases


3 Parallel Databases

4 Distributed Databases



Parallel DB Concepts

- Uses multiple processors to speed up query execution.
- Exploits **parallelism** in data access and computation.
- **Reduces query response time** for big datasets.



①
independent
read
update



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on **value intervals**.




Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** *str. like min* based on **value intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**.
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on **value intervals**.



Partitioning Strategies

- **Horizontal:** Distributes **rows**. *files → data*
- **Vertical:** Splits **columns**.
- **Hash Partitioning:** Uniform distribution via **hash**.
- **Range Partitioning:** Based on value **intervals**.

index

1-10	pr1
20-20	pr2
20-30	pr3



Parallel Query Cost Model

- **Startup cost (S):** Fixed overhead.
- **Communication cost (C):** Inter-node data transfer.
- **Computation cost (T):** Local processing time.
- Total: $T_{total} = S + C + T$




Parallel DB Architectures

- **Shared Memory:** Easy communication, **poor scalability.**


- **Shared Disk:** Easier fault-tolerance, contention risk.

- **Shared Nothing:** Best scalability, harder coordination.



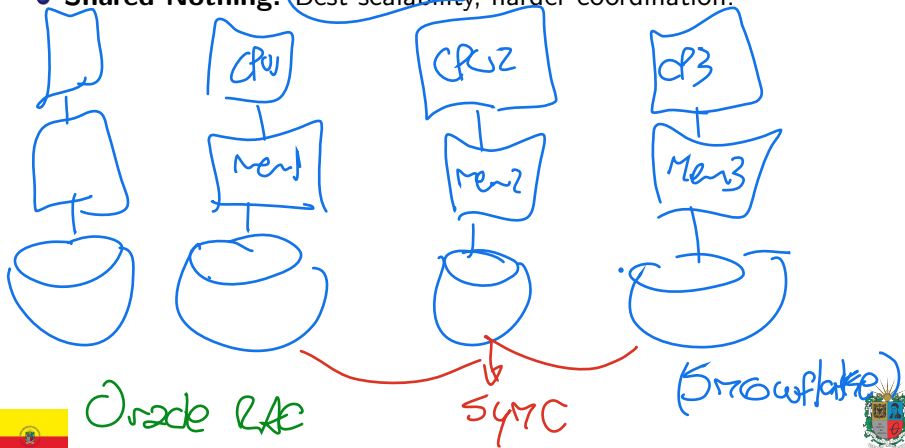
Parallel DB Architectures

- **Shared Memory:** Easy communication, poor scalability.
- **Shared Disk:** Easier fault-tolerance, contention risk.
- **Shared Nothing:** Best scalability, harder coordination.



Parallel DB Architectures

- **Shared Memory:** Easy communication, poor scalability.
- **Shared Disk:** Easier fault-tolerance, contention risk.
- **Shared Nothing:** Best scalability, harder coordination.



Case Study: Greenplum


- Open-source parallel DB based on PostgreSQL.
- MPP architecture (massively parallel processing).
- Good for OLAP and analytics workloads.

relational

10-100
cores

online

superset (SQL
Builder)




Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases**



What is a Distributed DB?

- Data stored across multiple physical nodes.
- Appears as a single logical database.
- Must ensure consistency, availability, and partition tolerance.



Transparency Goals

USER doesn't know


- **Location:** Hide physical location of data.
- **Replication:** Hide duplication.
- **Fragmentation:** Hide data partitioning.

→ Scalability (redundant)




Distributed DBMS Architecture

- Client-Server Model: Clients query; servers store data.
 - Federated Model: Semi-autonomous DBs collaborate.
 - Peer-to-Peer: Nodes act as both client and server.
- Keywords: Global schema, transaction manager, local engines.




Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.




Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- Layers: Global schema, transaction manager, local engines.



Distributed DBMS Architecture

- **Client-Server Model:** Clients query; servers store data.
- **Federated Model:** Semi-autonomous DBs collaborate.
- **Peer-to-Peer:** Nodes act as both client and server.
- **Layers:** Global schema, transaction manager, local engines.



Two-Phase Commit Protocol (2PC)

- 1 **Prepare Phase:** Coordinator asks all participants to prepare.
- 2 **Commit Phase:** If all vote yes, coordinator commits.

Issue: Blocking if coordinator crashes during commit.

DBMS



Paxos (Simplified Consensus)

- Needed for agreement in **distributed systems**.
- **Roles:**
 - *Proposer*: Suggests a value.
 - *Acceptor*: Votes.
 - *Learner*: Learns chosen value.
- Ensures consistency despite node failures.

Consensus algorithm
↓
blockchain



Case Study: Google Spanner

- Globally distributed RDBMS.
- Uses *TrueTime API* (atomic clocks + GPS).
- Offers external consistency and global transactions.

→ consistency



Conclusion

- **Data systems** have evolved for **scalability** and **complexity**.
- Choosing the **right DB** model depends on workload needs.
- Understanding design **trade-offs** is key for architects.

write
another
user

quality metric
(security)

↑

quality metric
(scalability)

↓



Outline

- 1 Object-Oriented Databases (OODB)
- 2 NoSQL Databases
- 3 Parallel Databases
- 4 Distributed Databases



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/databases-ii>

