

DATA BASE SYSTEMS ARCHITECTURE

Databases III

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2025-III



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



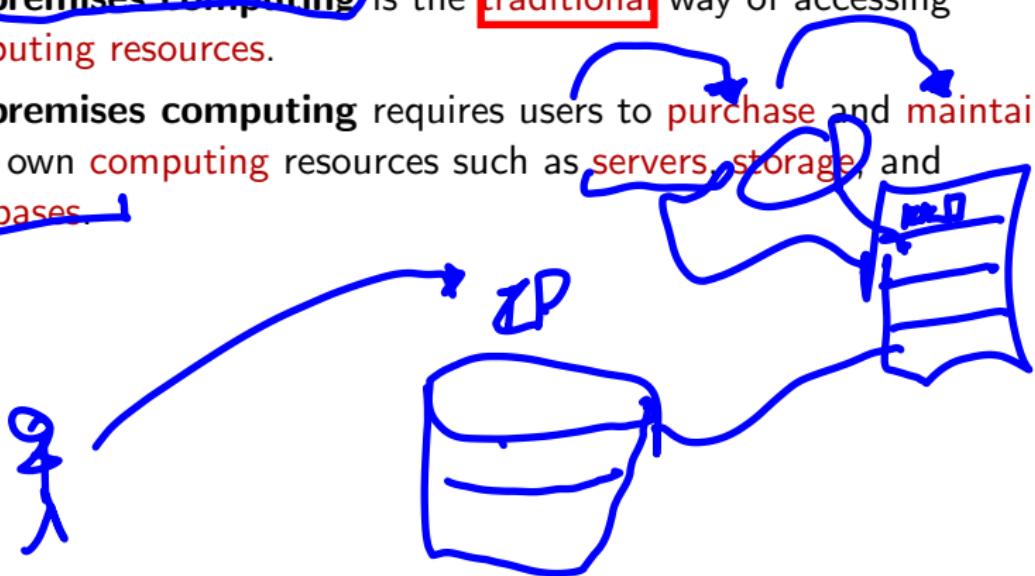
Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



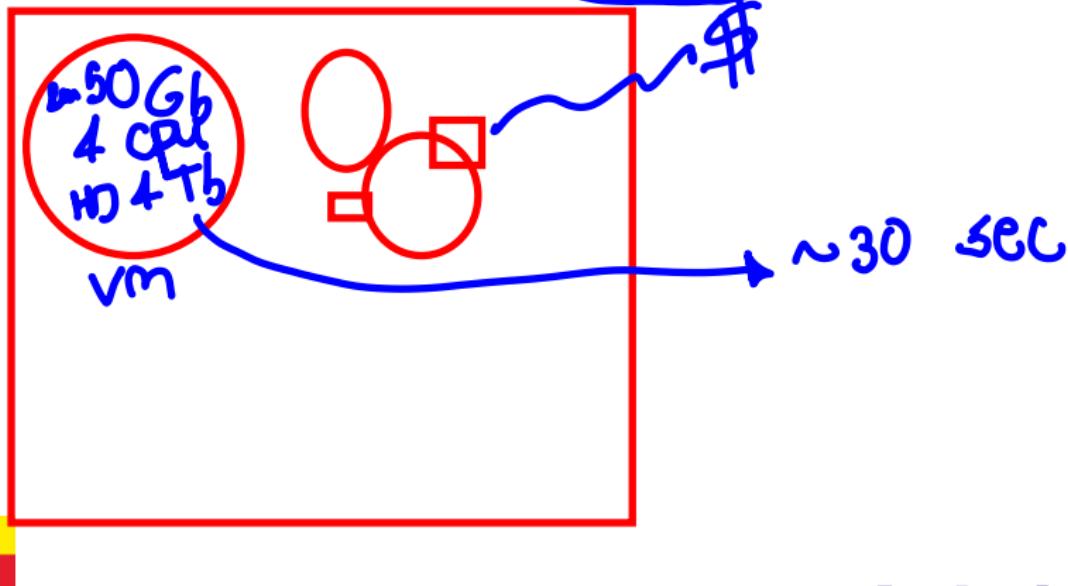
What is On-Premises Computing?

- **On-premises computing** is the **traditional** way of accessing computing resources.
- **On-premises computing** requires users to **purchase** and **maintain** their own **computing** resources such as **servers**, **storage**, and **databases**.



What is the Cloud Computing?

- **Cloud computing** is the delivery of **computing services** over the **internet**.
- **Cloud computing** allows users to **access computing resources** such as servers, storage, and databases **on demand**.

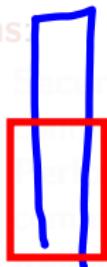


Pros & Cons of Cloud Computing

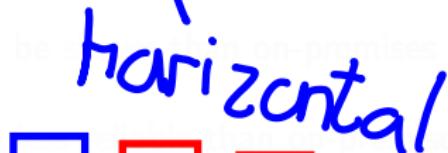
- Pros:

- **Cost-Effective:** Cloud computing is a **cost-effective** way to access computing resources.
- **Scalable:** Cloud computing is a **scalable** way to access computing resources.
- **Flexible:** Cloud computing is a **flexible** way to access computing resources.

- Cons:



vertical



horizontal



Pros & Cons of Cloud Computing

• Pros:

- **Cost-Effective:** Cloud computing is a cost-effective way to access computing resources.
- **Scalable:** Cloud computing is a scalable way to access computing resources.
- **Flexible:** Cloud computing is a flexible way to access computing resources.

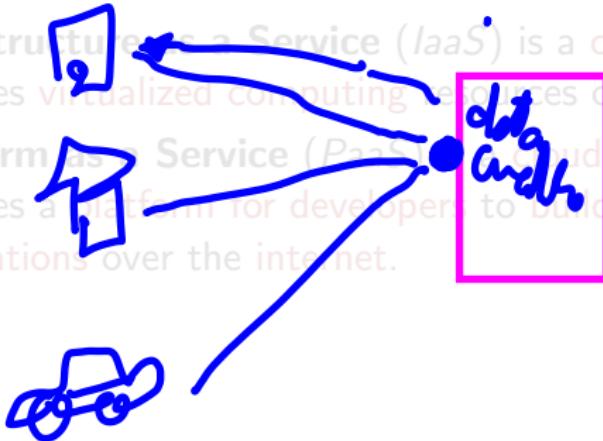
• Cons:

- **Security:** Cloud computing can be less secure than on-premises computing.
- **Performance:** Cloud computing can be slower than on-premises computing.
- **Reliability:** Cloud computing can be less reliable than on-premises computing.



SaaS Vs. IaaS Vs. PaaS

- **Software as a Service (SaaS)** is a **software distribution** model in which a **third-party** provider **hosts applications** and makes them available to customers over the **internet**.
- **Infrastructure as a Service (IaaS)** is a **cloud computing** model that provides **virtualized computing resources** over the **internet**.
- **Platform as a Service (PaaS)** is a **cloud computing** model that provides a **platform** for developers to **build, deploy, and manage applications** over the **internet**.



SaaS Vs. IaaS Vs. PaaS

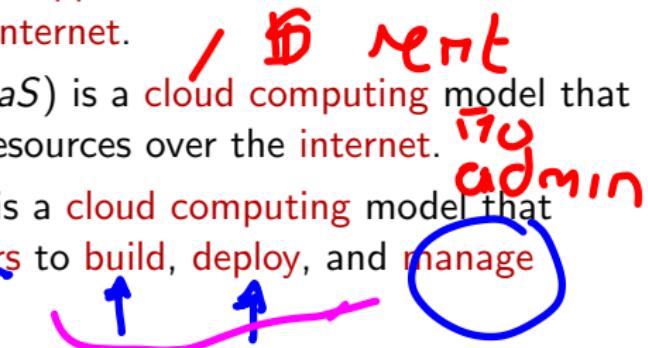
- **Software as a Service (SaaS)** is a **software distribution** model in which a **third-party** provider **hosts** applications and makes them available to customers over the **internet**.
- **Infrastructure as a Service (IaaS)** is a **cloud computing** model that provides **virtualized computing** resources over the **internet**.
- **Platform as a Service (PaaS)** is a **cloud computing** model that provides a platform for developers to build, deploy and manage applications over the **internet**.

IaC → terraform



SaaS Vs. IaaS Vs. PaaS

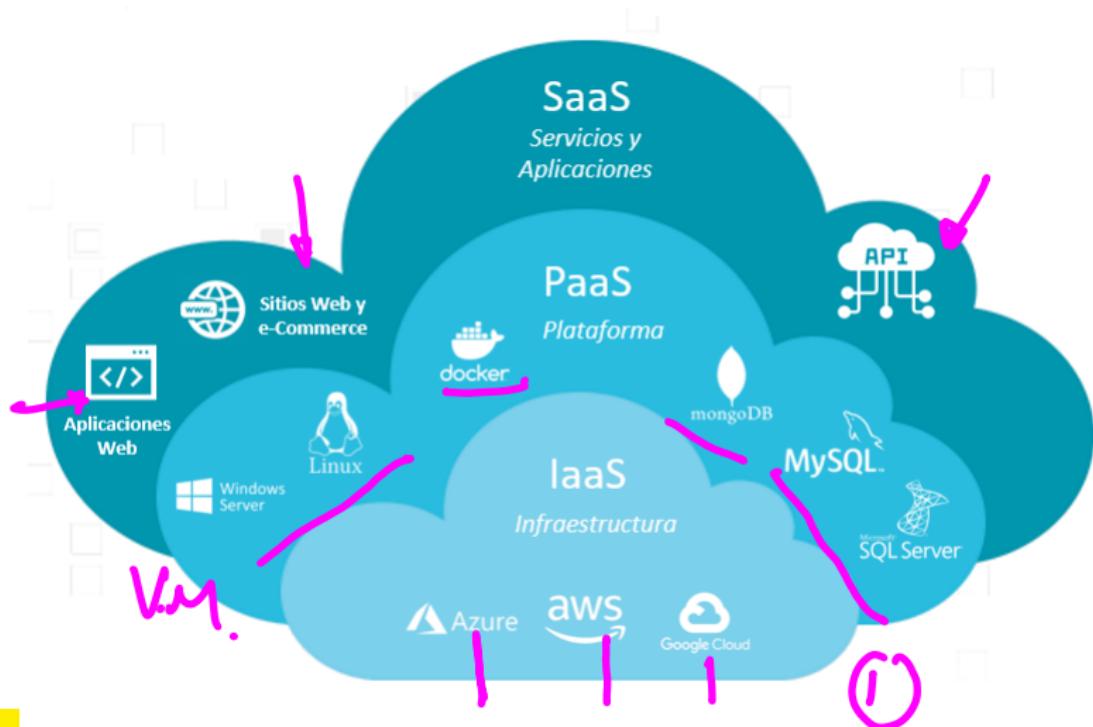
- **Software as a Service (SaaS)** is a **software distribution** model in which a **third-party** provider **hosts** applications and makes them available to customers over the **internet**. */ \$ rent*
- **Infrastructure as a Service (IaaS)** is a **cloud computing** model that provides **virtualized computing** resources over the **internet**. *to admin*
- **Platform as a Service (PaaS)** is a **cloud computing** model that provides a **platform for developers** to build, deploy, and **manage** applications over the **internet**.



Docker - K8s
Artifacts
ML/AI



Cloud Levels



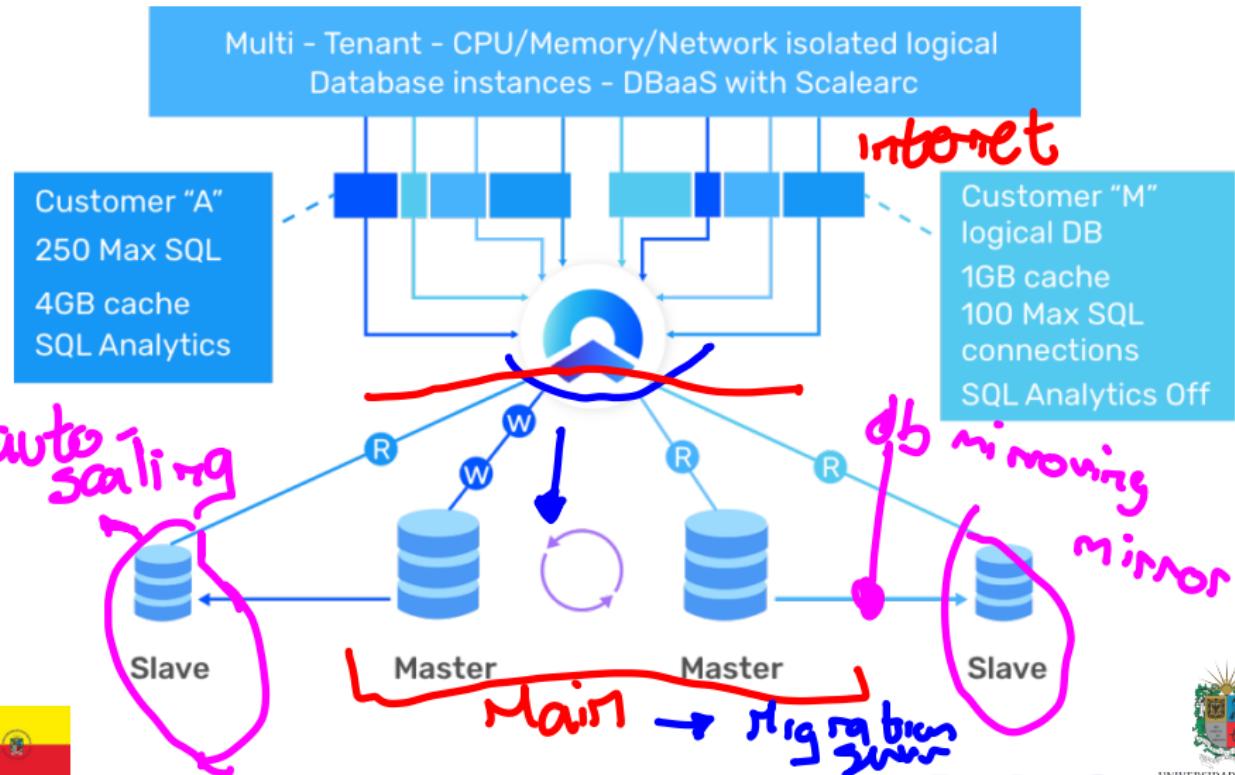
DataBases as a Service

Database as a Service (DBaaS) is a cloud computing model that provides database services over the internet.

- ① DBMS { NoSQL } ready to use
- ② Security : Firewall, Duty, Logs, Policies
- ③ Migration (Snowball → mobile)

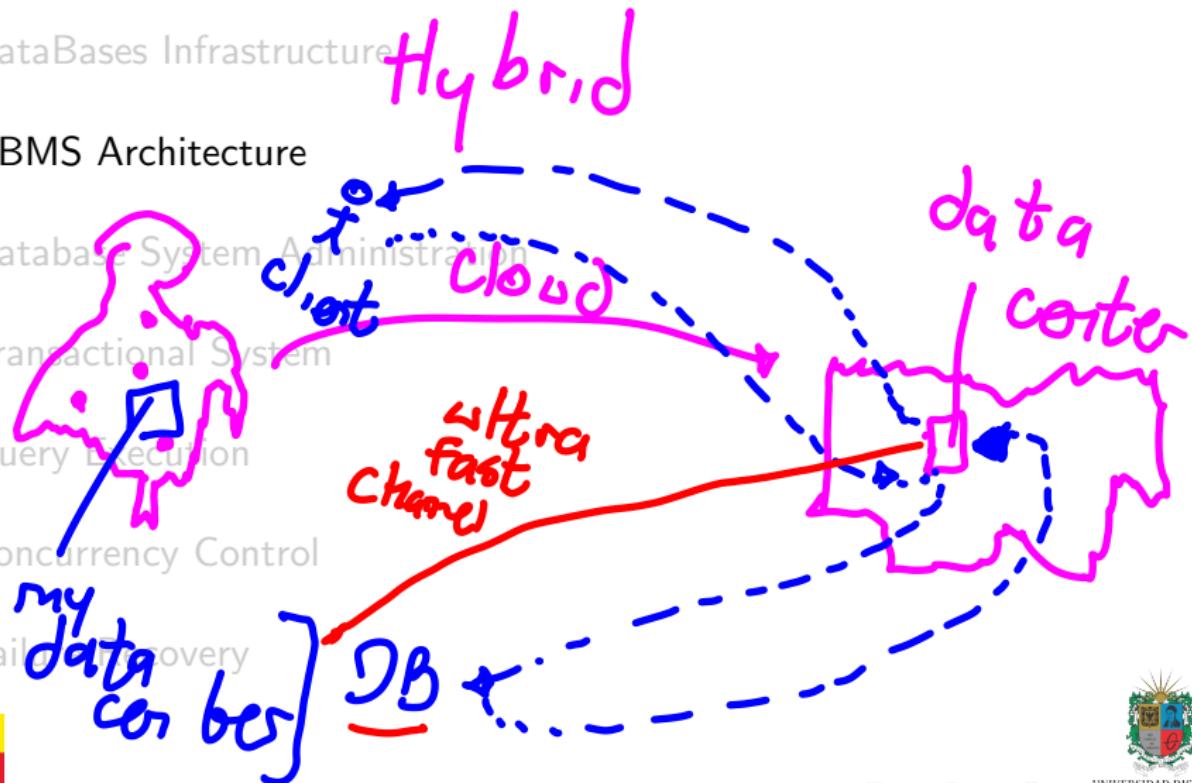


Case Study: DBaaS Custom for Clients



Outline

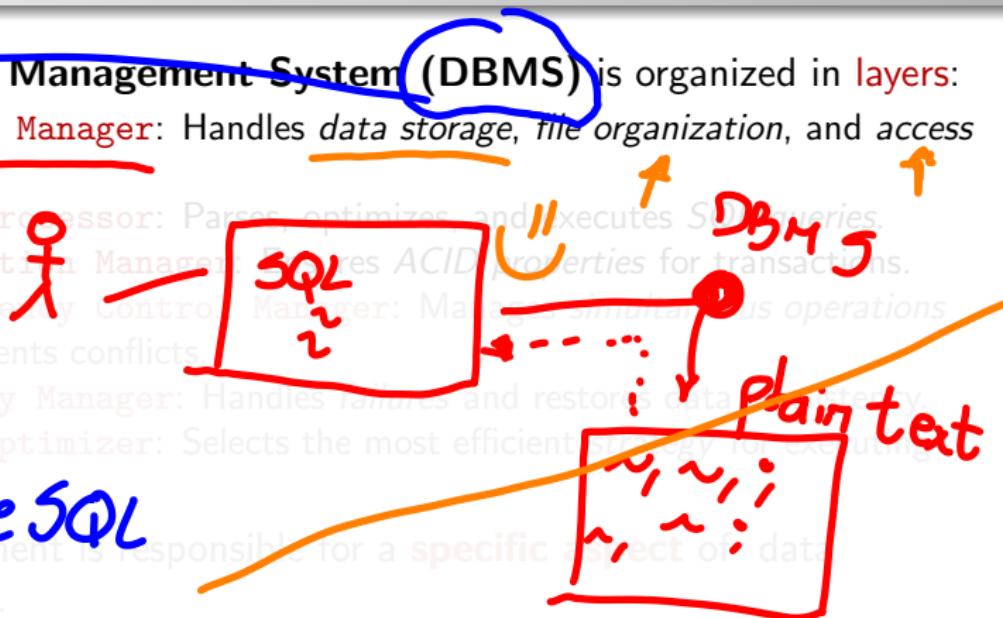
- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



DBMS Architecture Overview

- A Database Management System (DBMS) is organized in layers:
 - Storage Manager:** Handles data storage, file organization, and access methods.
 - Query Processor:** Parses, optimizes, and executes SQL queries.
 - Transaction Manager:** Ensures ACID properties for transactions.
 - Concurrency Control Manager:** Manages simultaneous operations and prevents conflicts.
 - Recovery Manager:** Handles failures and restores data integrity.
 - Optimizer:** Selects the most efficient query execution plan for executing queries.

MySQL
PostgreSQL
Oracle



DBMS Architecture Overview

- A Database Management System (DBMS) is organized in layers:
 - Storage Manager: Handles *data storage, file organization, and access methods.*
 - **Query Processor**: Parses, optimizes, and executes *SQL queries.*
 - Transaction Manager: Ensures *ACID properties for transactions.*
 - Concurrency Control Manager: Manages *simultaneous operations* and prevents conflicts.
 - Recovery Manager: Handles *failures and restores data consistency.*
 - Query Optimizer: Selects the most efficient *strategy for executing queries.*
- Each component is responsible for a *specific aspect* of data management.

graph

Compiler

tree

language

to interact

with storage



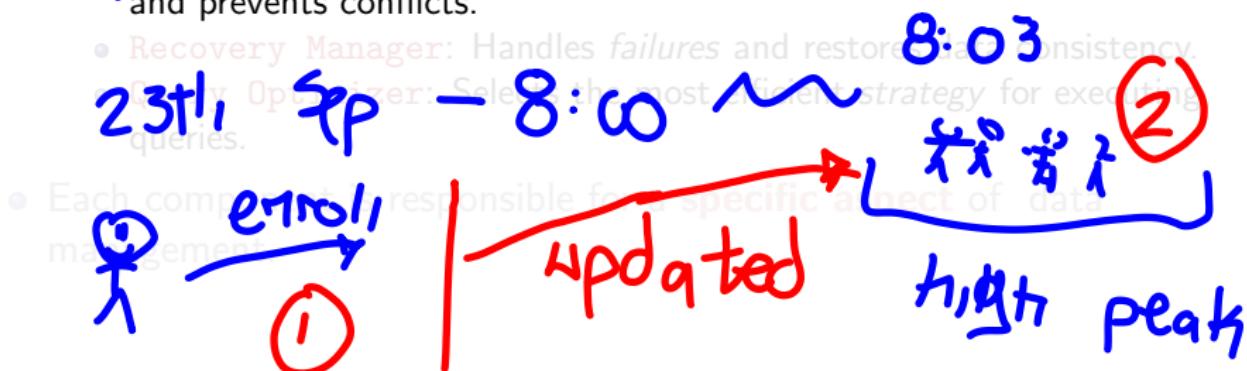
DBMS Architecture Overview

- A **Database Management System (DBMS)** is organized in **layers**:
 - **Storage Manager**: Handles *data storage, file organization, and access methods*.
 - **Query Processor**: Parses, optimizes, and executes *SQL queries*.
 - **Transaction Manager**: Ensures *ACID properties* for transactions.
 - **Synchronization Control Manager**: Manages simultaneous operations and prevents conflicts.
 - **Recovery Manager**: Handles *failures* and restores data consistency.
 - **Query Optimizer**: Selects the most efficient *strategy* for executing queries.
- Each component is responsible for a **specific aspect** of data management.



DBMS Architecture Overview

- A Database Management System (DBMS) is organized in layers:
 - Storage Manager: Handles *data storage, file organization, and access methods.*
 - Query Processor: Parses, optimizes, and executes *SQL queries.*
 - Transaction Manager: Ensures *ACID properties* for transactions.
 - **Concurrency Control Manager**: Manages *simultaneous operations* and prevents conflicts.



DBMS Architecture Overview

- A Database Management System (DBMS) is organized in layers:
 - Storage Manager: Handles *data storage, file organization, and access methods.*
 - Query Processor: Parses, optimizes, and executes *SQL queries.*
 - Transaction Manager: Ensures *ACID properties* for transactions.
 - Concurrency Control Manager: Manages *simultaneous operations* and prevents conflicts.
 - Recovery Manager: Handles *failures* and restores data consistency.
 - Query Optimizer: Selects the most efficient strategy for executing queries.
- Each component is responsible for ~~specific aspect~~ *inherent* aspect of data management.

high level

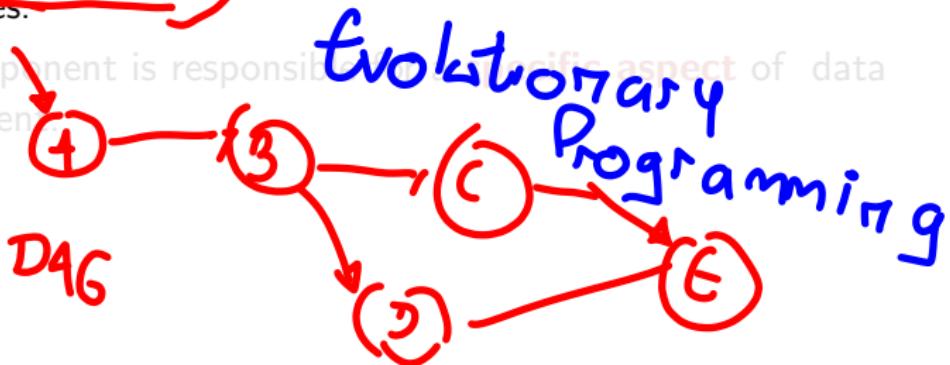
inherent

Backup



DBMS Architecture Overview

- A Database Management System (DBMS) is organized in layers:
 - Storage Manager: Handles *data storage, file organization, and access methods.*
 - Query Processor: Parses, optimizes, and executes *SQL queries.*
 - Transaction Manager: Ensures *ACID properties* for transactions.
 - Concurrency Control Manager: Manages *simultaneous operations* and prevents conflicts.
 - Recovery Manager: Handles *failures* and restores data consistency.
 - Query Optimizer: Selects the most efficient *strategy* for executing queries.
- Each component is responsible for a specific aspect of data management.



DBMS Architecture Overview

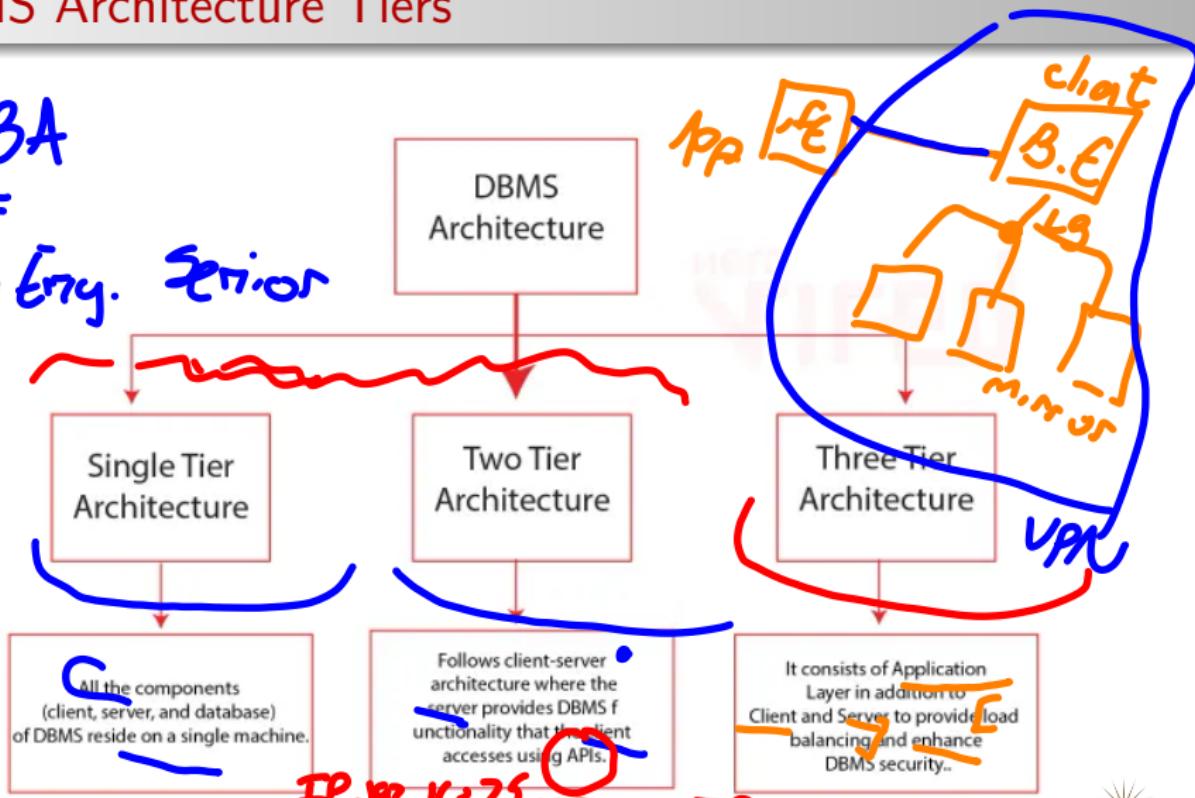
- A **Database Management System (DBMS)** is organized in **layers**:
 - **Storage Manager**: Handles *data storage, file organization, and access methods*.
 - **Query Processor**: Parses, optimizes, and executes *SQL queries*.
 - **Transaction Manager**: Ensures *ACID properties* for transactions.
 - **Concurrency Control Manager**: Manages *simultaneous operations* and prevents conflicts.
 - **Recovery Manager**: Handles *failures* and restores data consistency.
 - **Query Optimizer**: Selects the most efficient *strategy* for executing queries.
- Each component is responsible for a **specific aspect** of data management

all



DBMS Architecture Tiers

DBA
= Data Engr. Senior



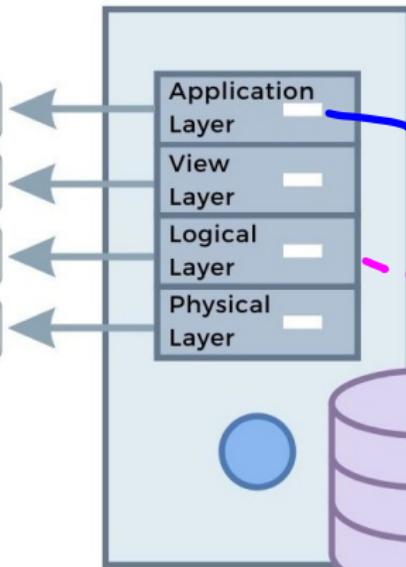
DBMS Architecture N-Tier

DBMS Architecture



DatabaseTown.com

- It is responsible for providing an interface for users.
- It is responsible for managing the different views of the data in the database.
- It is responsible for managing the logical organization of data in the database.
- It is responsible for managing the physical storage of data on disk.



Types of DBMS Architecture

There are several types of DBMS architectures:

- **Centralized DBMS:** All components are on a single server.
- Client-Server DBMS: Clients access the database through a server.
- Distributed DBMS: Data is distributed across multiple servers.
- Cloud DBMS: Database services are provided over the cloud.
- Hybrid DBMS: Combines features of centralized and distributed systems.
- Peer-to-Peer DBMS: Each node can act as a client and server.
- In memory DBMS: Data is stored in RAM for faster access.

1-tier
developers



Types of DBMS Architecture

There are several types of DBMS architectures:

- **Centralized DBMS:** All components are on a **single server**.
- **Client-Server DBMS:** Clients **access** the database through a **server**.
 - **Distributed DBMS:** Data is distributed across **multiple servers**.
 - **Cloud DBMS:** Database services are provided over the **cloud**.
 - **Hybrid DBMS:** Combines features of centralized and distributed systems.
- **Peer-to-Peer DBMS:** Each node can act as a **client** and **server**.
- **In memory DBMS:** Data is stored in **RAM** for faster access.



Types of DBMS Architecture

There are several types of DBMS architectures:

- **Centralized DBMS:** All components are on a **single server**.
- **Client-Server DBMS:** Clients access the database through a **server**.
- **Distributed DBMS:** Data is distributed across **multiple servers**.
- **Cloud DBMS:** Database services are provided over the **cloud**.
- **Hybrid DBMS:** Combines features of centralized and distributed systems.
- **Peer-to-Peer DBMS:** Each node can act as a client and server.
- **In memory DBMS:** Data stored in RAM for faster access.

parallelism

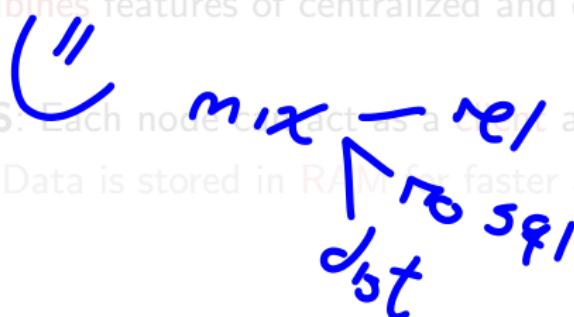
to replicas



Types of DBMS Architecture

There are several types of **DBMS** architectures:

- **Centralized DBMS**: All components are on a **single server**.
- **Client-Server DBMS**: Clients **access** the database through a **server**.
- **Distributed DBMS**: Data is distributed across **multiple servers**.
- **Cloud DBMS**: Database services are provided over the **cloud**.
- **Hybrid DBMS**: Combines features of centralized and distributed systems.
- **Peer-to-Peer DBMS**: Each node **acts** as a **client** and **server**.
- **In memory DBMS**: Data is stored in **RAM** for faster access.



Types of DBMS Architecture

There are several types of DBMS architectures:

- **Centralized DBMS**: All components are on a **single server**.
- **Client-Server DBMS**: Clients **access** the database through a **server**.
- **Distributed DBMS**: Data is distributed across **multiple servers**.
- **Cloud DBMS**: Database services are provided over the **cloud**.
- **Hybrid DBMS** **Combines** features of centralized and distributed systems.
- Peer-to-Peer DBMS: Each node can act as both client and server.
- In memory DBMS: Data is stored in **RAM** for faster access.



Types of DBMS Architecture

There are several types of **DBMS** architectures:

- **Centralized DBMS**: All components are on a **single server**.
- **Client-Server DBMS**: Clients **access** the database through a **server**.
- **Distributed DBMS**: Data is distributed across **multiple servers**.
- **Cloud DBMS**: Database services are provided over the **cloud**.
- **Hybrid DBMS**: **Combines** features of centralized and distributed systems.
- **Peer-to-Peer DBMS**: Each node can act as a **client** and **server**.
- **In memory DBMS**: Data is stored in **RAM** for faster access.



Types of DBMS Architecture

There are several types of DBMS architectures:

- **Centralized DBMS**: All components are on a **single server**.
- **Client-Server DBMS**: Clients **access** the database through a **server**.
- **Distributed DBMS**: Data is distributed across **multiple servers**.
- **Cloud DBMS**: Database services are provided over the **cloud**.
- **Hybrid DBMS**: **Combines** features of centralized and distributed systems.
- **Peer-to-Peer DBMS**: Each node can act as a **client** and **server**.
- **In memory DBMS**: Data is stored in **RAM** for faster access.

DB - 100-200ms

Cache - 5-10 ms

Cache
Lifetime (sec-min)

max 30
hours



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
- **Database administrators (DBAs)** are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (DBAs) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (DBAs) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (DBAs) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (DBAs) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (**DBAs**) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- Database administrators (**DBAs**) are **essential** for modern organizations.



Database System Administration

- **Database system administration** is the discipline of **managing, configuring, and maintaining database systems** to ensure their reliability, performance, security, and availability.
- Key responsibilities include:
 - Data storage management
 - System architecture design
 - Transaction and query processing
 - Concurrency control
 - Backup and recovery
 - Security and access control
- **Database administrators (DBAs)** are **essential** for modern organizations.



Record Storage Concepts

- A **record** (or **row/tuple**) is the **basic unit** of data storage in a *database table*.
- Efficient record storage is crucial for fast data retrieval and **update**.
- Storage techniques:
 - Heap storage: Unordered, Fast inserts.
 - B-tree storage: Ordered, Fast range queries.
 - Hash indexes for fast lookups.



Record Storage Concepts

- A **record** (or **row/tuple**) is the **basic unit** of data storage in a *database table*.
- **Efficient record storage** is crucial for **fast data retrieval** and **update**.
- Storage techniques:
 - Hash storage: Unordered, fast inserts.
 - Sequential storage: Ordered, fast range queries.
 - Using indexes for fast lookups.



Record Storage Concepts

- A **record** (or **row/tuple**) is the **basic unit** of data storage in a *database table*.
- **Efficient record storage** is crucial for **fast** data retrieval and **update**.
- Storage techniques:
 - **Heap storage**: Unordered, fast inserts.
 - **Sequential storage**: Ordered, fast range queries.
 - **Indexed storage**: Uses indexes for fast lookups.



Record Storage Concepts

- A **record** (or **row/tuple**) is the **basic unit** of data storage in a *database table*.
- **Efficient record storage** is crucial for **fast** data retrieval and **update**.
- Storage techniques:
 - **Heap storage**: Unordered, fast inserts.
 - **Sequential storage**: Ordered, fast range queries.
 - **Indexed storage**: Uses indexes for fast lookups.



Record Storage Concepts

- A **record** (or **row/tuple**) is the **basic unit** of data storage in a *database table*.
- **Efficient record storage** is crucial for **fast** data retrieval and **update**.
- Storage techniques:
 - **Heap storage**: Unordered, fast inserts.
 - **Sequential storage**: Ordered, fast range queries.
 - **Indexed storage**: Uses indexes for fast lookups.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- Block size and layout affect I/O performance.
- Records may be packed, slotted, or may span multiple blocks.
- Free space management is important for updates and inserts.
- Fragmentation management ensures efficient use of space.
- Data compression can also improve storage efficiency.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- **Block size** and **layout** affect **I/O performance**.
 - Records may be packed, slotted, or may span multiple blocks.
 - **Free space management** is important for **updates** and **inserts**.
 - **Fragmentation management** ensures efficient use of space.
 - **Data compression** can also improve storage efficiency.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- **Block size** and **layout** affect **I/O performance**.
- Records may be **packed**, **slotted**, or may **span multiple blocks**.
- Free space management is important for **updates** and **inserts**.
- Fragmentation management ensures efficient use of space.
- Data compression can also improve storage efficiency.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- **Block size** and **layout** affect **I/O performance**.
- Records may be **packed**, **slotted**, or may **span multiple blocks**.
- **Free space management** is important for **updates** and **inserts**.
- **Fragmentation management** ensures efficient use of space.
- **Data compression** can also improve storage efficiency.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- **Block size** and **layout** affect **I/O performance**.
- Records may be **packed**, **slotted**, or may **span multiple blocks**.
- **Free space management** is important for **updates** and **inserts**.
- **Fragmentation management** ensures efficient use of space.
- **Data compression** can also improve storage efficiency.



Block and Page Organization

- Data is stored in **blocks** (pages) on disk.
- **Block size** and **layout** affect **I/O performance**.
- Records may be **packed**, **slotted**, or may **span multiple blocks**.
- **Free space management** is important for **updates** and **inserts**.
- **Fragmentation management** ensures efficient use of space.
- **Data compression** can also improve storage efficiency.

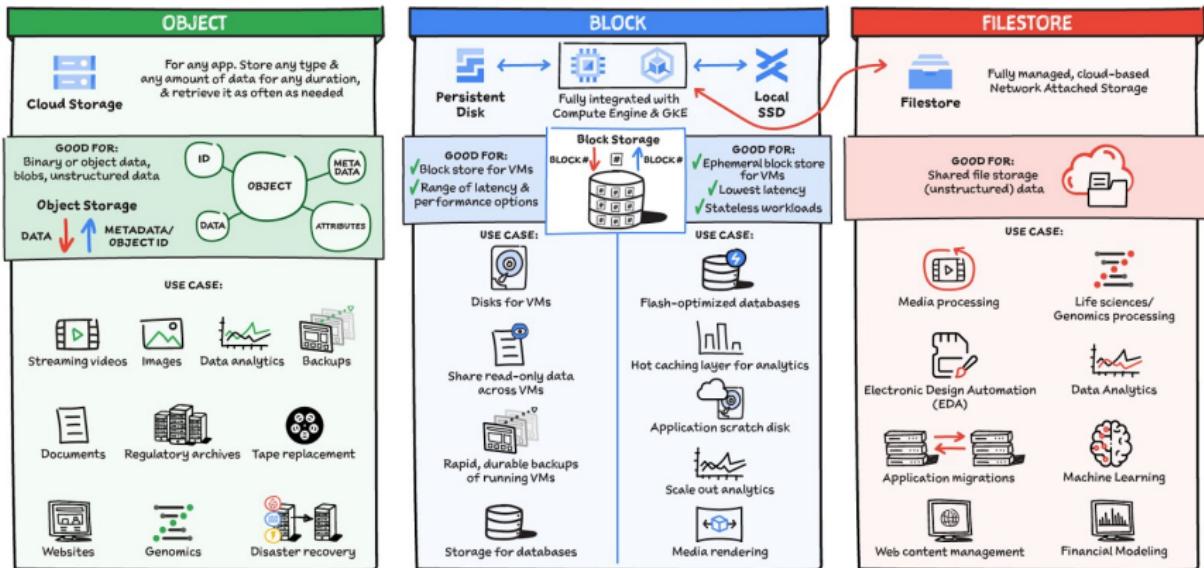


Record Storage: Use Cases

#GAPSketchnote
 @PVERGADIA
THECLOUDGIRL.DEV
 04.23.2021



Which Storage Should I Use?



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



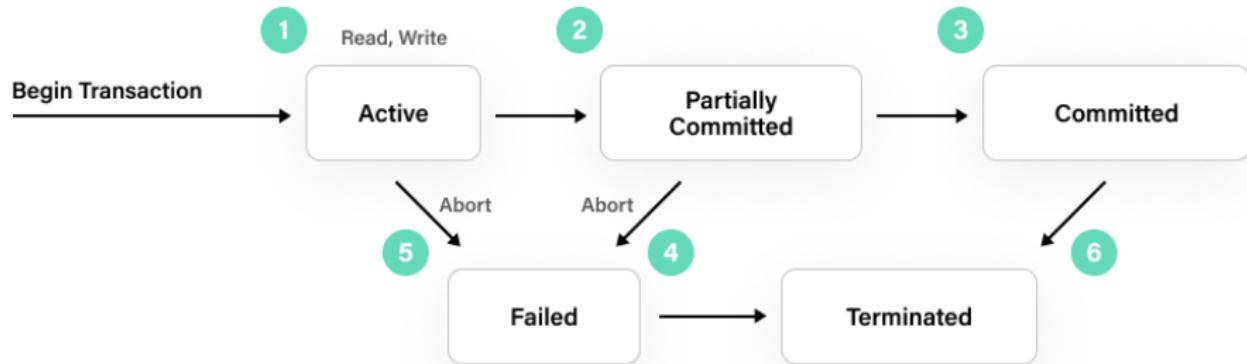
Transactional System Concepts

- A **transaction** is a sequence of operations performed as a **single logical unit of work**.
- Transactions must satisfy the **ACID** properties:
 - **Atomicity**: All or nothing.
 - **Consistency**: Preserves database integrity.
 - **Isolation**: Transactions do not interfere.
 - **Durability**: Results persist after completion.



Transaction Lifecycle

- **Begin:** Transaction starts.
- **Read/Write:** Operations are performed.
- **Commit:** Changes are made permanent.
- **Rollback:** Changes are undone if an error occurs.
- **Savepoints** can be used for partial rollbacks.



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery

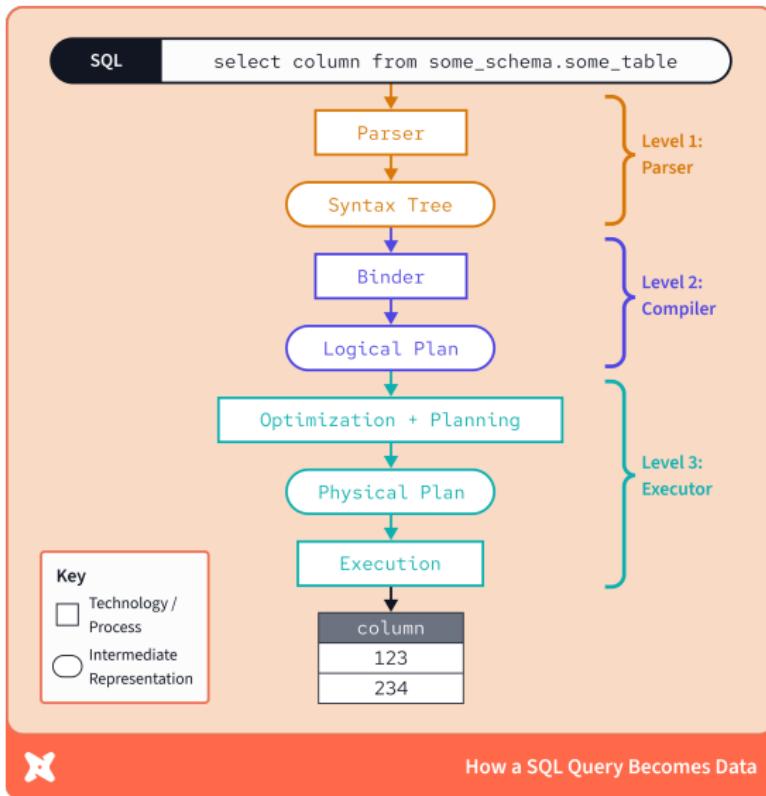


Query Execution Process

- **Query execution** is the process of **interpreting** and **running** database queries.
- Steps:
 - **Parsing**: Analyzing query syntax.
 - **Optimization**: Choosing the best execution plan.
 - **Execution**: Retrieving and processing data.
- **Efficient execution** is critical for **performance**.



Query Execution Flow: Full Transaction



How a SQL Query Becomes Data

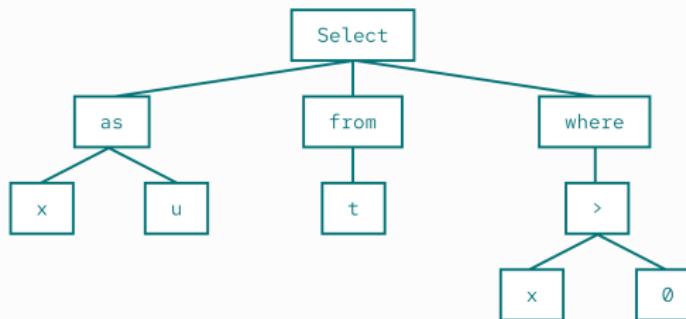
Query Execution Flow: Syntax Tree

SQL

select x as u from t where x > 0

Parser

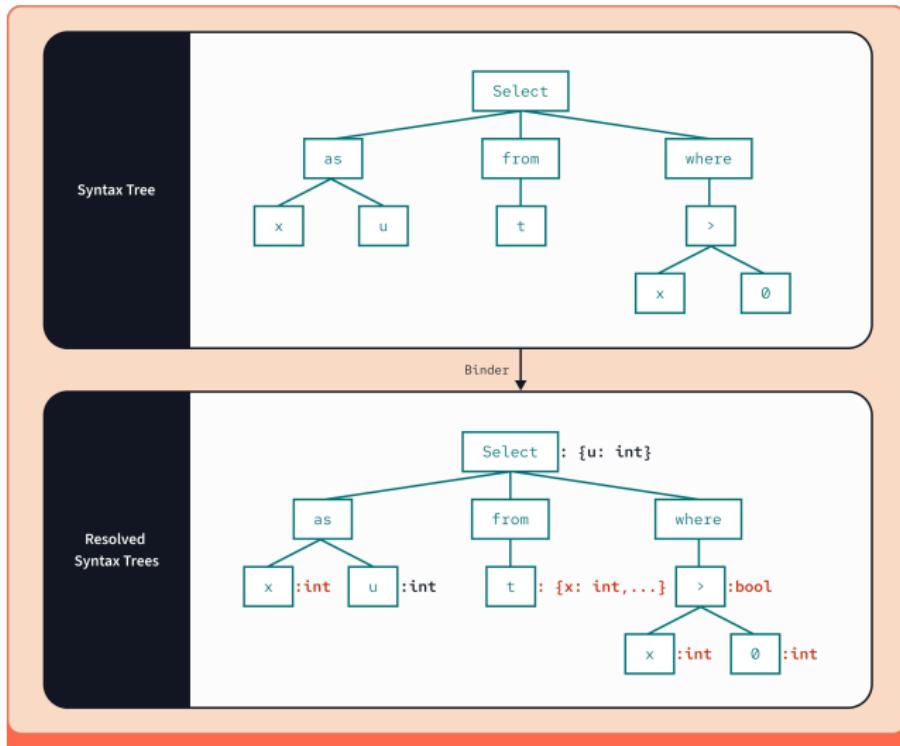
Syntax Tree



Parsers Recognize the Structure of the Query



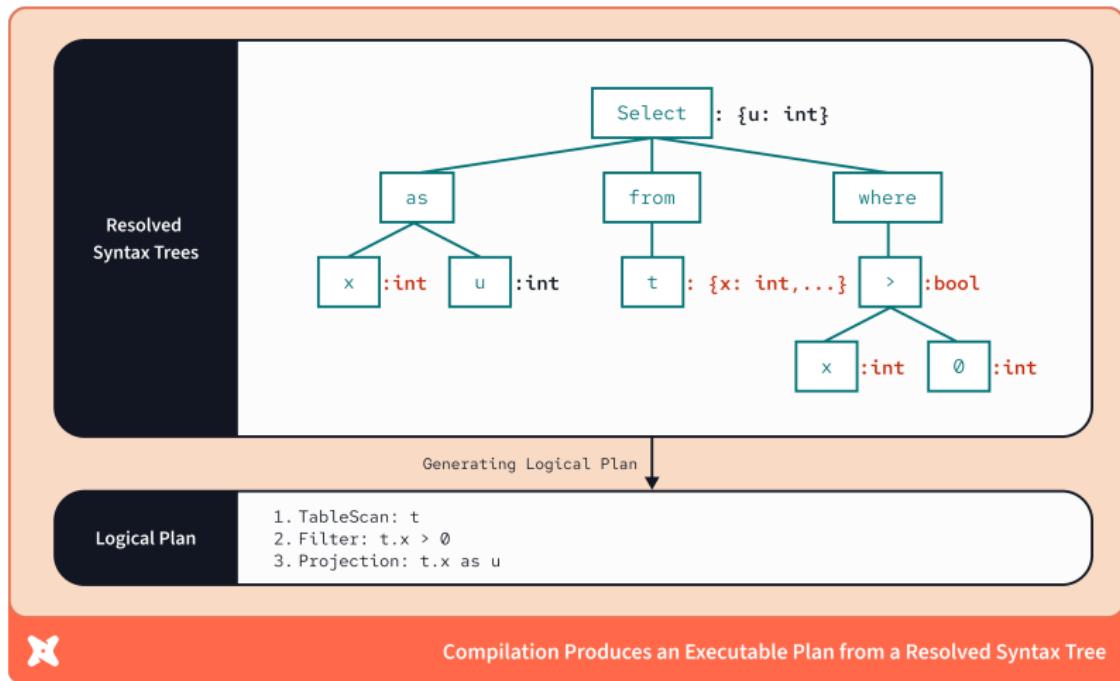
Query Execution Flow: Compilation



Binders Resolve Names and Check Types



Query Execution Flow: Logical Plan



Query Optimization

- The **query optimizer** selects the most efficient strategy for executing a query.
- Considers indexes, join methods, and data distribution.
- May rewrite queries for better performance.
- Cost-based and rule-based optimization approaches.



Query Optimization

- The **query optimizer** selects the most efficient strategy for executing a query.
- Considers **indexes**, **join methods**, and **data distribution**.
- May rewrite queries for better performance.
- Cost-based and rule-based optimization approaches.



Query Optimization

- The **query optimizer** selects the most efficient strategy for executing a query.
- Considers **indexes**, **join methods**, and **data distribution**.
- May **rewrite queries** for better performance.
- **Cost-based** and **rule-based** optimization approaches.

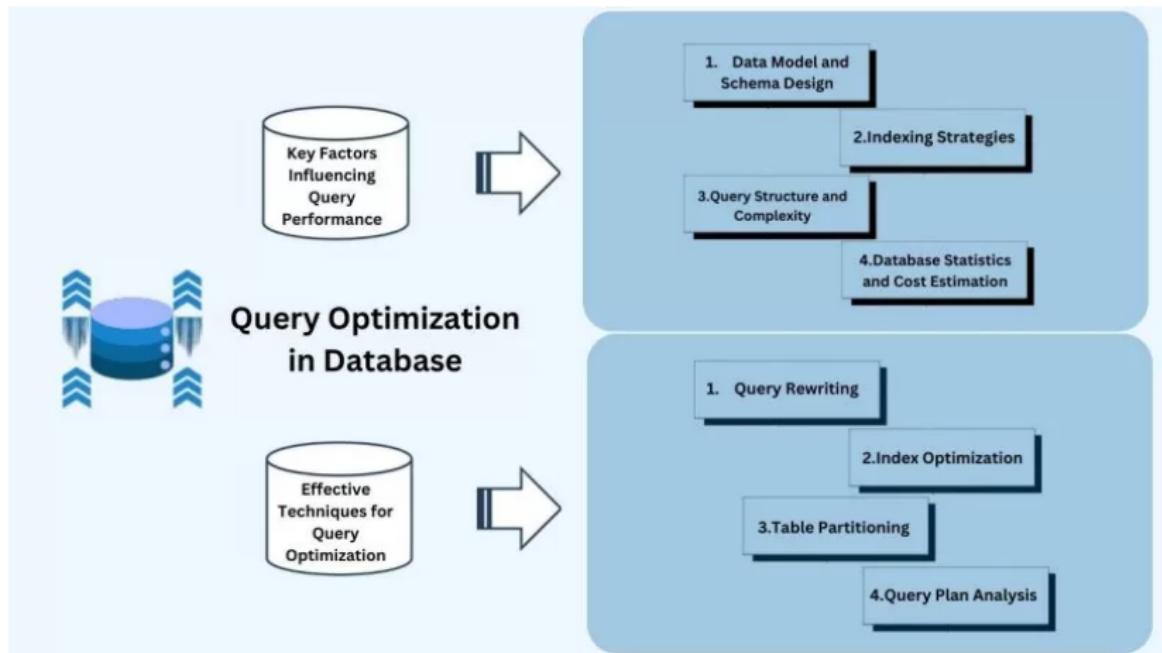


Query Optimization

- The **query optimizer** selects the most efficient strategy for executing a query.
- Considers **indexes**, join methods, and **data distribution**.
- May **rewrite queries** for better performance.
- **Cost-based** and **rule-based** optimization approaches.



Query Optimization Factors



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



Concurrency Control Concepts

- **Concurrency control** ensures correct results when **multiple transactions** execute **simultaneously**.
- Prevents problems like **lost updates**, **dirty reads**, and **deadlocks**.
- Techniques:
 - Locking protocols: Use locks to control access (e.g., two-phase locking).
 - Timestamp ordering: Assigns timestamps to transactions.
 - Optimistic concurrency control: Assumes no conflicts.
 - Pessimistic concurrency control: Assumes conflicts may occur.



Concurrency Control Concepts

- **Concurrency control** ensures correct results when **multiple transactions** execute **simultaneously**.
- Prevents problems like **lost updates**, **dirty reads**, and **deadlocks**.
- Techniques:
 - **Locking protocols**: Use locks to control access (e.g., two-phase locking).
 - **Timestamp ordering**: Assigns timestamps to transactions.
 - **Optimistic concurrency control**: Assumes no conflicts.
 - **Pessimistic concurrency control**: Assumes conflicts may occur.



Concurrency Control Concepts

- **Concurrency control** ensures correct results when **multiple transactions** execute **simultaneously**.
- Prevents problems like **lost updates**, **dirty reads**, and **deadlocks**.
- Techniques:
 - **Locking protocols**: Use locks to control access (e.g., two-phase locking).
 - **Timestamp ordering**: Assigns timestamps to transactions.
 - **Optimistic concurrency control**: Assumes no conflicts.
 - **Pessimistic concurrency control**: Assumes conflicts may occur.



Concurrency Control Concepts

- **Concurrency control** ensures correct results when **multiple transactions** execute **simultaneously**.
- Prevents problems like **lost updates**, **dirty reads**, and **deadlocks**.
- Techniques:
 - **Locking protocols**: Use locks to control access (e.g., two-phase locking).
 - **Timestamp ordering**: Assigns timestamps to transactions.
 - **Optimistic concurrency control**: Assumes no conflicts.
 - **Pessimistic concurrency control**: Assumes conflicts may occur.



Concurrency Control Concepts

- **Concurrency control** ensures correct results when **multiple transactions** execute **simultaneously**.
- Prevents problems like **lost updates**, **dirty reads**, and **deadlocks**.
- Techniques:
 - **Locking protocols**: Use locks to control access (e.g., two-phase locking).
 - **Timestamp ordering**: Assigns timestamps to transactions.
 - **Optimistic concurrency control**: Assumes no conflicts.
 - **Pessimistic concurrency control**: Assumes conflicts may occur.



Locking and Deadlocks

- **Locks** prevent **conflicting operations** on the same data.
- Deadlocks occur when transactions **wait indefinitely** for each other.
- DBMS uses deadlock detection and **resolution strategies**.
- **Isolation levels** (e.g., **Read Committed**, **Serializable**) control visibility of changes.



Locking and Deadlocks

- **Locks** prevent **conflicting operations** on the same data.
- **Deadlocks** occur when transactions **wait indefinitely** for each other.
- DBMS uses deadlock detection and **resolution strategies**.
- **Isolation levels** (e.g., **Read Committed**, **Serializable**) control visibility of changes.



Locking and Deadlocks

- **Locks** prevent **conflicting operations** on the same data.
- **Deadlocks** occur when transactions **wait indefinitely** for each other.
- DBMS uses **deadlock detection** and **resolution strategies**.
- **Isolation levels** (e.g., **Read Committed**, **Serializable**) control visibility of changes.



Locking and Deadlocks

- **Locks** prevent **conflicting operations** on the same data.
- **Deadlocks** occur when transactions **wait indefinitely** for each other.
- DBMS uses **deadlock detection** and **resolution strategies**.
- **Isolation levels** (e.g., **Read Committed**, **Serializable**) control visibility of changes.



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



Failure Recovery Concepts

- **Failure recovery** restores the database to a **consistent state** after a failure.
- Types of failures:
 - Transaction failure: Only one transaction fails.
 - System crash: All active transactions are lost.
 - Media failure: Disk or storage device fails.
- Recovery techniques:
 - Write-ahead logging (WAL):
 - Log all writes to disk before they are applied to the data pages.
 - If a failure occurs, roll back incomplete transactions and redo completed ones.



Failure Recovery Concepts

- **Failure recovery** restores the database to a **consistent state** after a failure.
- Types of failures:
 - **Transaction failure:** Only one transaction fails.
 - **System crash:** All active transactions are lost.
 - **Media failure:** Disk or storage device fails.
- Recovery techniques:
 - Write-ahead logging (WAL)
 - Checkpointing
 - Redo log
 - Undo log



Failure Recovery Concepts

- **Failure recovery** restores the database to a **consistent state** after a failure.
- Types of failures:
 - **Transaction failure:** Only one transaction fails.
 - **System crash:** All active transactions are lost.
 - **Media failure:** Disk or storage device fails.
- Recovery techniques:
 - **Write-ahead logging (WAL):**
 - Logs changes before applying them.
 - Ensures durability and atomicity.
 - **Checkpointing:**
 - Periodically saves the database state.
 - Reduces recovery time.
 - **Shadow paging:**
 - Maintains copies of data pages.
 - Allows quick recovery to a consistent state.



Failure Recovery Concepts

- **Failure recovery** restores the database to a **consistent state** after a failure.
- Types of failures:
 - **Transaction failure:** Only one transaction fails.
 - **System crash:** All active transactions are lost.
 - **Media failure:** Disk or storage device fails.
- Recovery techniques:
 - **Write-ahead logging (WAL):**
 - Logs changes before applying them.
 - Ensures durability and atomicity.
 - **Checkpointing:**
 - Periodically saves the database state.
 - Reduces recovery time.
 - **Shadow paging:**
 - Maintains copies of data pages.
 - Allows quick recovery to a consistent state.



Failure Recovery Concepts

- **Failure recovery** restores the database to a **consistent state** after a failure.
- Types of failures:
 - **Transaction failure:** Only one transaction fails.
 - **System crash:** All active transactions are lost.
 - **Media failure:** Disk or storage device fails.
- Recovery techniques:
 - **Write-ahead logging (WAL):**
 - Logs changes before applying them.
 - Ensures durability and atomicity.
 - **Checkpointing:**
 - Periodically saves the database state.
 - Reduces recovery time.
 - **Shadow paging:**
 - Maintains copies of data pages.
 - Allows quick recovery to a consistent state.



Write-Ahead Logging and Checkpoints

- **Write-ahead logging (WAL)**: All changes are **logged before** being applied to the database.
- **Checkpoints**: Periodically **save** the database state to speed up recovery.
- **Shadow paging**: Maintains **copies** of data **pages**.
- Recovery ensures atomicity and durability.



Write-Ahead Logging and Checkpoints

- **Write-ahead logging (WAL)**: All changes are **logged before** being applied to the database.
- **Checkpoints**: Periodically **save** the database **state** to speed up recovery.
- **Shadow paging**: Maintains **copies** of data **pages**.
- Recovery ensures atomicity and durability.



Write-Ahead Logging and Checkpoints

- **Write-ahead logging (WAL)**: All changes are **logged before** being applied to the database.
- **Checkpoints**: Periodically **save** the database **state** to speed up recovery.
- **Shadow paging**: Maintains **copies** of data **pages**.
- Recovery ensures atomicity and durability.



Write-Ahead Logging and Checkpoints

- **Write-ahead logging (WAL)**: All changes are **logged before** being applied to the database.
- **Checkpoints**: Periodically **save** the database **state** to speed up recovery.
- **Shadow paging**: Maintains **copies** of data **pages**.
- Recovery ensures atomicity and durability.



Outline

- 1 DataBases Infrastructure
- 2 DBMS Architecture
- 3 Database System Administration
- 4 Transactional System
- 5 Query Execution
- 6 Concurrency Control
- 7 Failure Recovery



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/databases-ii>

