

DEEP LEARNING

Advanced Topics: GANs, Transfer Learning & Continual Learning

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2026-I



UNIVERSIDAD DISTRITAL
FRANCISCO JOSÉ DE CALDAS

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning



Outline

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning



Generative vs. Discriminative Models

Discriminative

- Model $P(y \mid \mathbf{x})$
- Map input to *label*
- Classes 1–6 focus here
- Examples: ResNet, BERT

Generative

- Model $P(\mathbf{x})$ or $P(\mathbf{x} \mid y)$
- **Synthesise** new samples
- **Dense representation learning**
- Examples: VAE, GAN, Diffusion

Why Generative Models?

- Data augmentation for **rare classes**
- Anomaly detection via *likelihood*
- Creative applications (*image/text/audio synthesis*)
- Bridge to **Challenge 4** (GAIL): **adversarial reward** in RL



Generative vs. Discriminative Models

Discriminative

- Model $P(y \mid \mathbf{x})$
- Map input to *label*
- Classes 1–6 focus here
- Examples: ResNet, BERT

Generative

- Model $P(\mathbf{x})$ or $P(\mathbf{x} \mid y)$
- **Synthesise** new samples
- **Dense representation learning**
- Examples: VAE, GAN, Diffusion

Why Generative Models?

- Data augmentation for **rare classes**
- Anomaly detection via *likelihood*
- Creative applications (*image/text/audio synthesis*)
- Bridge to **Challenge 4** (GAIL): **adversarial reward** in RL




The GAN Framework I

Goodfellow *et al.*, 2014

Two networks play a **minimax game**:

- Generator G : maps *noise* $\mathbf{z} \sim p_z(\mathbf{z})$ to **fake data** $\tilde{\mathbf{x}} = G(\mathbf{z})$
- Discriminator D : outputs probability that input is **real**



The GAN Framework II


Value Function

$$\min_G \max_D \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))]$$

- **Nash Equilibrium:** G generates the **true data distribution** $\Rightarrow D(\mathbf{x}) = 0.5$ everywhere
- In practice: G maximises $\log D(G(\mathbf{z}))$ (**non-saturating loss**)



GAN Training Algorithm



① For each *training iteration*:

- Sample m real examples $\{\mathbf{x}^{(i)}\}$ from p_{data}
- Sample m noise vectors $\{\mathbf{z}^{(i)}\}$ from p_z

② **Update D :** maximise $\frac{1}{m} \sum_i [\log D(\mathbf{x}^{(i)}) + \log(1 - D(G(\mathbf{z}^{(i)})))]$

③ **Update G :** minimise $\frac{1}{m} \sum_i \log(1 - D(G(\mathbf{z}^{(i)})))$ (equivalently maximise $\log D(G(\mathbf{z}))$)



GAN Training Challenges

Mode Collapse

G finds a small subset of outputs that reliably fool D and **ignores the rest of the distribution**. Generated images look **similar to each other**.

Non-Convergence

D wins early (loss $\rightarrow 0$) $\Rightarrow$ gradient signal to G **vanishes**. Training **oscillates** or diverges.

Common Remedies

- **Label smoothing**: *real labels* = 0.9, not 1.0
- **Feature matching**: train G to match D 's **intermediate statistics**
- **Minibatch discrimination**: D penalises lack of **diversity**
- **Spectral normalisation**: constrain D 's **Lipschitz constant**



GAN Training Challenges

Mode Collapse

G finds a small subset of outputs that reliably fool D and **ignores the rest of the distribution**. Generated images look **similar to each other**.

Non-Convergence

D wins early (loss $\rightarrow 0$) $\Rightarrow$ gradient signal to G **vanishes**. Training **oscillates** or diverges.

Common Remedies

- **Label smoothing**: *real labels* = 0.9, not 1.0
- **Feature matching**: train G to match D 's **intermediate statistics**
- **Minibatch discrimination**: D penalises lack of **diversity**
- **Spectral normalisation**: constrain D 's **Lipschitz constant**



Wasserstein GAN

Arjovsky *et al.*, 2017 — *Earth Mover Distance*

Replace the **JS divergence** with the **Wasserstein-1** (*Earth Mover*) distance:

$$W(p_r, p_g) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim p_r}[f(x)] - \mathbb{E}_{x \sim p_g}[f(x)]$$

WGAN Objective

Critic: $\max_{\|f\|_L \leq 1} \mathbb{E}[f(\mathbf{x})] - \mathbb{E}[f(G(\mathbf{z}))]$

Generator: $\min_G -\mathbb{E}[f(G(\mathbf{z}))]$

• Wasserstein distance always provides meaningful gradients — no more vanishing signal

• *Weight clipping* (original) or *gradient penalty* (WGAN-GP) enforces Lipschitz constraint

• WGAN-GP: most reliable GAN variant in practice



Wasserstein GAN

Arjovsky *et al.*, 2017 — *Earth Mover Distance*

Replace the **JS divergence** with the **Wasserstein-1** (*Earth Mover*) distance:

$$W(p_r, p_g) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim p_r}[f(x)] - \mathbb{E}_{x \sim p_g}[f(x)]$$

WGAN Objective

Critic: $\max_{\|f\|_L \leq 1} \mathbb{E}[f(\mathbf{x})] - \mathbb{E}[f(G(\mathbf{z}))]$

Generator: $\min_G -\mathbb{E}[f(G(\mathbf{z}))]$

- Wasserstein distance **always provides meaningful gradients** — no more **vanishing signal**
- *Weight clipping* (original) or **gradient penalty** (WGAN-GP) enforces *Lipschitz* constraint



WGAN-GP: **most reliable** GAN variant in practice



Evaluating GANs

Fréchet Inception Distance (FID)

$$\text{FID} = \|\boldsymbol{\mu}_r - \boldsymbol{\mu}_g\|^2 + \text{Tr}\left(\boldsymbol{\Sigma}_r + \boldsymbol{\Sigma}_g - 2\sqrt{\boldsymbol{\Sigma}_r\boldsymbol{\Sigma}_g}\right)$$

Compare Inception feature statistics of *real* vs. *generated* images. **Lower is better.**

Inception Score (IS)

$\text{IS} = \exp(\mathbb{E}_{\mathbf{x}} [D_{\text{KL}}(p(y|\mathbf{x})||p(y))])$ **Higher is better**; measures *diversity* and *quality* simultaneously.

- FID is the standard benchmark (more reliable than IS)
- Visual inspection remains essential: FID alone can be gamed



Evaluating GANs

Fréchet Inception Distance (FID)

$$\text{FID} = \|\boldsymbol{\mu}_r - \boldsymbol{\mu}_g\|^2 + \text{Tr}\left(\boldsymbol{\Sigma}_r + \boldsymbol{\Sigma}_g - 2\sqrt{\boldsymbol{\Sigma}_r\boldsymbol{\Sigma}_g}\right)$$

Compare Inception feature statistics of *real* vs. *generated* images. **Lower is better.**

Inception Score (IS)

$\text{IS} = \exp(\mathbb{E}_{\mathbf{x}} [D_{\text{KL}}(p(y|\mathbf{x})||p(y))])$ **Higher is better**; measures *diversity* and *quality* simultaneously.

- FID is the **standard benchmark** (more reliable than IS)
- *Visual inspection* remains essential: FID alone can be gamed



Connection to Challenge 4 (GAIL)


Generative Adversarial Imitation Learning (Ho & Ermon, 2016)

GAIL applies the GAN framework to Reinforcement Learning:

- Generator $\equiv$ Policy π_θ : produces state-action trajectories
- Discriminator D_ϕ : distinguishes expert vs. policy trajectories
- Adversarial reward: $r_{\text{adv}}(s, a) = -\log(1 - D_\phi(s, a))$

Why GAN + RL?

GAIL learns to imitate without reward function — the discriminator replaces the environment's reward. This directly links the GAN material in this class to Challenge 4.



Connection to Challenge 4 (GAIL)

Generative Adversarial Imitation Learning (Ho & Ermon, 2016)

GAIL applies the GAN framework to **Reinforcement Learning**:

- Generator $\equiv$ Policy π_θ : produces **state-action trajectories**
- Discriminator D_ϕ : distinguishes **expert vs. policy** trajectories
- **Adversarial reward**: $r_{\text{adv}}(s, a) = -\log(1 - D_\phi(s, a))$

Why GAN + RL?

GAIL learns to imitate without a reward function — the discriminator **replaces** the environment's reward. This directly links the GAN material in this class to Challenge 4.



Practical Example: Vanilla GAN on MNIST

- Generator: *noise* $\mathbf{z} \in \mathbb{R}^{100} \rightarrow$ FC layers $\rightarrow$ image 28×28
- Discriminator: *image* $\rightarrow$ FC layers $\rightarrow$ scalar
- Track: *discriminator loss, generator loss, and FID every 1k steps*
- Observe: training *instability phases, mode collapse* if it occurs
- Experiment: compare JS-loss vs. WGAN-GP loss on *stability*
- Key insight: curves that diverge in JS-GAN often converge with WGAN-GP



Practical Example: Vanilla GAN on MNIST

- Generator: *noise* $\mathbf{z} \in \mathbb{R}^{100} \rightarrow$ FC layers $\rightarrow$ image 28×28
- Discriminator: *image* $\rightarrow$ FC layers $\rightarrow$ scalar
- **Track:** *discriminator loss, generator loss, and FID every 1k steps*
- **Observe:** training **instability phases, mode collapse** if it occurs
- **Experiment:** compare JS-loss vs. WGAN-GP loss on *stability*
- **Key insight:** curves that diverge in JS-GAN often converge with WGAN-GP



Practical Example: Vanilla GAN on MNIST

- Generator: *noise* $\mathbf{z} \in \mathbb{R}^{100} \rightarrow$ FC layers $\rightarrow$ image 28×28
- Discriminator: *image* $\rightarrow$ FC layers $\rightarrow$ scalar
- **Track:** *discriminator loss*, *generator loss*, and FID every $1k$ steps
- **Observe:** training **instability phases**, **mode collapse** if it occurs
- **Experiment:** compare JS-loss vs. WGAN-GP loss on *stability*
- **Key insight:** curves that diverge in JS-GAN often converge with WGAN-GP



DCGAN: Deep Convolutional GAN

Radford et al., 2015

Replace fully-connected layers with convolutional layers for image synthesis.

Generator


- FC $\rightarrow$ reshape to 4×4
- Transposed convolutions (upsample)
- BatchNorm after each conv
- ReLU hidden, tanh output

Discriminator

- Strided convolutions (downsample)
- LeakyReLU (slope 0.2)
- BatchNorm (except first/last)
- Sigmoid output

DCGAN $4 \times 4 \rightarrow 8 \times 8 \rightarrow 16 \times 16 \rightarrow 32 \times 32 \rightarrow$ final image

No pooling (use strided conv), BatchNorm everywhere except first and output layers.



DCGAN: Deep Convolutional GAN

Radford *et al.*, 2015

Replace **fully-connected** layers with **convolutional** layers for *image synthesis*.

Generator

- FC $\rightarrow$ reshape to 4×4
- **Transposed convolutions** (*upsample*)
- BatchNorm after each conv
- **ReLU hidden**, **tanh** output

Discriminator

- **Strided convolutions** (*downsample*)
- **LeakyReLU** (slope 0.2)
- BatchNorm (except *first/last*)
- **Sigmoid** output

DCGAN Rules of Thumb

No **pooling** (use strided conv), **BatchNorm** **everywhere** except **input** and **output** layers.



Conditional GAN (cGAN)


generate without $z \leftarrow \text{control}$

Mirza & Osindero, 2014

Condition both G and D on auxiliary information y (class label, text, etc.):

$$\min_G \max_D \mathbb{E}[\log D(\mathbf{x} | \mathbf{y})] + \mathbb{E}[\log(1 - D(G(\mathbf{z} | \mathbf{y})))]$$

- y concatenated with z in G and with feature maps in D
- Enables class-conditional generation: generate a specific digit, object, etc.
- Foundation for image-to-image translation: condition on an entire input image



Pix2Pix: Image-to-Image Translation

Isola et al., 2017

Learn a mapping from input image x to output image y (paired training data required):


$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{cGAN}}(G, D) + \lambda \mathcal{L}_{L_1}(G)$$

- The L_1 term forces pixel-level faithfulness to the input
- D operates on image patches (PatchGAN) rather than full images
- Applications:** sketch $\rightarrow$ photo, semantic map $\rightarrow$ scene, day $\rightarrow$ night, grayscale $\rightarrow$ colour

Fidelity

style

Cha. 7B



CycleGAN: Unpaired Image Translation

Zhu et al., 2017

No paired training data. Two generators $G_{X \rightarrow Y}$ and $F_{Y \rightarrow X}$ with cycle-consistency loss:

$$\mathcal{L}_{\text{cyc}} = \mathbb{E}[\|F(G(\mathbf{x})) - \mathbf{x}\|_1] + \mathbb{E}[\|G(F(\mathbf{y})) - \mathbf{y}\|_1]$$

- Cycle-consistency: one translation followed by its inverse recovers the original
- **Applications**: photo ↔ painting, horse ↔ zebra, summer ↔ winter
- **Connection to Challenge 7 Part B**: the neural style transfer operation is related to — but *distinct from* — CycleGAN



Progressive and Style GANs

ProGAN (Karras et al., 2018)

Train at progressively larger resolutions:

$4 \times 4 \rightarrow 8 \times 8 \rightarrow \dots \rightarrow 1024 \times 1024$. Add new Conv layers as resolution increases. Stable training, photorealistic faces.

StyleGAN (Karras et al., 2019)

- Latent code $z \rightarrow$ mapping network $\rightarrow w \in \mathcal{W}$
- w controls per-layer style via Adaptive Instance Normalisation (AdaIN)
- Style mixing: interpolate styles at different layers for controllable synthesis
- Unprecedented photorealism; enabled deepfake research and detection



Progressive and Style GANs

ProGAN (Karras et al., 2018)

Train at progressively larger resolutions:

$4 \times 4 \rightarrow 8 \times 8 \rightarrow \dots \rightarrow 1024 \times 1024$. Add new Conv layers as resolution increases. Stable training, photorealistic faces.

StyleGAN (Karras et al., 2019)


- Latent code $\mathbf{z} \rightarrow$ mapping network $\rightarrow \mathbf{w} \in \mathcal{W}$
- $\mathbf{w}$ controls per-layer style via Adaptive Instance Normalisation (AdaIN) $\rightarrow$ low resolution
- Style mixing: interpolate styles at different layers for controllable synthesis $\rightarrow$ textures
- Unprecedented photorealism: enabled deepfake research and detection



Diffusion Models (*Modern Context*)

Denoising Diffusion Probabilistic Models (DDPM)

A newer paradigm that outperforms GANs on image quality (Ho et al., 2020). Gradually add Gaussian noise to data, then learn to reverse the diffusion.



Dall-E
Midjourney

HD images

↳ slower than GAN

Diffusion Models (*Modern Context*)

Denoising Diffusion Probabilistic Models (DDPM)

A **newer paradigm** that **outperforms** GANs on **image quality** (Ho *et al.*, 2020). Gradually add **Gaussian noise** to data, then learn to **reverse** the *diffusion*.

- Training objective: predict the added noise at each step
- Stable training (no adversarial instability)
 - Enables: DALL-E 2, Stable Diffusion, Midjourney
 - **Trade-off**: better sample quality; slower inference than GAN

Context for This Course

GANs remain the **conceptual foundation** and are **central** to Challenge 4 (GAIL). Diffusion models are the **frontier** — referenced here as *motivation for further study*.



Diffusion Models (*Modern Context*)

Denoising Diffusion Probabilistic Models (DDPM)

A **newer paradigm** that **outperforms** GANs on **image quality** (Ho *et al.*, 2020). Gradually add **Gaussian noise** to data, then learn to **reverse** the *diffusion*.

- Training objective: predict the **added noise** at each step
- Stable training (no **adversarial instability**)
- Enables: DALL-E 2, Stable Diffusion, Midjourney
- **Trade-off**: better sample quality; slower inference than GAN

Context for This Course

GANs remain the **conceptual foundation** and are **central** to Challenge 4 (GAIL). Diffusion models are the **frontier** — referenced here as *motivation for further study*.



Diffusion Models (*Modern Context*)

Denoising Diffusion Probabilistic Models (DDPM)

A **newer paradigm** that **outperforms** GANs on **image quality** (Ho *et al.*, 2020). Gradually add **Gaussian noise** to data, then learn to **reverse** the *diffusion*.

- Training objective: predict the **added noise** at each step
- Stable training (no **adversarial instability**)
- Enables: DALL-E 2, Stable Diffusion, Midjourney
- **Trade-off**: **better sample quality**; **slower inference** than GAN

Context for This Course

GANs remain the **conceptual foundation** and are **central** to Challenge 4 (GAIL). Diffusion models are the **frontier** — referenced here as *motivation for further study*.

images (videos)

GAN Evaluation Checklist

- **Loss curves:** generator and discriminator should be roughly **balanced**
- **FID:** compute on held-out real set vs. **10k** generated samples
- **Visual diversity:** plot a grid of 8×8 generated samples; check for *mode collapse*
- **Discriminator accuracy:** above **85%** early $\Rightarrow G$ is not learning *fast enough*
- **Interpolation:** linear interpolation in **z-space** should produce *smooth transitions*
- **For Challenge 4 (GAIL):** discriminator accuracy over training reveals how fast the policy matches the *expert occupancy measure*

FID < 50 ✓
< 10 ✓



Practical Example: DCGAN on CIFAR-10

- Generator: noise $\mathbf{z} \in \mathbb{R}^{128} \rightarrow \text{reshape} \rightarrow 4 \times 4 \rightarrow 4 \text{ ConvTranspose blocks} \rightarrow 32 \times 32 \times 3$
- Discriminator: $32 \times 32 \times 3 \rightarrow 4 \text{ strided Conv blocks} \rightarrow \text{scalar}$
- Experiment A: DCGAN with BCE loss; record FID at *intervals*
- Experiment B: swap to WGAN-GP loss; compare *stability*
- Experiment C: add *class labels* to produce a cGAN
- Expected: WGAN-GP < FID than standard DCGAN; cGAN generates *recognisable class-conditional* samples



Practical Example: DCGAN on CIFAR-10

- Generator: noise $\mathbf{z} \in \mathbb{R}^{128} \rightarrow \text{reshape} \rightarrow 4 \times 4 \rightarrow 4 \text{ ConvTranspose blocks} \rightarrow 32 \times 32 \times 3$
- Discriminator: $32 \times 32 \times 3 \rightarrow 4 \text{ strided Conv blocks} \rightarrow \text{scalar}$
- **Experiment A:** DCGAN with BCE loss; record FID at *intervals*
- **Experiment B:** swap to WGAN-GP loss; compare **stability**
- **Experiment C:** add *class labels* to produce a cGAN
- **Expected:** WGAN-GP < FID than standard DCGAN; cGAN generates **recognisable class-conditional** samples



Practical Example: DCGAN on CIFAR-10

- Generator: noise $\mathbf{z} \in \mathbb{R}^{128} \rightarrow \text{reshape} \rightarrow 4 \times 4 \rightarrow 4 \text{ ConvTranspose blocks} \rightarrow 32 \times 32 \times 3$
- Discriminator: $32 \times 32 \times 3 \rightarrow 4 \text{ strided Conv blocks} \rightarrow \text{scalar}$
- **Experiment A:** DCGAN with BCE loss; record FID at *intervals*
- **Experiment B:** swap to WGAN-GP loss; compare **stability**
- **Experiment C:** add *class labels* to produce a cGAN
- **Expected:** WGAN-GP < FID than standard DCGAN; cGAN generates **recognisable class-conditional** samples



Outline

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning



The Transfer Learning Paradigm

Core Idea

Reuse representations learned on a large *source task* (e.g. ImageNet, 1.2M images, 1000 classes) to solve a **target task** with much less data.

- **Why it works:** *low-level layers* learn **universal features** (edges, textures, shapes) that transfer across *vision tasks*
- **Data efficiency:** reach **high accuracy** with 50–500 examples per class vs. requiring *tens of thousands* from scratch
- **Compute efficiency:** *fine-tuning* takes *minutes to hours*; training from scratch takes *days to weeks* on a GPU cluster



What is Transferred?

Layer-by-Layer Analysis

- **Early layers (layers 1–3):** Gabor-like edge detectors, colour blobs — highly transferable across all *visual tasks*
- **Middle layers (layers 4–6):** texture patterns, object parts — moderately transferable; more *task-specific*
- **Top layers (layers 7+):** class-specific concept detectors — least transferable, replace for new tasks

• Visualisation: maximise activation of individual neurons (feature visualisation)

• Practical rule: the more different source and target domains are, the fewer layers to transfer




What is Transferred?

Layer-by-Layer Analysis

- **Early layers (layers 1–3):** Gabor-like edge detectors, colour blobs — highly transferable across all *visual tasks*
- **Middle layers (layers 4–6):** texture patterns, object parts — moderately transferable; more *task-specific*
- **Top layers (layers 7+):** class-specific concept detectors — least transferable; replace for *new tasks*
- **Visualisation:** maximise activation of individual neurons (feature visualisation)
- **Practical rule:** the more different source and target domains are, the fewer layers to *transfer*



Two Transfer Strategies I




Strategy 1 — Feature Extraction (*Frozen Backbone*)

Freeze all pretrained layers; replace and train only the new classification head.

- Fast: only head is updated
- Works well when source and target domains are similar
- Recommended when target dataset is small (< 1k images/class)



Two Transfer Strategies II



Strategy 2 — Fine-Tuning

Unfreeze some or all pretrained layers; train with a small learning rate (e.g. 10^{-5} – 10^{-4}) to avoid catastrophic overwriting.

- Slower; requires more data; risks overfitting on small datasets
- Best results when source and target domains differ substantially



When to Use Which Strategy

	Similar Domain	Different Domain
Small dataset	Feature extraction (high accuracy)	Feature extraction (reasonable accuracy)
Large dataset	Fine-tune few top layers	Fine-tune full network (with careful LR)

Practical Guideline for Challenge 7

Start with frozen backbone; if validation accuracy *plateaus*, progressively unfreeze *top residual blocks*. Use smaller LR for deeper layers (differential LR).



Popular Pretrained Backbones

- ResNet-50 / ResNet-101: standard baseline; strong accuracy–speed trade-off
- EfficientNet-B0 to B7: compound scaling; mobile-friendly; comparable or better than ResNet with fewer parameters
- MobileNet-V3: designed for edge devices; 3–5× faster than ResNet-50
- ViT (Vision Transformer): patch-based; excels when data is large; less inductive bias
- CLIP (OpenAI): text+image pretraining; zero-shot classification possible

For Challenge 7

Use ResNet-50 as the primary backbone; compare with EfficientNet-B0 as an alternative. Both available in `torchvision.models` with IMAGENET1K_V2 weights.

Popular Pretrained Backbones

- ResNet-50 / ResNet-101: standard baseline; strong accuracy–speed trade-off
- EfficientNet-B0 **to** B7: *compound scaling*; mobile-friendly; comparable or better than ResNet with fewer parameters
- MobileNet-V3: designed for *edge devices*; $3\text{--}5\times$ faster than ResNet-50
- ViT (**Vision Transformer**): *patch-based*; excels when data is large; less inductive bias
- CLIP (**OpenAI**): *text+image pretraining*; zero-shot classification possible

For Challenge 7

Use ResNet-50 as the primary backbone; compare with EfficientNet-B0 as an alternative. Both available in torchvision.models with IMAGENET1K_V2 weights.

Few-Shot Learning

Setup

Learn a new task from only K labelled examples per class ($K = \underline{1}, \underline{5}, \underline{10}, \underline{50}$). This is **impossible** for training from scratch but **feasible** with pretrained backbones.

Prototypical Networks (Snell et al., 2017)

Classify by distance to class prototypes in embedding space:

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}, y) \in S_k} f_{\theta}(\mathbf{x}), \quad p(y = k | \mathbf{x}) = \frac{\exp(-d(f_{\theta}(\mathbf{x}), \mathbf{c}_k))}{\sum_j \exp(-d(f_{\theta}(\mathbf{x}), \mathbf{c}_j))}$$

- f_{θ} is the **pretrained backbone** (fixed)
- No **fine-tuning** needed: the frozen features are already **discriminative**

near
 $\begin{matrix} & C_1 & \\ \swarrow & & \searrow \\ C_2 & & C_3 \end{matrix}$
 $\begin{matrix} \times \\ \times \\ \times \end{matrix}$



Few-Shot Learning

Setup

Learn a new task from only K *labelled examples* per class ($K = 1, 5, 10, 50$). This is **impossible** for training from scratch but **feasible** with pretrained backbones.

Prototypical Networks (Snell et al., 2017)

Classify by **distance to class prototypes** in *embedding space*:

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}, y) \in S_k} \underline{f_\theta(\mathbf{x})}, \quad p(y = k \mid \mathbf{x}) = \frac{\exp(-d(f_\theta(\mathbf{x}), \mathbf{c}_k))}{\sum_j \exp(-d(f_\theta(\mathbf{x}), \mathbf{c}_j))}$$

- f_θ is the **pretrained backbone** (*fixed*)
- No *fine-tuning* needed: the *frozen features* are already **discriminative**



ImageNet Preprocessing Standard

Required for All Pretrained PyTorch Models

$$\text{Normalise}(\mathbf{x}) = \frac{\mathbf{x} - \mu}{\sigma}$$

$$\mu = [\underline{0.485}, \underline{0.456}, \underline{0.406}], \quad \sigma = [\underline{0.229}, \underline{0.224}, \underline{0.225}]$$

→ ImageNet

- Resize to 224×224 (for ResNet, EfficientNet, ViT families)
- Apply the same normalisation to target domain images — a common mistake in Challenge 1 is forgetting this step for cartoon/sketch images
- Training augmentations: random horizontal flip, random crop, color jitter → overfitting
- Validation/test: deterministic centre crop at 224×224; no augmentation



Challenge 7 Part A Connection

Few-Shot Classification on DomainNet (Part A)

- **Budget:** only 50 labelled images/class $\times$ 6 classes = 300 training images
- **Backbone:** ResNet-50 with ImageNet weights
- **Strategies:** frozen backbone vs. fine-tuning vs. from-scratch baseline
- **Domains:** each group assigned a *source* $\rightarrow$ *target* pair (Real, Clipart, Sketch, Painting)

Expected Outcome

From-scratch will severely underfit. Frozen backbone will already achieve strong accuracy. Fine-tuning (if done *carefully*) gives the best result. The gap demonstrates the value of pretraining in low-data regimes.



Practical Example: ResNet-50 Feature Extraction and Fine-Tuning

- **Setup:** DomainNet “Real” photos, 6 animal categories, 50 images/class
- **Run A:** frozen backbone + linear head; train 30 epochs, $LR = 10^{-3}$
- **Run B:** unfreeze layer4 + head; train 30 epochs, $LR = 10^{-4}$
- **Run C:** full fine-tune; $LR = 5 \times 10^{-5}$; risk of overfitting
- **Metrics:** top-1 accuracy on source test set; confusion matrix
- **Visualisation:** t-SNE of backbone features — do classes separate even without *fine-tuning*?



Practical Example: ResNet-50 Feature Extraction and Fine-Tuning

- **Setup:** DomainNet “Real” *photos*, 6 **animal categories**, 50 images/class
- **Run A:** **frozen backbone** + *linear head*; train 30 epochs, $\text{LR} = 10^{-3}$
- **Run B:** unfreeze layer4 + head; train 30 epochs, $\text{LR} = 10^{-4}$
- **Run C:** **full fine-tune**; $\text{LR} = 5 \times 10^{-5}$; **risk of overfitting**
- **Metrics:** top-1 accuracy on *source test set*; confusion matrix
- **Visualisation:** t-SNE of *backbone features* — do classes separate even without *fine-tuning*?



Practical Example: ResNet-50 Feature Extraction and Fine-Tuning

- **Setup:** DomainNet “Real” *photos*, 6 **animal categories**, 50 images/class
- **Run A:** **frozen backbone** + *linear head*; train 30 epochs, $LR = 10^{-3}$
- **Run B:** unfreeze layer4 + head; train 30 epochs, $LR = 10^{-4}$
- **Run C:** **full fine-tune**; $LR = 5 \times 10^{-5}$; **risk of overfitting**
- **Metrics:** **top-1 accuracy** on *source test set*; **confusion matrix**
- **Visualisation:** t-SNE of **backbone features** — do classes **separate** even without *fine-tuning*?



Outline

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning




Neural Style Transfer

Gatys, Ecker & Bethge, 2015

Synthesise a new image that has the **content** of one image and the **style** of another — **without training any network**. We **optimise the pixel values** of the **generated image**.

- **Content**: encoded in **deep VGG-19 feature maps** (relu4_2)
- **Style**: encoded in **Gram matrices** of **shallow-to-mid** feature maps (relu1_1, relu2_1, relu3_1, relu4_1)
- Use a **frozen** VGG-19; the **network** is a **differentiable** feature extractor only



Gram Matrix: Encoding Style

↳ style

Definition

For feature map $F^l \in \mathbb{R}^{C_l \times H_l \times W_l}$, the Gram matrix is:

$$G_{ij}^l = \frac{1}{C_l H_l W_l} \sum_k F_{ik}^l F_{jk}^l$$



$i \sim j \rightarrow$ pattern

- Captures pairwise correlations between feature channels
- Insensitive to spatial layout — it is a pure texture descriptor
- Style loss: $\mathcal{L}_{\text{style}} = \sum_l w_l \|G^l - \hat{G}^l\|_F^2$ where $\hat{G}^l$ is the target style image's Gram matrix

Firsts layer $\rightarrow$ colors + textures (at level)
 middle $\rightarrow$ image patterns + texture
 texture



Style Transfer Total Loss and Optimisation

Total Loss

$$\mathcal{L}_{\text{total}} = \alpha \mathcal{L}_{\text{content}} + \beta \mathcal{L}_{\text{style}}$$

→ texture

$$\mathcal{L}_{\text{content}} = \frac{1}{2} \sum_l \|F^l - P^l\|^2$$

ADAM

- Typical: $\alpha = 1$, $\beta = 10^3 - 10^6$ (*style usually dominates*)
- **Optimiser:** L-BFGS (*quasi-Newton*) converges in ~300 iterations
- **Initialisation:** start from the content image for stable results
- **Output:** a new image with content + style; directly serves as synthetic target-domain data in Challenge 7 Part C

70
70/30



Domain Shift

Definition

Domain shift (covariate shift) occurs when the **input distribution** changes between **training** (**source**) and **deployment** (**target**):

$$p_S(\mathbf{x}) \neq p_T(\mathbf{x}) \text{ yet } p_S(y|\mathbf{x}) \approx p_T(y|\mathbf{x})$$

- **Domain shift penalty:** $\Delta_{\text{shift}} = \text{Acc}_{\text{source}} - \text{Acc}_{\text{target}}$
- Real examples: scanner A vs. scanner B (medical imaging), daylight vs. night (autonomous driving), simulation vs. real robot
- **Challenge 7:** *Real photos* $\rightarrow$ *Clipart / Sketch / Painting* — a clean controlled domain shift scenario
- Measure Δ_{shift} for both the *frozen backbone* and *fine-tuned model*



Domain Shift

Definition


Domain shift (**covariate shift**) occurs when the **input distribution** changes between *training* (**source**) and *deployment* (**target**):

$$p_S(\mathbf{x}) \neq p_T(\mathbf{x}) \text{ yet } p_S(y|\mathbf{x}) \approx p_T(y|\mathbf{x})$$


- **Domain shift penalty:** $\Delta_{\text{shift}} = \text{Acc}_{\text{source}} - \text{Acc}_{\text{target}}$
- Real examples: *scanner A* vs. *scanner B* (medical imaging), *daylight* vs. *night* (autonomous driving), *simulation* vs. *real robot*
- **Challenge 7:** *Real photos* $\rightarrow$ *Clipart* / *Sketch* / *Painting* — a clean controlled domain shift scenario
- Measure Δ_{shift} for both the *frozen backbone* and *fine-tuned model*




Three Adaptation Strategies I (Challenge 7 Part C)



$$\Delta_{\text{shift}} = \text{large}$$



$$\Delta_{\text{shift}} = \text{medium}$$



$$\Delta_{\text{shift}} = \text{small}$$



Three Adaptation Strategies II (Challenge 7 Part C)

- 1 **Baseline** (*no adaptation*): apply the Part A **source-trained** classifier directly to target domain; measure Δ_{shift}
- 2 **Target fine-tuning** with *few labels*: *fine-tune* the *top layers* using a *small set* of labelled *target-domain* examples ($\sim 50/\text{class}$); compare with Part A *fine-tuning*
- 3 **Synthetic data augmentation** (using Part B output): mix *style-transferred images* into the *source training set*; *retrain* and evaluate on *target domain*

Expected Ranking

Synthetic augmentation $\geq$ Target fine-tuning $\gg$ No adaptation (though results vary by group and domain pair)



Three Adaptation Strategies II (Challenge 7 Part C)

- 1 **Baseline** (*no adaptation*): apply the Part A **source-trained** classifier directly to target domain; measure Δ_{shift}
- 2 **Target fine-tuning** with *few labels*: *fine-tune* the *top layers* using a **small set** of labelled *target-domain* examples ($\sim 50/\text{class}$); compare with Part A *fine-tuning*
- 3 **Synthetic data augmentation** (using Part B output): mix style-transferred images into the *source training set*; *retrain* and evaluate on *target domain*

Expected Ranking

Synthetic augmentation $\geq$ Target fine-tuning $\gg$ No adaptation (though results vary by group and domain pair)



Three Adaptation Strategies II (Challenge 7 Part C)

- 1 **Baseline** (*no adaptation*): apply the Part A **source-trained** classifier directly to target domain; measure Δ_{shift}
- 2 **Target fine-tuning** with *few labels*: *fine-tune* the *top layers* using a **small set** of labelled *target-domain* examples ($\sim 50/\text{class}$); compare with Part A *fine-tuning*
- 3 **Synthetic data augmentation** (using Part B output): mix style-transferred images into the source training set; retrain and evaluate on *target domain*

Expected Ranking

Synthetic augmentation $\geq$ Target fine-tuning $\gg$ No adaptation (though results vary by group and domain pair)



Three Adaptation Strategies II (Challenge 7 Part C)

- 1 **Baseline** (*no adaptation*): apply the Part A **source-trained** classifier directly to target domain; measure Δ_{shift}
- 2 **Target fine-tuning** with *few labels*: *fine-tune* the *top layers* using a **small set** of labelled *target-domain* examples ($\sim 50/\text{class}$); compare with Part A *fine-tuning*
- 3 **Synthetic data augmentation** (using Part B output): mix **style-transferred images** into the *source training set*; *retrain* and evaluate on *target domain*

Expected Ranking

Synthetic augmentation $\geq$ *Target fine-tuning* $\gg$ **No adaptation** (though results vary by *group* and *domain pair*)



Maximum Mean Discrepancy (MMD) complements

Measuring Domain Shift in Feature Space

Δ -shift

$$\text{MMD}^2(p, q) = \|\mathbb{E}_{\mathbf{x} \sim p} [\phi(\mathbf{x})] - \mathbb{E}_{\mathbf{y} \sim q} [\phi(\mathbf{y})]\|_{\mathcal{H}}^2$$

where ϕ maps to a Reproducing Kernel Hilbert Space (typically Gaussian RBF kernel).

- Compute MMD between source and target feature embeddings from the backbone
- High MMD: hard domain shift; adaptation more difficult
- Track MMD before and after adaptation to quantify the alignment improvement
- Complement Δ_{shift} (accuracy-based) with MMD (distribution-based)



Domain-Adversarial Neural Networks (DANN)

Ganin et al., 2016 — Gradient Reversal

Add a **domain classifier head** and a **gradient reversal layer (GRL)**. During *backprop*, the GRL multiplies the gradient by $-\lambda$, making the *feature extractor domain-invariant*.

Combined Loss

$$\mathcal{L} = \mathcal{L}_{\text{class}} - \lambda \mathcal{L}_{\text{domain}}$$

- $\mathcal{L}_{\text{class}}$: *cross-entropy on labelled source samples*
- $\mathcal{L}_{\text{domain}}$: *binary CE on source/target domain labels (0/1)*
- GRL forces features that **confuse** the *domain classifier* $\Rightarrow$ **domain-invariant**
- Optional in Challenge 7; recommended for groups seeking *distinct*



Domain-Adversarial Neural Networks (DANN)

Ganin *et al.*, 2016 — Gradient Reversal

Add a **domain classifier head** and a **gradient reversal layer (GRL)**. During *backprop*, the GRL multiplies the gradient by $-\lambda$, making the *feature extractor* **domain-invariant**.

Combined Loss

$$\mathcal{L} = \mathcal{L}_{\text{class}} - \lambda \mathcal{L}_{\text{domain}}$$

- $\mathcal{L}_{\text{class}}$: *cross-entropy* on *labelled source* samples
- $\mathcal{L}_{\text{domain}}$: binary CE on source/target *domain labels* (0/1)
- GRL forces features that **confuse** the *domain classifier* $\Rightarrow$ **domain-invariant**
- **Optional** in Challenge 7; recommended for groups seeking *distinction*



Practical Example: Challenge 7 Full Pipeline

- **Data:** DomainNet Real (*source*) $\rightarrow$ Clipart (*target*), 6 classes
- **Part A:** frozen ResNet-50 on 50 img/class; accuracy on *source test set*
- **Part B:** Style-transfer 30 images/class with *target-domain style*
- **Part C:** evaluate *source model* on *target*; record Δ_{shift}
- **Adaptation:** retrain with *synthetic data*; compare to *no-adaptation baseline*
- **Visualise:** UMAP of *source* and *target features* before and *after adaptation* (*source* = blue, *target* = red); *clusters* should *mix after adaptation*



Practical Example: Challenge 7 Full Pipeline

- **Data:** DomainNet Real (*source*) $\rightarrow$ Clipart (*target*), 6 classes
- **Part A:** **frozen** ResNet-50 on 50 img/class; **accuracy** on **source test set**
- **Part B:** **Style-transfer** 30 images/class with **target-domain style**
- **Part C:** **evaluate source model** on *target*; record Δ_{shift}
- **Adaptation:** **retrain** with **synthetic data**; compare to **no-adaptation baseline**
- **Visualise:** UMAP of **source** and **target features** before and *after adaptation* (*source* = blue, *target* = red); **clusters** should **mix after adaptation**



Practical Example: Challenge 7 Full Pipeline

- **Data:** DomainNet Real (*source*) → Clipart (*target*), 6 classes
- **Part A:** *frozen* ResNet-50 on 50 img/class; *accuracy* on *source test set*
- **Part B:** *Style-transfer* 30 images/class with *target-domain style*
- **Part C:** *evaluate source model* on *target*; record Δ_{shift}
- **Adaptation:** *retrain* with *synthetic data*; compare to *no-adaptation baseline*
- **Visualise:** UMAP of *source* and *target features* before and after *adaptation* (*source* = blue, *target* = red); *clusters* should *mix after adaptation*



Outline

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning



The Online Learning Paradigm

Stream vs. Batch

wrong inputs

In **online learning**, data arrives as a stream: each *sample* (or *mini-batch*) is seen once, then *discarded*. The model must **adapt without storing** all past data.

- **Motivation:** financial fraud detection, sensor streams, recommendation systems, robotics with a continuously changing environment
- **Key constraint:** memory budget $\ll$ full training set
- **Key challenge:** distribution shift over time; older observations may no longer reflect the current distribution



Catastrophic Forgetting

forget (overwrite)

McCloskey & Cohen, 1989 -- The Stability-Plasticity Dilemma

When a neural network is trained **sequentially** on Task A, then Task B, it **forgets** Task A. New gradient updates **overwrite** weights **critical** for Task A.

- **Example:** a ResNet trained on Cats/Dogs, then retrained on Cats, will score **near chance** on Dogs
 → Forget
- The network must be simultaneously **stable** (*retain old knowledge*) and **plastic** (*acquire new knowledge*)
- **Connection to RL:** switching Atari games between Challenges 1, 3, 4 triggers **catastrophic forgetting** if the same network is reused



Catastrophic Forgetting

McCloskey & Cohen, 1989 — The Stability--Plasticity Dilemma

When a neural network is trained **sequentially** on *Task A*, then *Task B*, it **forgets** *Task A*. New gradient updates **overwrite** weights **critical** for *Task A*.

- **Example:** a ResNet trained on *Cats/Dogs*, then retrained on *Cars*, will score **near-chance** on *Dogs*
- The network must be simultaneously **stable** (*retain old knowledge*) and **plastic** (*acquire new knowledge*)
- **Connection to RL:** switching Atari games between Challenges 1, 3, 4 triggers **catastrophic forgetting** if the same network is reused



Three Approaches to Continual Learning

Regularisation

- Penalise changes to important weights
- EWC, SI, MAS
- No extra memory for *data*

Replay

- Store or generate past examples
- Mix into current training
- ER, DER,
- Generative Replay

GAW

Architecture

- Allocate new capacity per *task*
- No forgetting by *design*
- PNN, PackNet
- Grows with tasks

↓
not
scalable



Elastic Weight Consolidation (EWC) I


Kirkpatrick *et al.*, 2017

After *Task A*, compute the **Fisher Information** F_i for each weight θ_i : F_i measures how much θ_i **contributed** to *Task A* performance. Then regularise training on *Task B*:

$$\mathcal{L}_{\text{EWC}} = \mathcal{L}_B(\theta) + \frac{\lambda}{2} \sum_i F_i (\theta_i - \theta_{A,i}^*)^2$$



Elastic Weight Consolidation (EWC) II



- θ_A^* : optimal weights after Task A
- High $F_i \Rightarrow$ important for Task A $\Rightarrow$ highly constrained
- **Approximation:** diagonal of the Fisher matrix (*per-weight importance*)
- Effective for a *small number* of sequential tasks; **degrades** as tasks accumulate



Experience Replay

Core Idea

Maintain an **episodic memory buffer** $\mathcal{M}$ of *past examples*. When training on new data, **interleave** replayed examples:

$$\mathcal{L} = \mathcal{L}_{\text{new}}(\theta) + \alpha \mathcal{L}_{\text{replay}}(\theta; \mathcal{M})$$

- **Reservoir sampling**: maintains a *uniform random sample* of past data as new data arrives
- DER (*Dark Experience Replay*): replay **soft logits** from the past checkpoint (**knowledge distillation**); more informative than *hard labels*
- **Generative replay**: use a GAN or VAE to *synthesise* past experiences — no raw data storage needed
- **Memory buffer size** is the key *computational trade-off*



Self-Supervised and Contrastive Learning

Pretext Tasks

Create **free supervision** from the data itself: predict *rotation angle*, reconstruct *masked patches* (MAE), solve *jigsaw puzzles*.

SimCLR (Chen *et al.*, 2020)

Apply two *random augmentations* to the same image; **pull** the two views **together**, **push** views of different images **apart**:

$$\mathcal{L}_{\text{NT-Xent}} = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k \neq i} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}$$

• Self-supervised pre-training + linear probe: competitive with supervised pre-training

• Enables transfer learning without any labels in pretraining

• Connection: CLIP uses contrastive learning across image-text



Self-Supervised and Contrastive Learning

Pretext Tasks

Create **free supervision** from the data itself: predict *rotation angle*, reconstruct *masked patches* (MAE), solve *jigsaw puzzles*.

SimCLR (Chen *et al.*, 2020)

Apply two *random augmentations* to the same image; **pull** the two views **together**, **push** views of different images **apart**:

$$\mathcal{L}_{\text{NT-Xent}} = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k \neq i} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}$$

- *Self-supervised* pre-training + **linear probe**: competitive with **supervised** pre-training
- Enables **transfer learning without any labels** in pretraining

Connection: CLIP uses **contrastive learning** across *image-text pairs*



Practical Example: EWC on Sequential MNIST variants

- **Task A:** classify digits 0–4 on standard MNIST
- **Task B:** classify digits 5–9 on **permuted** MNIST (*pixels shuffled*)
- **Baseline:** *sequential fine-tuning* — observe **sharp accuracy drop** on *Task A*
- EWC: add **Fisher-weighted regularisation**; measure *retention*
- **Replay:** compare EWC with a 500-sample *memory buffer*
- **Expected:** **Replay** $\geq$ EWC $\gg$ naive *sequential fine-tuning*
- **Metric:** **backward transfer** $= \frac{1}{T-1} \sum_{i < T} (A_{T,i} - A_{i,i})$ (*average accuracy drop on prior tasks after training on all tasks*)



Practical Example: EWC on Sequential MNIST variants

- **Task A:** classify digits 0–4 on standard MNIST
- **Task B:** classify digits 5–9 on **permuted** MNIST (*pixels shuffled*)
- **Baseline:** *sequential fine-tuning* — observe **sharp accuracy drop** on *Task A*
- EWC: add **Fisher-weighted regularisation**; measure *retention*
- **Replay:** compare EWC with a 500-sample *memory buffer*
- **Expected:** **Replay** $\geq$ EWC $\gg$ naive *sequential fine-tuning*
- **Metric:** **backward transfer** = $\frac{1}{T-1} \sum_{i < T} (A_{T,i} - A_{i,i})$ (*average accuracy drop on prior tasks after training on all tasks*)



Practical Example: EWC on Sequential MNIST variants

- **Task A:** classify digits 0–4 on standard MNIST
- **Task B:** classify digits 5–9 on **permuted** MNIST (*pixels shuffled*)
- **Baseline:** *sequential fine-tuning* — observe **sharp accuracy drop** on *Task A*
- EWC: add **Fisher-weighted regularisation**; measure *retention*
- **Replay:** compare EWC with a 500-sample *memory buffer*
- **Expected:** **Replay** $\geq$ EWC $\gg$ naive *sequential fine-tuning*
- **Metric:** **backward transfer** $= \frac{1}{T-1} \sum_{i < T} (A_{T,i} - A_{i,i})$ (*average accuracy drop on prior tasks after training on all tasks*)



Outline

- 1 Generative Adversarial Networks
- 2 Transfer Learning I: Pretrained Models & Few-Shot Classification
- 3 Transfer Learning II: Neural Style Transfer & Domain Adaptation
- 4 Online & Continual Learning



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/machine-learning>

