

NEURAL NETWORKS

Foundations: Neural Networks to Representation Learning

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2026-I



UNIVERSIDAD DISTRITAL
FRANCISCO JOSÉ DE CALDAS

Outline

1 Neural Network Fundamentals

2 Optimization & Regularization



Outline

1 Neural Network Fundamentals

2 Optimization & Regularization




What is a Neural Network?

Definition

A neural network is a parametric function $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$ composed of simple, interconnected computational units (neurons).

↳ ~ 1948

- **Neuron:** weighted sum + nonlinear activation
- **Layer:** neurons operating in parallel
- **Depth:** key enabler of hierarchical features
- **Training:** adjust θ so $f_{\theta}(x) \approx y$



Different $\theta \Rightarrow$ different f_{θ} .

Training finds the best fit.



What is a Neural Network?

⊖ ~~Not~~ convex


Definition

A neural network is a **parametric function** $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$ composed of simple, interconnected computational units (*neurons*).

- **Neuron:** weighted sum + nonlinear activation
- **Layer:** neurons operating in parallel

- **Depth:** layers of hierarchical features
- **Training:** adjust θ so $f_{\theta}(x) \approx y$

⊖ ~~Not~~ convex
 $w_1 + w_2 = 1$



Different $\theta \Rightarrow$ different f_{θ} .

Training finds the best fit.

$$f(x_1, x_2) = w_1 \cdot x_1 + w_2 \cdot x_2$$



What is a Neural Network?


Definition

A neural network is a **parametric function** $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$ composed of simple, interconnected computational units (*neurons*).

- **Neuron:** weighted sum + nonlinear activation
- **Layer:** neurons operating in parallel

- **Depth:** key enabler of hierarchical features
- **Training:** adjust θ so $f_{\theta}(x) \approx y$

Parallel



Different $\theta \Rightarrow$ different f_{θ} .

Training finds the best fit.




What is a Neural Network?

Definition


A neural network is a **parametric function** $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$ composed of simple, interconnected computational units (*neurons*).

- **Neuron:** weighted sum + nonlinear activation
- **Layer:** neurons operating in parallel
- **Depth:** key enabler of hierarchical features
- **Training:** adjust θ so $f_{\theta}(x) \approx y$



Different $\theta \Rightarrow$ different f_{θ} .

Training finds the best fit.




What is a Neural Network?

Definition

A neural network is a **parametric function** $f_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$ composed of simple, interconnected computational units (*neurons*).

- **Neuron:** **weighted sum** + nonlinear activation
- **Layer:** neurons operating in parallel
- **Depth:** key enabler of **hierarchical features**
- **Training:** adjust θ so $f_{\theta}(x) \approx y$



Different $\theta \Rightarrow$ different f_{θ} .

Training finds the best fit.


$w_1, w_2 \rightarrow 4$
 $\theta_3 \rightarrow b_i$
 Points



The Perceptron

Rosenblatt (1958)

$$\hat{y} = \text{sign}(\sum_i w_i x_i + b)$$



Strengths

- Binary classification
- Convergence guaranteed (linearly separable data)

Limitations

- Cannot solve XOR
- Only linearly separable
- Hard threshold — not differentiable

Solution


Stack layers with smooth activations $\Rightarrow$ Multi-Layer Perceptron



The Perceptron

Rosenblatt (1958)

$$\hat{y} = \text{sign}(\sum_i w_i x_i + b)$$



AND-OR

Solution

Stack layers with smooth activations $\Rightarrow$ Multi-Layer Perceptron

Strengths

- Binary classification ✓
- Convergence guarantee (linearly separable data) ✓

Limitations

- Cannot solve XOR
- Only linearly separable
- Hard threshold — not differentiable


Linear



The Perceptron

Rosenblatt (1958)

$$\hat{y} = \text{sign}(\sum_i w_i x_i + b)$$



Strengths

- Binary classification
- Convergence guarantee (linearly separable data)

Limitations

- Cannot solve XOR
- Only linearly separable
- Hard threshold — not differentiable

Solution

Stack layers with smooth activations $\Rightarrow$ Multi-Layer Perceptron ✓



Multi-Layer Perceptron (MLP) — Architecture


Architecture

Input layer → Hidden layer(s) → Output layer

For each layer l :

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}, \quad \mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)})$$

at least
3 layers



• Universal Approximation (Hornik, 1991): one hidden layer with enough neurons approximates any continuous function
 • Depth is more parameter-efficient than extreme width




Multi-Layer Perceptron (MLP) — Architecture

Architecture

Input layer → Hidden layer(s) → Output layer

For each layer l :

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}, \quad \mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)})$$




Multi-Layer Perceptron (MLP) — Architecture

Architecture

Input layer → Hidden layer(s) → Output layer

For each layer l :

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)}, \quad \mathbf{a}^{(l)} = \sigma(\mathbf{z}^{(l)})$$



Universal Approximation (Hornik, 1991): one hidden layer with enough neurons approximates any continuous function

Depth is more parameter-efficient than extreme width



Multi-Layer Perceptron (MLP) — Output & Softmax




Output layer choices

- **Linear:** regression
- **Sigmoid:** binary classification
- **Softmax:** multi-class → probability distribution over K classes

Softmax

Fuzzy logic


$$\hat{y}_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}}, \quad \hat{y}_k \in (0, 1), \quad \sum_k \hat{y}_k = 1$$



$z = [2.1, 0.3, -1.2] \xrightarrow{\text{softmax}} y = [0.65, 0.27, 0.08] \text{ — prediction: Cat}$



- **Linear**: regression **Sigmoid**: binary classification
- **Softmax**: multi-class $\rightarrow$ probability distribution over K classes

$$\hat{y}_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}}, \quad \hat{y}_k \in (0, 1), \quad \sum_k \hat{y}_k = 1$$



Multi-Layer Perceptron (MLP) — Output & Softmax

Output layer choices

- **Linear**: regression **Sigmoid**: binary classification
- **Softmax**: multi-class → probability distribution over K classes

Softmax

$$\hat{y}_k = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}}, \quad \hat{y}_k \in (0, 1), \quad \sum_k \hat{y}_k = 1$$



Fuzzy logic

$\mathbf{z} = [2.1, 0.3, -1.2] \xrightarrow{\text{softmax}} \hat{\mathbf{y}} = [0.65, 0.27, 0.08] \text{ — prediction: Cat}$



Activation Functions

Sigmoid:

$$\sigma(x) = \frac{1}{1 + e^{-x}} \in (0, 1)$$

Tanh:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \in (-1, 1)$$

ReLU: $\max(0, x)$ — **default today**

GELU: $x \Phi(x)$ — used in Transformers

factor → neuron output

Why ReLU dominates

- No saturation for positive inputs
- Sparse activation (many zeros)
- Computationally trivial
- Mitigates vanishing gradient

$(w_1 x_1 + w_2 x_2) \cdot \text{af}$



Activation Functions

Sigmoid:

$$\sigma(x) = \frac{1}{1 + e^{-x}} \in (0, 1)$$

Tanh:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \in (-1, 1)$$

ReLU: $\max(0, x)$ — **default today**

GELU: $x \cdot \Phi(x)$ — used in Transformers

Why ReLU dominates


- No saturation for positive inputs
- Sparse activation (many zeros)
- Computationally trivial
- Mitigates vanishing gradient



Activation Functions — Visual Overview

Sigmoid


$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Output layer


ReLU

$$\max(0, x)$$




Tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



GELU


$$x \cdot \Phi(x)$$



Loss Functions

The Objective of Training

Minimise $\mathcal{L}(\theta)$ — a measure of disagreement between predictions and targets over the training set.




Loss Functions

Mean Squared Error (regression)

$$\mathcal{L}_{\text{MSE}} = \frac{1}{N} \sum_{i=1}^N \|\hat{y}_i - y_i\|^2$$

Categorical Cross Entropy (classification)

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log \hat{y}_{ik}$$



- Cross-entropy with softmax is the standard loss for multiclass classification
- The choice of loss determines the metric that the network receives




Loss Functions

Mean Squared Error (regression)

$$\mathcal{L}_{\text{MSE}} = \frac{1}{N} \sum_{i=1}^N \|\hat{y}_i - y_i\|^2$$

Categorical Cross-Entropy (classification)

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log \hat{y}_{ik}$$



Loss Functions

Mean Squared Error (regression)

$$\mathcal{L}_{\text{MSE}} = \frac{1}{N} \sum_{i=1}^N \|\hat{y}_i - y_i\|^2$$

Categorical Cross-Entropy (classification)

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log \hat{y}_{ik}$$

- Cross-entropy + softmax is the standard for multiclass classification
- The choice of loss determines the gradient signal the network receives



Backpropagation

Core Idea

Apply the chain rule repeatedly to compute $\frac{\partial \mathcal{L}}{\partial \mathbf{W}^{(l)}}$ for every layer, from output back to input.


Two-pass algorithm:

• Forward pass: compute all $\mathbf{z}^{(l)}$ and store activations

• Backward pass: compute all $\frac{\partial \mathcal{L}}{\partial \mathbf{z}^{(l)}}$ when

Time complexity is roughly the same as forward pass

Made



Backpropagation

Core Idea

Apply the **chain rule** repeatedly to compute $\partial \mathcal{L} / \partial \mathbf{W}^{(l)}$ for every layer, from output back to input.

Two-pass algorithm:

① **Forward pass:** compute $\mathbf{z}^{(l)}$, $\mathbf{a}^{(l)}$ for all l ; store activations

② **Backward pass:** compute error signal $\delta^{(l)}$, then

$$\nabla_{\mathbf{W}^{(l)}} \mathcal{L} = \delta^{(l)} (\mathbf{a}^{(l-1)})^\top$$

• Time complexity: $O(|\theta|)$ — same order as a forward pass

• Made efficient by automatic differentiation (PyTorch, TensorFlow)



Backpropagation

Core Idea

Apply the **chain rule** repeatedly to compute $\partial \mathcal{L} / \partial \mathbf{W}^{(l)}$ for every layer, from output back to input.

Two-pass algorithm:

① **Forward pass:** compute $\mathbf{z}^{(l)}, \mathbf{a}^{(l)}$ for all l ; store activations

② **Backward pass:** compute error signal $\delta^{(l)}$, then

$$\nabla_{\mathbf{W}^{(l)}} \mathcal{L} = \delta^{(l)} (\mathbf{a}^{(l-1)})^\top$$

- Time complexity: $O(|\theta|)$ — same order as a forward pass
- Made efficient by automatic differentiation (PyTorch, TensorFlow)



Backpropagation

Core Idea

Apply the **chain rule** repeatedly to compute $\partial \mathcal{L} / \partial \mathbf{W}^{(l)}$ for every layer, from output back to input.

Two-pass algorithm:

- 1 **Forward pass:** compute $\mathbf{z}^{(l)}, \mathbf{a}^{(l)}$ for all l ; store activations
- 2 **Backward pass:** compute error signal $\delta^{(l)}$, then

$$\nabla_{\mathbf{W}^{(l)}} \mathcal{L} = \delta^{(l)} (\mathbf{a}^{(l-1)})^\top$$
 - Time complexity: $O(|\theta|)$ — same order as a forward pass
 - Made efficient by automatic differentiation (PyTorch, TensorFlow)



The Training Loop

- 1 **Initialise** weights θ_0 (random)
- 2 Split data: train / validation / test
- 3 For each **epoch**:
 - Sample a mini-batch $\{(x_i, y_i)\}_{i=1}^B$
 - Forward pass $\rightarrow$ compute $\mathcal{L}$
 - Backward pass $\rightarrow$ compute $\nabla \mathcal{L}$
 - Update $\theta \leftarrow \theta - \eta \nabla \mathcal{L}$
- 4 Monitor **validation loss**; stop when it increases (early stopping)

Golden Rule

Never tune hyperparameters using the **test set**. The test set measures final generalisation only.



The Training Loop

- 1 **Initialise** weights θ_0 (random)
- 2 Split data: **train** / validation / test
- 3 For each **epoch**:
 - Sample a mini-batch $\{(x_i, y_i)\}_{i=1}^B$
 - Forward pass $\rightarrow$ compute $\mathcal{L}$
 - Backward pass $\rightarrow$ compute $\nabla \mathcal{L}$
 - Update: $\theta \leftarrow \theta - \alpha \nabla \mathcal{L}$
- 4 Monitor **validation loss**; stop when it increases (early stopping)

Golden Rule

Never tune hyperparameters using the **test set**. The test set measures final generalisation only.



The Training Loop

- 1 **Initialise** weights θ_0 (random)
- 2 Split data: **train** / **validation** / **test**
- 3 For each **epoch**:


- Sample a mini-batch $\{(x_i, y_i)\}_{i=1}^B$
- Forward pass $\rightarrow$ compute $\mathcal{L}$
- Backward pass $\rightarrow$ compute $\nabla \mathcal{L}$
- Update: $\theta \leftarrow \theta - \alpha \nabla \mathcal{L}$

over all train data

- 4 Monitor **validation loss**: stop when it increases (early stopping)

Golden Rule

Never tune hyperparameters using the **test set**. The test set measures final generalisation only.



The Training Loop

- 1 **Initialise** weights θ_0 (random)
- 2 Split data: **train** / **validation** / **test**
- 3 For each **epoch**:
 - Sample a mini-batch $\{(x_i, y_i)\}_{i=1}^B$
 - Forward pass $\rightarrow$ compute $\mathcal{L}$
 - Backward pass $\rightarrow$ compute $\nabla \mathcal{L}$
 - Update: $\theta \leftarrow \theta - \alpha \nabla \mathcal{L}$
- 4 Monitor **validation loss**; stop when it increases (early stopping)

Golden Rule

Never tune hyperparameters using the **test set**. The test set measures final generalisation only.



The Training Loop

- ① **Initialise** weights θ_0 (random)
- ② Split data: **train** / **validation** / **test**
- ③ For each **epoch**:
 - Sample a mini-batch $\{(x_i, y_i)\}_{i=1}^B$
 - Forward pass $\rightarrow$ compute $\mathcal{L}$
 - Backward pass $\rightarrow$ compute $\nabla \mathcal{L}$
 - Update: $\theta \leftarrow \theta - \alpha \nabla \mathcal{L}$
- ④ Monitor **validation loss**; stop when it increases (early stopping)

Golden Rule

Never tune hyperparameters using the **test set**. The test set measures final generalisation only.

overfitting



Practical Example: MLP on Two-Moons

- **Problem:** binary classification on non-linear boundary
- **Architecture:** $2 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU, sigmoid output
- **Experiment:** vary number of hidden layers and width
- **Observation:** deep narrow network generalises better than wide shallow one
- **Workshop connection:** understanding this experiment is the foundation for all classification tasks in Challenges 2, 5, and 7



Practical Example: MLP on Two-Moons

- **Problem:** binary classification on non-linear boundary
- **Architecture:** $2 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU, sigmoid output
- **Experiment:** vary number of hidden layers and width
- **Observation:** deep narrow network generalises better than wide shallow one
- **Workshop connection:** understanding this experiment is the foundation for all classification tasks in Challenges 2, 5, and 7



Practical Example: MLP on Two-Moons

- **Problem:** binary classification on non-linear boundary
- **Architecture:** $2 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU, sigmoid output
- **Experiment:** vary number of hidden layers and width
- **Observation:** deep narrow network generalises better than wide shallow one
- **Workshop connection:** understanding this experiment is the foundation for all classification tasks in Challenges 2, 5, and 7



Practical Example: MLP on Two-Moons

- **Problem:** binary classification on non-linear boundary
- **Architecture:** $2 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU, sigmoid output
- **Experiment:** vary number of hidden layers and width
- **Observation:** deep narrow network generalises better than wide shallow one
- **Workshop connection:** understanding this experiment is the foundation for all classification tasks in Challenges 2, 5, and 7



Practical Example: MLP on Two-Moons

- **Problem:** binary classification on non-linear boundary
- **Architecture:** $2 \rightarrow 32 \rightarrow 32 \rightarrow 1$, ReLU, sigmoid output
- **Experiment:** vary number of hidden layers and width
- **Observation:** deep narrow network generalises better than wide shallow one
- **Workshop connection:** understanding this experiment is the foundation for all classification tasks in Challenges 2, 5, and 7



Outline

- 1 Neural Network Fundamentals
- 2 Optimization & Regularization



Gradient Descent

Update Rule

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$$

Full-batch GD

$$\theta_0 = w_1, w_2, w_3, \dots, w_n$$

Deterministic

$$\theta_1 = w_1, w_2, w_3, \dots, w_n$$

Mini-batch SGD

- Noisy gradient estimates $\rightarrow$ random
- Constant step size per step
- Noise helps escape local minima $\rightarrow \nabla_{\theta}$
- Standard in practice

- Batch size B : typical values 32–512; affects convergence and memory
- Learning rate α is the most critical hyperparameter

$$\nabla_{\theta} \approx \theta$$



Gradient Descent

Update Rule

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$$

Full-batch GD

- Exact gradient
- Expensive per step
- Deterministic

Mini-batch SGD

- Noisy gradient estimate
- Cheap per step
- Noise helps escape local minima
- Standard in practice



Gradient Descent

Update Rule

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$$

hyperparameter optimization

Full-batch GD

- Exact gradient
- Expensive per step
- Deterministic

Mini-batch SGD

- Noisy gradient estimate
- Cheap per step
- Noise helps escape local minima

$\bar{\nabla} \neq 0$

32-64-128-256-512

• Standard in practice

- Batch size B : typical values 32-512; affects convergence and memory
- Learning rate α is the most critical hyperparameter

0.0003
0.0001

0.003
0.001

0.03
0.01


0.3
0.1



Momentum and Nesterov *~ 80's*

SGD with Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} - \alpha \nabla \mathcal{L}, \quad \theta \leftarrow \theta + \mathbf{v}$$



Intuition

Momentum is like a ball rolling down a hill that builds *inertia*. Nesterov is a ball that *looks ahead* before deciding where to roll.



Momentum and Nesterov

SGD with Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} - \alpha \nabla \mathcal{L}, \quad \theta \leftarrow \theta + \mathbf{v}$$

- Typical $\beta = 0.9$; accumulates gradient **history**
- Dampens oscillations; accelerates in consistent directions
- **Nesterov**: compute gradient at the **look-ahead** position $\theta + \beta \mathbf{v}$ instead of current θ — slightly better convergence

Intuition

Momentum is like a ball rolling down a hill that builds **inertia**. Nesterov is a ball that **looks ahead** before deciding where to roll.



Momentum and Nesterov

SGD with Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} - \alpha \nabla \mathcal{L}, \quad \theta \leftarrow \theta + \mathbf{v}$$

- Typical $\beta = 0.9$; accumulates gradient **history**
- Dampens oscillations; accelerates in consistent directions
- **Nesterov**: compute gradient at the **look ahead** position $\theta + \beta \mathbf{v}$ instead of current θ — slightly better convergence

Intuition

Momentum is like a ball rolling down a hill that builds **inertia**. Nesterov is a ball that **looks ahead** before deciding where to roll.



Momentum and Nesterov

SGD with Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} - \alpha \nabla \mathcal{L}, \quad \theta \leftarrow \theta + \mathbf{v}$$

- Typical $\beta = 0.9$; accumulates gradient **history**
- Dampens oscillations; accelerates in consistent directions
- **Nesterov**: compute gradient at the **look-ahead** position $\theta + \beta \mathbf{v}$ instead of current θ — slightly better convergence

Intuition

Momentum is like a ball rolling down a hill that builds **inertia**. Nesterov is a ball that **looks ahead** before deciding where to roll.



Momentum and Nesterov

SGD with Momentum

$$\mathbf{v} \leftarrow \beta \mathbf{v} - \alpha \nabla \mathcal{L}, \quad \theta \leftarrow \theta + \mathbf{v}$$

- Typical $\beta = 0.9$; accumulates gradient **history**
- Dampens oscillations; accelerates in consistent directions
- **Nesterov**: compute gradient at the **look-ahead** position $\theta + \beta \mathbf{v}$ instead of current θ — slightly better convergence

Intuition

Momentum is like a ball rolling down a hill that builds **inertia**. Nesterov is a ball that looks ahead before deciding where to roll.

↓
snowball



Adam Optimizer $\rightarrow$ XGBoost

Adaptive Moment Estimation (Kingma & Ba, 2015)

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \underline{g}_t$$

$\swarrow$ *weighted avg.*

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\theta \leftarrow \theta - \underline{\alpha} \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$\swarrow$ *scaling*

$\swarrow$ *momentum*

$\swarrow$ *adaptive*

Convex Combination
 $\alpha + (1 - \alpha)$



Adam Optimizer

Adaptive Moment Estimation (Kingma & Ba, 2015)

$$m_t = \underline{\beta_1} m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \underline{\beta_2} v_{t-1} + (1 - \beta_2) g_t^2$$

$$\theta \leftarrow \theta - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

- Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$, $\alpha = 10^{-3}$
0.001
- Per-parameter adaptive learning rates
- **AdamW**: Adam + decoupled weight decay — preferred in current practice
- Adam dominates in deep learning; SGD+Momentum preferred in CV research



Adam Optimizer

Adaptive Moment Estimation (Kingma & Ba, 2015)

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$$\theta \leftarrow \theta - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)$$

- Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$, $\alpha = 10^{-3}$
- Per-parameter **adaptive learning rates**
- **AdamW**: Adam + decoupled weight decay — preferred in current practice
- Adam dominates in deep learning; SGD+Momentum preferred in CV research



Adam Optimizer

Adaptive Moment Estimation (Kingma & Ba, 2015)

$$\begin{aligned}
 m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t & \hat{m}_t &= m_t / (1 - \beta_1^t) \\
 v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 & \hat{v}_t &= v_t / (1 - \beta_2^t) \\
 \theta &\leftarrow \theta - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)
 \end{aligned}$$

- Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$, $\alpha = 10^{-3}$
- Per-parameter **adaptive learning rates**
- **AdamW**: Adam + decoupled weight decay — preferred in current practice
- Adam dominates in deep learning; SGD+Momentum preferred in CV research



Adam Optimizer

Adaptive Moment Estimation (Kingma & Ba, 2015)

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$$\theta \leftarrow \theta - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)$$

- Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$, $\alpha = 10^{-3}$
- Per-parameter **adaptive learning rates**
- **AdamW**: Adam + decoupled weight decay — preferred in current practice
- Adam dominates in deep learning; SGD + Momentum preferred in CV research



- ReduceLROnPlateau: halve LR when validation loss stops improving

Modern practice: **linear warmup** + **cosine decay**. Used in LLaMA, BERT, and GPT. Especially important with large batch sizes.

Converge
to fast

Optimal



Learning Rate Scheduling

- **Step decay:** multiply LR by factor γ every k epochs
- **Cosine annealing** $\alpha_t = \alpha_{\min} + \frac{1}{2}(\alpha_{\max} - \alpha_{\min})(1 + \cos(\pi t/T))$
- **Linear warmup:** ramp from near-zero for first w steps, then decay
- **ReduceLROnPlateau:** halve LR when validation loss stops improving

Warmup and Decay

Modern practice: linear warmup + cosine decay. Used in ResNets, BERT, and GPT. Especially important with large batch sizes.



Learning Rate Scheduling

- **Step decay:** multiply LR by factor γ every k epochs
- **Cosine annealing:** $\alpha_t = \alpha_{\min} + \frac{1}{2}(\alpha_{\max} - \alpha_{\min})(1 + \cos(\pi t/T))$
- **Linear warmup:** ramp from near-zero for first w steps, then decay
- **ReduceLROnPlateau:** halve LR when validation loss stops improving

Warmup and Decay

Modern practice: **linear warmup** + **cosine decay**. Used in ResNets, BERT, and GPT. Especially important with large batch sizes.



Learning Rate Scheduling

- **Step decay:** multiply LR by factor γ every k epochs
- **Cosine annealing:** $\alpha_t = \alpha_{\min} + \frac{1}{2}(\alpha_{\max} - \alpha_{\min})(1 + \cos(\pi t/T))$
- **Linear warmup:** ramp from near-zero for first w steps, then decay
- **ReduceLROnPlateau:** halve LR when validation loss stops improving

Warmup and Decay

Modern practice: **linear warmup** + **cosine decay**. Used in ResNets, BERT, and GPT. Especially important with large batch sizes.




Learning Rate Scheduling

- **Step decay:** multiply LR by factor γ every k epochs
- **Cosine annealing:** $\alpha_t = \alpha_{\min} + \frac{1}{2}(\alpha_{\max} - \alpha_{\min})(1 + \cos(\pi t/T))$
- **Linear warmup:** ramp from near-zero for first w steps, then decay
- **ReduceLROnPlateau:** halve LR when validation loss stops improving

Warmup and Decay

Modern practice: linear warmup + cosine decay. Used in ResNets, BERT, and GPT. Especially important with large batch sizes.



Google



Bias-Variance Trade-Off

Underfitting (high bias)


- Model too simple
- High training error
- Cannot capture data structure
- Fix: increase capacity

Overfitting (high variance)

- Model too complex
- Low training, high val error
- Memorises noise
- Fix: regularise

Practical Guideline

Start with a model that slightly overfits, then regularise until train and val errors converge. Never start with a model that underfits.



Bias-Variance Trade-Off

Underfitting (high bias)

- Model too simple
- High training error
- Cannot capture data structure
- Fix: increase **capacity**

Overfitting (high variance)

- Model too complex
- Low training, high val error
- Memorises noise
- Fix: **regularise**

Practical Guideline

Start with a model that **slightly overfits**, then regularise until train and val errors converge. Never start with a model that underfits.

α -medium $\Leftarrow$ Control



L1 and L2 Regularization

→ training time

Augmented Objective

$$\mathcal{L}_{\text{reg}}(\theta) = \mathcal{L}(\theta) + \Omega(\theta)$$

 $\mathcal{L} \rightarrow \Omega(\theta)$

L2 Regularization (Weight Decay / Ridge)

$$\Omega(\theta) = \frac{\lambda}{2} \|\mathbf{W}\|_F^2 \quad \text{drives weights small but rarely zero}$$

 $\mathcal{L} \rightarrow -\Omega(\theta)$

L1 Regularization (Lasso)

$$\Omega(\theta) = \lambda \|\mathbf{W}\|_1 \quad \text{promotes sparse weights (many exact zeros)}$$



L1 and L2 Regularization

Augmented Objective

$$\mathcal{L}_{\text{reg}}(\theta) = \mathcal{L}(\theta) + \Omega(\theta)$$

L2 Regularization (Weight Decay / Ridge)

$$\Omega(\theta) = \frac{\lambda}{2} \|\mathbf{W}\|_F^2 \quad \text{drives weights small but rarely zero}$$

L1 Regularization (Lasso)

$$\Omega(\theta) = \lambda \|\mathbf{W}\|_1 \quad \text{promotes sparse weights (many exact zeros)}$$

- In deep learning: L2 (weight decay) is the standard choice
- Elastic Net: combines both — used in linear models and sparse NNs



L1 and L2 Regularization

Augmented Objective

$$\mathcal{L}_{\text{reg}}(\theta) = \mathcal{L}(\theta) + \Omega(\theta)$$

L2 Regularization (Weight Decay / Ridge)

$$\Omega(\theta) = \frac{\lambda}{2} \|\mathbf{W}\|_F^2 \quad \text{drives weights small but rarely zero}$$

L1 Regularization (Lasso)

$$\Omega(\theta) = \lambda \|\mathbf{W}\|_1 \quad \text{promotes sparse weights (many exact zeros)}$$

edge to size



L1 and L2 Regularization

Augmented Objective

$$\mathcal{L}_{\text{reg}}(\theta) = \mathcal{L}(\theta) + \Omega(\theta)$$

L2 Regularization (Weight Decay / Ridge)

$$\Omega(\theta) = \frac{\lambda}{2} \|\mathbf{W}\|_F^2 \quad \text{drives weights small but rarely zero}$$

L1 Regularization (Lasso)

$$\Omega(\theta) = \lambda \|\mathbf{W}\|_1 \quad \text{promotes sparse weights (many exact zeros)}$$

- In deep learning: L2 (weight decay) is the standard choice
- Elastic Net: combines both — used in linear models and sparse NNs



Dropout

Definition (Srivastava et al., 2014)

During training, independently zero out each activation with probability p . At inference, multiply activations by $(1 - p)$ (*inverted dropout* scales during training).

- Forces redundant representations — no single neuron can be relied upon
- Equivalent to averaging over 2^n thinned networks
- Typical $p = 0.5$ (hidden layers); $p = 0.1$ after conv layers
- Less effective with $p > 0.5$ (neurons on the same layer)



Dropout

Definition (Srivastava et al., 2014)

During training, independently zero out each activation with probability p . At inference, multiply activations by $(1 - p)$ (*inverted dropout* scales during training).

- Forces **redundant representations** — no single neuron can be relied upon
- Equivalent to averaging over 2^n thinned networks
- Typical $p = 0.5$ (hidden layers); $p = 0.1$ after conv layers
- Less effective with **Batch Normalization** on the same layer



Dropout

Definition (Srivastava et al., 2014)

During training, independently zero out each activation with probability p . At inference, multiply activations by $(1 - p)$ (*inverted dropout* scales during training).

- Forces **redundant representations** — no single neuron can be relied upon
- Equivalent to averaging over 2^n thinned networks
 - Typical $p = 0.5$ (hidden layers); $p = 0.1$ after conv layers
 - Less effective with **Batch Normalization** on the same layer



Dropout

Definition (Srivastava et al., 2014)

During training, independently zero out each activation with probability p . At inference, multiply activations by $(1 - p)$ (*inverted dropout* scales during training).

- Forces **redundant representations** — no single neuron can be relied upon
- Equivalent to averaging over 2^n thinned networks
- Typical $p = 0.5$ (hidden layers); $p = 0.1$ after conv layers
- Less effective with **Batch Normalization** on the same layer



Dropout

Definition (Srivastava et al., 2014)

During training, independently zero out each activation with probability p . At inference, multiply activations by $(1 - p)$ (*inverted dropout* scales during training).

- Forces **redundant representations** — no single neuron can be relied upon
- Equivalent to averaging over 2^n thinned networks
- Typical $p = 0.5$ (hidden layers); $p = 0.1$ after conv layers
- Less effective with **Batch Normalization** on the same layer



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are learnable parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are **learnable** parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are **learnable** parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are **learnable** parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are **learnable** parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Batch Normalization

Definition (Ioffe & Szegedy, 2015)

Normalise activations within a mini-batch, then **scale and shift**:

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \varepsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

- γ, β are **learnable** parameters
- Stabilises training: allows larger learning rates
- Reduces sensitivity to weight initialisation
- At inference: use running mean/variance accumulated during training
- Alternative: Layer Norm (Transformers), Group Norm (small batches)



Practical Example: Regularizing a Deep MLP

- **Problem:** overfit a 5-layer MLP on a small dataset
- **Experiment A:** add L2 weight decay $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}\}$
- **Experiment B:** add Dropout $p \in \{0.2, 0.4, 0.6\}$
- **Experiment C:** insert Batch Normalization before each activation
- **Metric:** validation accuracy and gap to training accuracy
- **Workshop connection:** regularisation setup is required in Challenges 3, 4, 5, and 6



Practical Example: Regularizing a Deep MLP

- **Problem:** overfit a 5-layer MLP on a small dataset
- **Experiment A:** add L2 weight decay $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}\}$
- **Experiment B:** add Dropout $p \in \{0.2, 0.4, 0.6\}$
- **Experiment C:** insert Batch Normalization before each activation
- **Metric:** validation accuracy and gap to training accuracy
- **Workshop connection:** regularisation setup is required in Challenges 3, 4, 5, and 6



Practical Example: Regularizing a Deep MLP

- **Problem:** overfit a 5-layer MLP on a small dataset
- **Experiment A:** add L2 weight decay $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}\}$
- **Experiment B:** add Dropout $p \in \{0.2, 0.4, 0.6\}$
- **Experiment C:** insert Batch Normalization before each activation
- **Metric:** validation accuracy and gap to training accuracy
- **Workshop connection:** regularisation setup is required in Challenges 3, 4, 5, and 6



Practical Example: Regularizing a Deep MLP

- **Problem:** overfit a 5-layer MLP on a small dataset
- **Experiment A:** add L2 weight decay $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}\}$
- **Experiment B:** add Dropout $p \in \{0.2, 0.4, 0.6\}$
- **Experiment C:** insert Batch Normalization before each activation
- **Metric:** validation accuracy and gap to training accuracy
- **Workshop connection:** regularisation setup is required in Challenges 3, 4, 5, and 6



Practical Example: Regularizing a Deep MLP

- **Problem:** overfit a 5-layer MLP on a small dataset
- **Experiment A:** add L2 weight decay $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}\}$
- **Experiment B:** add Dropout $p \in \{0.2, 0.4, 0.6\}$
- **Experiment C:** insert Batch Normalization before each activation
- **Metric:** validation accuracy and gap to training accuracy
- **Workshop connection:** regularisation setup is required in Challenges 3, 4, 5, and 6



Outline

- 1 Neural Network Fundamentals
- 2 Optimization & Regularization



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/machine-learning>

