

REINFORCEMENT LEARNING

From Fundamentals to Advanced Policy Methods

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Full-time Adjunct Professor
Computer Engineering Program
School of Engineering
Universidad Distrital Francisco José de Caldas

2026-I



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



What is Reinforcement Learning?

Reinforcement Learning

- **Reinforcement Learning** is a computational approach to learning whereby an **agent** tries to maximize the notion of cumulative **reward**.
- The agent learns to achieve a **goal** in an uncertain, potentially complex **environment**.

- **Agent:** The learner or decision maker
- **Environment:** Everything the agent interacts with



What is Reinforcement Learning?

Reinforcement Learning

- **Reinforcement Learning** is a computational approach to learning whereby an **agent** tries to maximize the notion of cumulative **reward**.
 - The agent learns to achieve a **goal** in an uncertain, potentially complex **environment**.
-
- **Agent:** The learner or **decision maker**
 - **Environment:** Everything the agent **interacts** with
 - **Actions:** Set of all possible **moves** the agent can make
 - **Rewards:** Feedback from environment to agent



What is Reinforcement Learning?

Reinforcement Learning

- **Reinforcement Learning** is a computational approach to learning whereby an **agent** tries to maximize the notion of cumulative **reward**.
 - The agent learns to achieve a **goal** in an uncertain, potentially complex **environment**.
-
- **Agent:** The learner or **decision maker**
 - **Environment:** Everything the agent **interacts** with
 - **Actions:** Set of all possible **moves** the agent can make
 - **Rewards:** **Feedback** from environment to agent



What is Reinforcement Learning?

Reinforcement Learning

- **Reinforcement Learning** is a computational approach to learning whereby an **agent** tries to maximize the notion of cumulative **reward**.
 - The agent learns to achieve a **goal** in an uncertain, potentially complex **environment**.
-
- **Agent:** The learner or **decision maker**
 - **Environment:** Everything the agent **interacts** with
 - **Actions:** Set of all possible **moves** the agent can make
 - **Rewards:** **Feedback** from environment to agent



What is Reinforcement Learning?

Reinforcement Learning

- **Reinforcement Learning** is a computational approach to learning whereby an **agent** tries to maximize the notion of cumulative **reward**.
 - The agent learns to achieve a **goal** in an uncertain, potentially complex **environment**.
-
- **Agent:** The learner or **decision maker**
 - **Environment:** Everything the agent **interacts** with
 - **Actions:** Set of all possible **moves** the agent can make
 - **Rewards:** **Feedback** from environment to agent



RL vs. Other Learning Paradigms

Supervised Learning:

- Uses **labeled** data
- Direct feedback
- Static datasets

Unsupervised Learning:

- No labels
- Pattern discovery
- Structure finding

Reinforcement Learning:

- Trial-and-error learning
- Delayed rewards
- Sequential decision making
- Exploration vs. exploitation
- Dynamic environments



RL vs. Other Learning Paradigms

Supervised Learning:

- Uses **labeled** data
- Direct feedback
- Static datasets

Unsupervised Learning:

- No labels
- Pattern discovery
- Structure finding

Reinforcement Learning:

- **Trial-and-error** learning
- **Delayed** rewards
- **Sequential** decision making
- **Exploration** vs. exploitation
- **Dynamic** environments



Markov Decision Processes (MDPs)

An MDP is defined by the tuple (S, A, P, R, γ) :

- **S**: Set of **states**
- **A**: Set of **actions**
- **P**: **Transition** probabilities $P(s'|s, a)$
- **R**: **Reward** function $R(s, a, s')$
- γ : **Discount factor** $[0, 1]$

Markov Property

The future is independent of the past given the **present state**.

$$P(s_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots) = P(s_{t+1}|s_t, a_t)$$



Markov Decision Processes (MDPs)

An MDP is defined by the tuple (S, A, P, R, γ) :

- **S**: Set of **states**
- **A**: Set of **actions**
- **P**: **Transition** probabilities $P(s'|s, a)$
- **R**: **Reward** function $R(s, a, s')$
- γ : **Discount factor** $[0, 1]$

Markov Property

The future is independent of the past given the **present state**.

$$P(s_{t+1}|s_t, a_t, s_{t-1}, a_{t-1}, \dots) = P(s_{t+1}|s_t, a_t)$$



Agent-Environment Interaction

The RL Loop:

- ① Agent observes current **state** s_t
- ② Agent selects **action** a_t based on policy $\pi(a|s)$
- ③ Environment transitions to new **state** s_{t+1}
- ④ Environment gives **reward** r_{t+1}
- ⑤ Agent updates its **knowledge**
- ⑥ Repeat...



Value Functions and Policies

Policy

A **policy** $\pi(a|s)$ is a mapping from states to **probability distributions** over actions.

Value Functions

- **State Value:** $V^\pi(s) = E[G_t|s_t = s]$
- **Action Value (Q-function):** $Q^\pi(s, a) = E[G_t|s_t = s, a_t = a]$

Where $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ is the **return**.



Value Functions and Policies

Policy

A **policy** $\pi(a|s)$ is a mapping from states to **probability distributions** over actions.

Value Functions

- **State Value:** $V^\pi(s) = E[G_t|s_t = s]$
- **Action Value (Q-function):** $Q^\pi(s, a) = E[G_t|s_t = s, a_t = a]$

Where $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ is the **return**.



Practical Example: Grid World Navigation

- **Environment:** 4×4 grid with walls and goal
- **States:** Agent's position (x, y)
- **Actions:** Up, Down, Left, Right
- **Rewards:** -1 per step, $+100$ at goal, -100 in pit
- **Goal:** Find shortest path to reward



Practical Example: Grid World Navigation

- **Environment:** 4×4 grid with walls and goal
- **States:** Agent's position (x, y)
- **Actions:** Up, Down, Left, Right
- **Rewards:** -1 per step, $+100$ at goal, -100 in pit
- **Goal:** Find shortest path to reward



Practical Example: Grid World Navigation

- **Environment:** 4×4 grid with walls and goal
- **States:** Agent's position (x, y)
- **Actions:** Up, Down, Left, Right
- **Rewards:** -1 per step, $+100$ at goal, -100 in pit
- **Goal:** Find shortest path to reward



Q-Learning Algorithm I

Definition

Q-Learning is a **model-free** reinforcement learning algorithm that learns the optimal action-value function $Q^*(s, a)$ directly from interactions with the environment.



Q-Learning Algorithm II

Q-Learning Update Rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

- α : Learning rate $[0, 1]$
- γ : Discount factor $[0, 1]$
- **Off-policy**: Learns optimal policy regardless of behavior policy
- **Model-free**: No need to know transition probabilities



Q-Learning Algorithm II

Q-Learning Update Rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

- α : Learning rate $[0, 1]$
- γ : Discount factor $[0, 1]$
- **Off-policy**: Learns optimal policy regardless of behavior policy
- **Model-free**: No need to know transition probabilities



Bellman Equation for Q-Values

Bellman Optimality Equation

$$Q^*(s, a) = E[r + \gamma \max_{a'} Q^*(s', a') | s, a]$$

- Q-Learning approximates this equation through temporal difference learning
- The TD error: $\delta = r + \gamma \max_a Q(s', a) - Q(s, a)$



Bellman Equation for Q-Values

Bellman Optimality Equation

$$Q^*(s, a) = E[r + \gamma \max_{a'} Q^*(s', a') | s, a]$$

- Q-Learning approximates this equation through **temporal difference learning**
- The **TD error**: $\delta = r + \gamma \max_a Q(s', a) - Q(s, a)$
- Converges to optimal Q-function under certain conditions



Bellman Equation for Q-Values

Bellman Optimality Equation

$$Q^*(s, a) = E[r + \gamma \max_{a'} Q^*(s', a') | s, a]$$

- Q-Learning approximates this equation through **temporal difference** learning
- The **TD error**: $\delta = r + \gamma \max_a Q(s', a) - Q(s, a)$
- Converges to optimal Q-function under certain conditions



Bellman Equation for Q-Values

Bellman Optimality Equation

$$Q^*(s, a) = E[r + \gamma \max_{a'} Q^*(s', a') | s, a]$$

- Q-Learning approximates this equation through **temporal difference** learning
- The **TD error**: $\delta = r + \gamma \max_a Q(s', a) - Q(s, a)$
- Converges to optimal Q-function under certain conditions



Exploration vs. Exploitation

The Exploration-Exploitation Dilemma

Should the agent choose the **best known action** (exploit) or try **something new** (explore)?

ϵ -Greedy Strategy

- With probability ϵ : choose random action (explore)
- With probability $1-\epsilon$: choose best known action (exploit)
- Often use **decaying ϵ** over time



Exploration vs. Exploitation

The Exploration-Exploitation Dilemma

Should the agent choose the **best known action** (exploit) or try **something new** (explore)?

ϵ -Greedy Strategy

- With probability ϵ : choose random action (explore)
- With probability $1-\epsilon$: choose best known action (exploit)
- Often use **decaying ϵ** over time



Practical Example: Q-Learning in Grid World

- Initialize $Q(s, a)$ arbitrarily (e.g., all zeros)
- For each episode:
 - Initialize state s
 - Repeat until episode ends:
 - Choose action a using ϵ -greedy policy
 - Take action a , observe reward r and new state s'
 - Update Q-value: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_a Q(s', a) - Q(s, a)]$
 - Update state: $s \leftarrow s'$



Practical Example: Q-Learning in Grid World

- Initialize $Q(s, a)$ arbitrarily (e.g., all zeros)
- For each episode:
 - Initialize state s
 - Repeat until episode ends:
 - Choose action a using ϵ -greedy policy
 - Take action a , observe reward r and new state s'
 - Update Q-value: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_a Q(s', a) - Q(s, a)]$
 - Update state: $s \leftarrow s'$



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



SARSA Algorithm

SARSA Update Rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

- **On-policy:** Updates based on the **actual policy** being followed
- **SARSA:** State-Action-Reward-State-Action sequence
 - More conservative than Q-Learning
 - Better for **safety-critical** applications



SARSA Algorithm

SARSA Update Rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

- **On-policy:** Updates based on the **actual policy** being followed
- **SARSA:** State-Action-Reward-State-Action sequence
- More **conservative** than Q-Learning
- Better for **safety-critical** applications



On-Policy vs. Off-Policy Learning

On-Policy (SARSA):

- Learns about policy **currently following**
- More **conservative**
- Accounts for exploration in updates
- Safer in risky environments

Off-Policy (Q-Learning):

- Learns about **optimal policy**
- More **aggressive**
- Ignores exploration in updates
- Faster convergence



SARSA Algorithm Mechanics

- Initialize $Q(s, a)$ arbitrarily (e.g., all zeros)
- For each episode:
 - Initialize state s
 - Choose action a using ϵ -greedy policy
 - Repeat until episode ends:
 - Take action a , observe reward r and new state s'
 - Choose next action a' using ϵ -greedy policy
 - Update Q-value: $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$
 - Update state and action: $s \leftarrow s'$, $a \leftarrow a'$



Practical Example: Cliff Walking Problem

- **Environment:** Agent walks along a cliff edge
- **Cliff:** Large negative reward (-100)
- **Goal:** Reach the end safely
- **SARSA:** Takes safer path (away from cliff)
- **Q-Learning:** Takes riskier optimal path (near cliff)



Practical Example: Cliff Walking Problem

- **Environment:** Agent walks along a cliff edge
- **Cliff:** Large negative reward (-100)
- **Goal:** Reach the end safely
- **SARSA:** Takes safer path (away from cliff)
- **Q-Learning:** Takes riskier optimal path (near cliff)



Practical Example: Cliff Walking Problem

- **Environment:** Agent walks along a cliff edge
- **Cliff:** Large negative reward (-100)
- **Goal:** Reach the end safely
- **SARSA:** Takes safer path (away from cliff)
- **Q-Learning:** Takes riskier optimal path (near cliff)



Function Approximation in RL

The Problem

- Large or **continuous state spaces**
- Cannot store **Q-values** for every state-action pair
- Need to **generalize** from seen to unseen states

Solution: Function Approximation

Use neural networks to approximate: $Q(s, a) \approx Q(s, a; \theta)$



Function Approximation in RL

The Problem

- Large or **continuous state spaces**
- Cannot store **Q-values** for every state-action pair
- Need to **generalize** from seen to unseen states

Solution: Function Approximation

Use neural networks to approximate: $Q(s, a) \approx Q(s, a; \theta)$



Deep Q Network (DQN) Architecture

Key Components

- **Deep Neural Network:** Approximates Q-function
- **Experience Replay:** Stores and samples past experiences
- **Target Network:** Stable Q-targets for training
- **Fixed Q-Targets:** Reduces correlation in updates

Loss Function

$$L(\theta) = E[(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta))^2]$$



Deep Q Network (DQN) Architecture

Key Components

- **Deep Neural Network:** Approximates Q-function
- **Experience Replay:** Stores and samples past experiences
- **Target Network:** Stable Q-targets for training
- **Fixed Q-Targets:** Reduces correlation in updates

Loss Function

$$L(\theta) = E[(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta))^2]$$



Experience Replay Mechanism

- **Store:** Experiences (s, a, r, s') in replay **buffer**
- **Sample:** Random minibatch for training
- **Benefits:**
 - Breaks temporal correlations
 - Improves data efficiency
 - Enables batch learning



Experience Replay Mechanism

- **Store:** Experiences (s, a, r, s') in replay **buffer**
- **Sample:** Random minibatch for training
- **Benefits:**
 - Breaks temporal correlations
 - Improves data efficiency
 - Enables batch learning



Experience Replay Mechanism

- **Store:** Experiences (s, a, r, s') in replay **buffer**
- **Sample:** Random minibatch for training
- **Benefits:**
 - Breaks **temporal correlations**
 - Improves **data efficiency**
 - Enables **batch learning**



DQN for Atari Breakout: Practical Example

Atari Breakout Implementation

- **Input:** Raw pixels $(210 \times 160 \times 3) \rightarrow$ preprocessed $(84 \times 84 \times 4)$
- **Network:** CNN with 3 conv layers + 2 dense layers
- **Actions:** 4 discrete actions (NOOP, FIRE, RIGHT, LEFT)

Preprocessing: Grayscale, resize, frame stacking (4 frames)

Reward: +1 for each brick destroyed, +5 for clearing level

Penalty: -1 for hitting paddle, -1 for hitting wall, -1 for hitting ball

Exploration: ε-greedy with ε decreasing over time

Training: Using experience replay and target network



DQN for Atari Breakout: Practical Example

Atari Breakout Implementation

- **Input:** Raw pixels ($210 \times 160 \times 3$) → preprocessed ($84 \times 84 \times 4$)
- **Network:** CNN with 3 conv layers + 2 dense layers
- **Actions:** 4 discrete actions (NOOP, FIRE, RIGHT, LEFT)
- **Preprocessing:** Grayscale, resize, frame stacking (4 frames)
- **Reward:** +1 for each brick destroyed, +5 for clearing level
- **Strategy:** Agent learns to tunnel through bricks for maximum score
- **Performance:** Achieves superhuman performance (400 points)



DQN for Atari Breakout: Practical Example

Atari Breakout Implementation

- **Input:** Raw pixels ($210 \times 160 \times 3$) → preprocessed ($84 \times 84 \times 4$)
 - **Network:** CNN with 3 conv layers + 2 dense layers
 - **Actions:** 4 discrete actions (NOOP, FIRE, RIGHT, LEFT)
-
- **Preprocessing:** Grayscale, resize, frame stacking (4 frames)
 - **Reward:** +1 for each brick destroyed, +5 for clearing level
 - **Strategy:** Agent learns to tunnel through bricks for maximum score
 - **Performance:** Achieves superhuman performance (400 points)



DQN for Atari Breakout: Practical Example

Atari Breakout Implementation

- **Input:** Raw pixels ($210 \times 160 \times 3$) → preprocessed ($84 \times 84 \times 4$)
- **Network:** CNN with 3 conv layers + 2 dense layers
- **Actions:** 4 discrete actions (NOOP, FIRE, RIGHT, LEFT)
- **Preprocessing:** Grayscale, resize, frame stacking (4 frames)
- **Reward:** +1 for each brick destroyed, +5 for clearing level
- **Strategy:** Agent learns to **tunnel through bricks** for maximum score
- **Performance:** Achieves **superhuman performance** (400 points)



DQN for Atari Breakout: Practical Example

Atari Breakout Implementation

- **Input:** Raw pixels ($210 \times 160 \times 3$) → preprocessed ($84 \times 84 \times 4$)
 - **Network:** CNN with 3 conv layers + 2 dense layers
 - **Actions:** 4 discrete actions (NOOP, FIRE, RIGHT, LEFT)
-
- **Preprocessing:** Grayscale, resize, frame stacking (4 frames)
 - **Reward:** +1 for each brick destroyed, +5 for clearing level
 - **Strategy:** Agent learns to **tunnel through bricks** for maximum score
 - **Performance:** Achieves **superhuman** performance (400 points)



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



Policy-Based vs. Value-Based Methods

Value-Based:

- Learn **value functions**
- Derive policy from values
- Discrete action spaces
- Q-Learning, SARSA

Policy-Based:

- Learn **policy directly**
- Parameterized policies
- **Continuous actions**
- Stochastic policies



Policy-Based vs. Value-Based Methods

Value-Based:

- Learn **value functions**
- Derive policy from values
- Discrete action spaces
- Q-Learning, SARSA

Policy-Based:

- Learn **policy directly**
- Parameterized policies
- **Continuous** actions
- Stochastic policies



Policy Gradient Theorem

Gradient of Expected Return

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a)]$$

↑ Policy parameters

- Intuition: Move parameters in direction that increases probability of good actions
- Key insight: We can estimate gradients through sampling



Policy Gradient Theorem

Gradient of Expected Return

$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a)]$$

- \hat{I}_{\cdot} : Policy parameters
- **Intuition:** Move parameters in direction that **increases probability** of good actions
- **Key insight:** We can estimate gradients through **sampling**



REINFORCE Algorithm

REINFORCE Update

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t$$

- Uses Monte Carlo return G_t
- High variance but unbiased
- Often use baseline to reduce variance
- $\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (G_t - b(s_t))$



REINFORCE Algorithm

REINFORCE Update

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t$$

- Uses **Monte Carlo** return G_t
- High **variance** but unbiased
- Often use **baseline** to reduce variance
- $\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (G_t - b(s_t))$



Actor-Critic Architecture

Two Components

- **Actor:** Policy $\pi_\theta(a|s)$ — chooses actions
- **Critic:** Value function $V_\phi(s)$ — evaluates actions

Benefits

- Lower variance than REINFORCE
- Online learning (no need to wait for episode end)
- Uses TD error: $\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$



Actor-Critic Architecture

Two Components

- **Actor:** Policy $\pi_\theta(a|s)$ — chooses actions
- **Critic:** Value function $V_\phi(s)$ — evaluates actions

Benefits

- Lower variance than REINFORCE
- Online learning (no need to wait for episode end)
- Uses TD error: $\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$



Proximal Policy Optimization (PPO) Motivation

Problems with Basic Policy Gradients

- Large policy updates can be **destructive**
- Hard to choose appropriate **step size**
- Poor **sample efficiency**

PPO Solution

Constrain policy updates to stay in a **trust region** around current policy



Proximal Policy Optimization (PPO) Motivation

Problems with Basic Policy Gradients

- Large policy updates can be **destructive**
- Hard to choose appropriate **step size**
- Poor **sample efficiency**

PPO Solution

Constrain policy updates to stay in a **trust region** around current policy



PPO Clipped Objective

Clipped Surrogate Objective

$$L^{CLIP}(\theta) = E[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)]$$

- $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ (probability ratio)
- \hat{A}_t is advantage estimate
- Clips ratio to $[1 - \epsilon, 1 + \epsilon]$
- Prevents **destructively large** policy updates



Practical Example: PPO on LunarLander-v2

LunarLander-v2 Environment

- **State:** Position, velocity, angle, angular velocity, leg contact (8D continuous)
- **Actions:** 4 discrete actions (NOOP, LEFT, MAIN, RIGHT)
- **Goal:** Land safely between flags (+100 to +140 points)

→ Reward Shaping: Distance, speed, angle penalties, fuel costs
→ PPO Advantages: Stable learning with sparse rewards
→ Policy Network: MLP ($R \rightarrow 128 \rightarrow 128 \rightarrow 4$)
→ Value Network: MLP ($R \rightarrow 128 \rightarrow 128 \rightarrow 1$)
→ Actor-Critic: PPO (Actor: Policy Network, Critic: Value Network)



Practical Example: PPO on LunarLander-v2

LunarLander-v2 Environment

- **State:** Position, velocity, angle, angular velocity, leg contact (8D continuous)
 - **Actions:** 4 discrete actions (NOOP, LEFT, MAIN, RIGHT)
 - **Goal:** Land safely between flags (+100 to +140 points)
-
- **Reward Shaping:** Distance, speed, angle penalties; fuel costs
 - **PPO Advantages:** Stable learning with sparse rewards
 - **Policy Network:** MLP ($8 \rightarrow 128 \rightarrow 128 \rightarrow 4$)
 - **Training:** Converges in 1000 – 3000 episodes
 - **Performance:** Solves environment (average reward > 200)



Practical Example: PPO on LunarLander-v2

LunarLander-v2 Environment

- **State:** Position, velocity, angle, angular velocity, leg contact (8D continuous)
 - **Actions:** 4 discrete actions (NOOP, LEFT, MAIN, RIGHT)
 - **Goal:** Land safely between flags (+100 to +140 points)
-
- **Reward Shaping:** Distance, speed, angle penalties; fuel costs
 - **PPO Advantages:** Stable learning with **sparse rewards**
 - **Policy Network:** MLP ($8 \rightarrow 128 \rightarrow 128 \rightarrow 4$)
 - **Training:** Converges in 1000 – 3000 episodes
 - **Performance:** Solves environment (**average reward > 200**)



Practical Example: PPO on LunarLander-v2

LunarLander-v2 Environment

- **State:** Position, velocity, angle, angular velocity, leg contact (8D continuous)
 - **Actions:** 4 discrete actions (NOOP, LEFT, MAIN, RIGHT)
 - **Goal:** Land safely between flags (+100 to +140 points)
-
- **Reward Shaping:** Distance, speed, angle penalties; fuel costs
 - **PPO Advantages:** Stable learning with **sparse rewards**
 - **Policy Network:** MLP ($8 \rightarrow 128 \rightarrow 128 \rightarrow 4$)
 - **Training:** Converges in 1000 – 3000 episodes
 - **Performance:** Solves environment (**average reward > 200**)



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



Trust Region Concept

Trust Region Optimization

- Optimize within a **trusted region** of current policy
- Ensures **monotonic improvement**
- Prevents policy **collapse**

Trust Region Policy Optimization (TRPO) Constraint

$$E[KL(\pi_{\theta_{old}}(\cdot|s), \pi_{\theta}(\cdot|s))] \leq \delta$$



Trust Region Concept

Trust Region Optimization

- Optimize within a **trusted region** of current policy
- Ensures **monotonic improvement**
- Prevents policy **collapse**

Trust Region Policy Optimization (TRPO) Constraint

$$E[KL(\pi_{\theta_{old}}(\cdot|s), \pi_{\theta}(\cdot|s))] \leq \delta$$



TRPO Mathematical Formulation

Constrained Optimization Problem

$$\text{maximize } L(\theta) \quad (1)$$

$$\text{subject to } \bar{D}_{KL}(\theta_{old}, \theta) \leq \delta \quad (2)$$

- Solved using conjugate gradient method
- More principled than PPO but computationally expensive
- Guarantees monotonic improvement



TRPO Mathematical Formulation

Constrained Optimization Problem

$$\text{maximize } L(\theta) \quad (1)$$

$$\text{subject to } \bar{D}_{KL}(\theta_{old}, \theta) \leq \delta \quad (2)$$

- Solved using **conjugate gradient** method
- More principled than PPO but **computationally expensive**
- Guarantees **monotonic improvement**



Natural Policy Gradients

Natural Policy Gradient

- Uses Fisher information matrix to scale gradients
- Accounts for geometry of policy space
- Update: $\theta \leftarrow \theta + \alpha F^{-1} \nabla_{\theta} J(\theta)$

- More stable updates than vanilla policy gradients
- Related to TRPO as it also considers trust regions
- Can be used in large action spaces



Natural Policy Gradients

Natural Policy Gradient

- Uses Fisher information matrix to scale gradients
- Accounts for geometry of policy space
- Update: $\theta \leftarrow \theta + \alpha F^{-1} \nabla_{\theta} J(\theta)$
- More stable updates than vanilla policy gradients
- Related to TRPO as it also considers trust regions
- Can be used in large action spaces



Practical Example: TRPO in Complex Continuous Control

Complex Continuous Control Tasks

- High-dimensional state and action spaces
- Requires precise control and coordination
- Examples: *robotic manipulation, locomotion*

• Modeling: Simulating dynamics and constraints
• Policy Network: Deep neural networks for representation
• Training: Sample-efficient with trust region updates
• Evaluation: Human-level performance



Practical Example: TRPO in Complex Continuous Control

Complex Continuous Control Tasks

- High-dimensional state and action spaces
- Requires precise control and coordination
- Examples: *robotic manipulation, locomotion*

- **TRPO Advantages:** Handles complex dynamics and constraints
- **Policy Network:** Deep neural networks for representation
- **Training:** Sample-efficient with trust region updates
- **Performance:** Achieves state-of-the-art results



Practical Example: TRPO in Complex Continuous Control

Complex Continuous Control Tasks

- High-dimensional state and action spaces
 - Requires precise control and coordination
 - Examples: *robotic manipulation, locomotion*
-
- **TRPO Advantages:** Handles complex dynamics and constraints
 - **Policy Network:** Deep neural networks for representation
 - **Training:** Sample-efficient with trust region updates
 - **Performance:** Achieves state-of-the-art results



Practical Example: TRPO in Complex Continuous Control

Complex Continuous Control Tasks

- High-dimensional state and action spaces
- Requires precise control and coordination
- Examples: *robotic manipulation, locomotion*
- **TRPO Advantages:** Handles complex dynamics and constraints
- **Policy Network:** Deep neural networks for representation
- **Training:** Sample-efficient with trust region updates
- **Performance:** Achieves state-of-the-art results



Multi-Agent Environments

Challenges

- Non-stationary environments
- Other agents are learning simultaneously
- Coordination and competition
- Communication between agents

- Cooperative: Agents share common goal
- Competitive: Zero-sum games
- Adversarial: One agent's gain is another's loss



Multi-Agent Environments

Challenges

- Non-stationary environments
 - Other agents are learning simultaneously
 - Coordination and competition
 - Communication between agents
-
- Cooperative: Agents share common goal
 - Competitive: Zero-sum games
 - Mixed: Combination of cooperation and competition



Multi-Agent Environments

Challenges

- Non-stationary environments
 - Other agents are learning simultaneously
 - Coordination and competition
 - Communication between agents
-
- **Cooperative:** Agents share common goal
 - **Competitive:** Zero-sum games
 - **Mixed:** Combination of cooperation and competition



Multi-Agent Environments

Challenges

- Non-stationary environments
 - Other agents are learning simultaneously
 - Coordination and competition
 - Communication between agents
-
- **Cooperative:** Agents share common goal
 - **Competitive:** Zero-sum games
 - **Mixed:** Combination of cooperation and competition



Nash Equilibrium in RL

Nash Equilibrium

A strategy profile where no agent can improve by **unilaterally changing their strategy**.

- Difficult to compute in practice
- May not always exist or be unique
- Agents may not converge to equilibrium
- Alternative: Find correlated equilibria



Nash Equilibrium in RL

Nash Equilibrium

A strategy profile where no agent can improve by **unilaterally changing** their strategy.

- Difficult to **compute** in practice
- May not always **exist** or be **unique**
- Agents may not **converge** to equilibrium
- Alternative: Find **correlated equilibria**



Nash Equilibrium in RL

Nash Equilibrium

A strategy profile where no agent can improve by **unilaterally changing** their strategy.

- Difficult to **compute** in practice
- May not always **exist** or be **unique**
- Agents may not **converge** to equilibrium
- Alternative: Find **correlated equilibria**



Centralized Training, Decentralized Execution

CTDE Paradigm

- **Training:** Use **global information** and centralized critic
- **Execution:** Each agent acts with only **local observations**

- Addresses central plausibility
- Enables coordination during training
- Maintains decentralized execution
- Examples: MADDPG, COMA, QMIX



Centralized Training, Decentralized Execution

CTDE Paradigm

- **Training:** Use **global information** and centralized critic
- **Execution:** Each agent acts with only **local observations**
- Addresses **partial observability**
- Enables **coordination** during training
- Maintains **decentralized execution**
- Examples: MADDPG, COMA, QMIX



Centralized Training, Decentralized Execution

CTDE Paradigm

- **Training:** Use **global information** and centralized critic
 - **Execution:** Each agent acts with only **local observations**
-
- Addresses **partial observability**
 - Enables **coordination** during training
 - Maintains **decentralized execution**
 - Examples: MADDPG, COMA, QMIX



Communication in Multi-Agent Systems

Communication Protocols

- **Explicit:** Agents exchange messages
- **Implicit:** Agents infer others' intentions from actions

- Communication can enhance coordination and performance.
- Challenges: Bandwidth constraints, noise, security



Communication in Multi-Agent Systems

Communication Protocols

- **Explicit:** Agents exchange messages
- **Implicit:** Agents infer others' intentions from actions
- Communication can enhance **coordination** and **performance**
- Challenges: Bandwidth constraints, noise, security
- Research area: Learning communication protocols end-to-end



Communication in Multi-Agent Systems

Communication Protocols

- **Explicit:** Agents exchange messages
- **Implicit:** Agents infer others' intentions from actions
- Communication can enhance **coordination** and **performance**
- Challenges: Bandwidth constraints, noise, security
- Research area: Learning communication protocols end-to-end



Communication in Multi-Agent Systems

Communication Protocols

- **Explicit:** Agents exchange messages
- **Implicit:** Agents infer others' intentions from actions
- Communication can enhance **coordination** and **performance**
- Challenges: Bandwidth constraints, noise, security
- Research area: Learning **communication protocols** end-to-end



Outline

- 1 Introduction to Reinforcement Learning Fundamentals
- 2 Advanced Reinforcement Learning Algorithms
- 3 Policy Optimization Methods
- 4 Advanced Policy Methods & Multi-Agent RL



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/reinforcement-learning>

