

INTRODUCTION TO MACHINE LEARNING

Systems Sciences Foundations

Author: Eng. Carlos Andrés Sierra, M.Sc.
cavirguezs@udistrital.edu.co

Lecturer
Department of Computer Engineering
School of Engineering
Universidad Distrital Francisco José de Caldas

2025-I



Outline

- 1 Fundamentals of Machine Learning
- 2 Python Tools for Machine Learning
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation



Outline

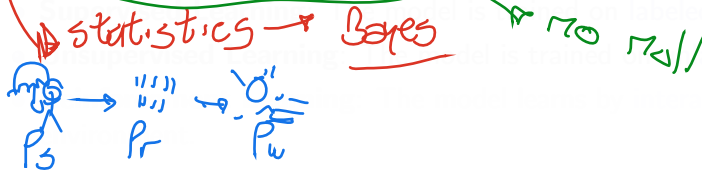
- 1 Fundamentals of Machine Learning
- 2 Python Tools for Machine Learning
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation



Key Concepts in Machine Learning

Machine Learning

- **Machine learning** is a method of ~~data~~ analysis that automates analytical model building.
- It is a **branch** of artificial intelligence based on the idea that systems can learn from data, identify patterns and make decisions with minimal human intervention.



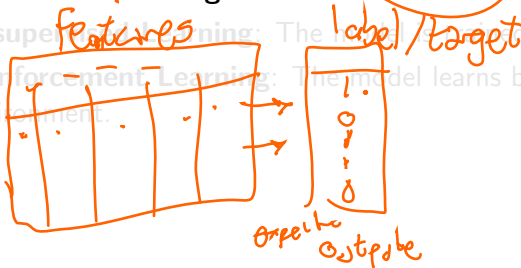
Key Concepts in Machine Learning

Machine Learning

- **Machine learning** is a method of data analysis that **automates** analytical model building.
- It is a **branch** of **artificial intelligence** based on the idea that systems can **learn from data**, **identify patterns** and **make decisions** with **minimal human intervention**.

- ① • **Supervised Learning**: The model is trained on labeled data.


- **Unsupervised Learning**: The model is trained on unlabeled data.
- **Reinforcement Learning**: The model learns by interacting with an environment.



Key Concepts in Machine Learning

Machine Learning

- **Machine learning** is a method of data analysis that **automates** analytical model building.
- It is a **branch** of **artificial intelligence** based on the idea that systems can **learn from data**, **identify patterns** and **make decisions** with minimal human intervention.
- **Supervised Learning**: The model is trained on **labeled data**.
- **Unsupervised Learning**: The model is trained on **unlabeled data**.
- **Reinforcement Learning**: The model learns by interacting with an environment.



Key Concepts in Machine Learning

Machine Learning

- **Machine learning** is a method of data analysis that **automates** analytical model building.
- It is a **branch** of **artificial intelligence** based on the idea that systems can **learn from data**, **identify patterns** and **make decisions** with minimal human intervention.

- **Supervised Learning**: The model is trained on **labeled data**.
- **Unsupervised Learning**: The model is trained on **unlabeled data**.
- **Reinforcement Learning**: The model learns by **interacting** with an environment.

1 step
mult. - step → cybernetics



Typical Machine Learning Problems

- **Classification:** Predicting a label. *targets → supervised*


- **Regression:** Predicting a continuous value.

- **Clustering:** Finding similar data points.

- **Dimensionality Reduction:** Reducing the number of features.

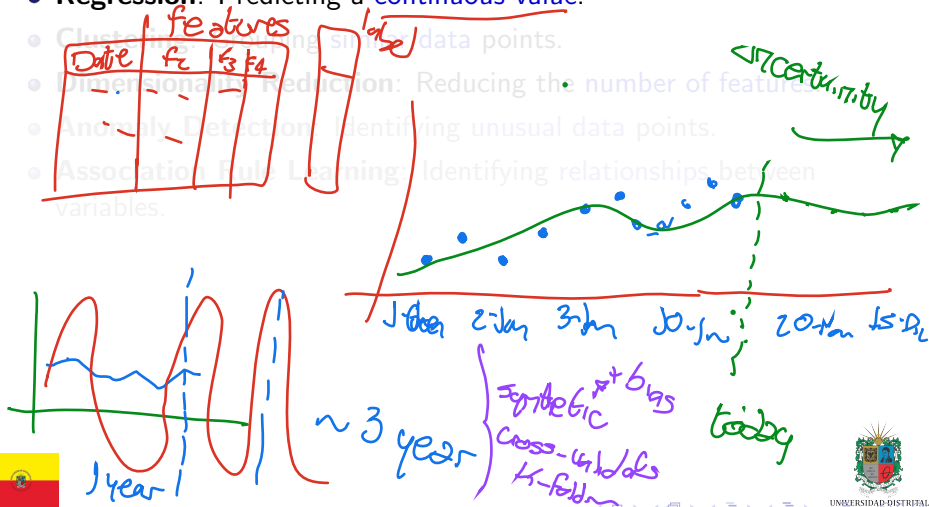
- **Anomaly Detection:** Identifying unusual data points.

- **Association Rule Learning:** Identifying relationships between variables.



Typical Machine Learning Problems

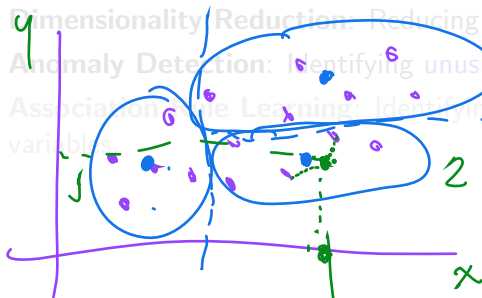
- **Classification:** Predicting a **label**.
- **Regression:** Predicting a **continuous value**.



Typical Machine Learning Problems

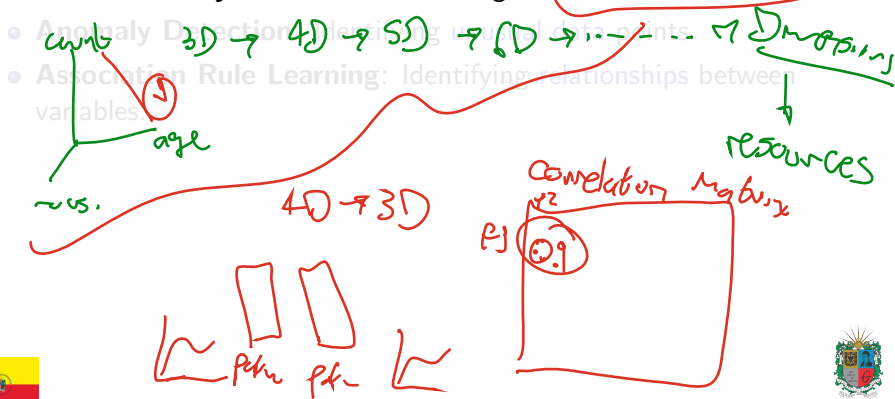
- **Classification:** Predicting a **label**.
- **Regression:** Predicting a **continuous value**.
- **Clustering:** Grouping **similar data** points.
- **Dimensionality Reduction:** Reducing the number of features.
- **Anomaly Detection:** Identifying unusual data points.
- **Association Rule Learning:** Identifying relationships between variables.

supervised
unsupervised



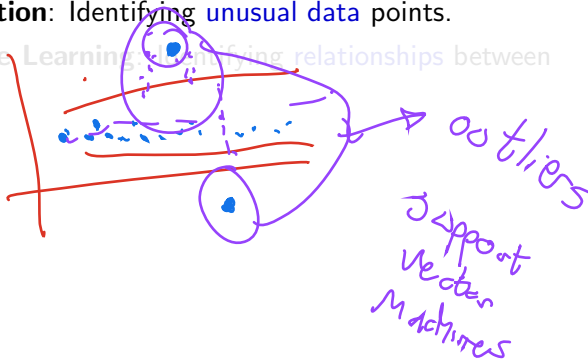
Typical Machine Learning Problems

- **Classification:** Predicting a **label**.
- **Regression:** Predicting a **continuous value**.
- **Clustering:** Grouping **similar data** points.
- **Dimensionality Reduction:** Reducing the **number of features**.



Typical Machine Learning Problems

- **Classification:** Predicting a **label**.
- **Regression:** Predicting a **continuous value**.
- **Clustering:** Grouping **similar data** points.
- **Dimensionality Reduction:** Reducing the **number of features**.
- **Anomaly Detection:** Identifying **unusual data** points.
- **Association Rule Learning:** Identifying relationships between variables.



Typical Machine Learning Problems

- **Classification:** Predicting a **label**.
- **Regression:** Predicting a **continuous value**.
- **Clustering:** Grouping **similar data** points.
- **Dimensionality Reduction:** Reducing the **number of features**.
- **Anomaly Detection:** Identifying **unusual data** points.
- **Association Rule Learning:** Identifying **relationships** between variables.

Language
↳ words
sequence

correlation

f_1
 f_2 } $\rightarrow f_3$



The Machine Learning Workflow

- **Data Collection:** Gathering the data. ①

- **Data Preprocessing:** Cleaning and preparing the data.

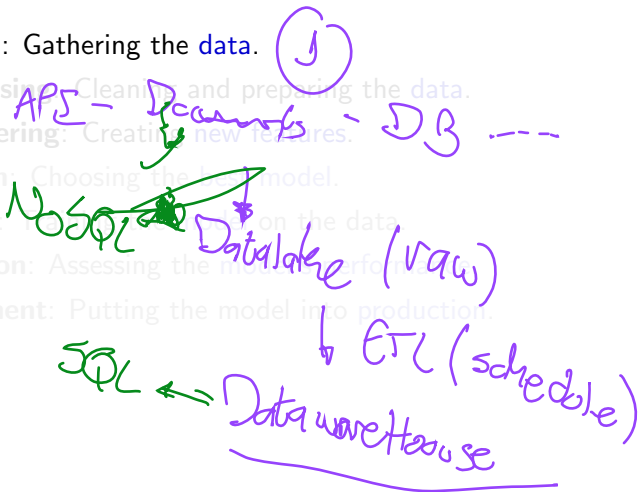
- **Feature Engineering:** Creating new features.

- **Model Selection:** Choosing the best model.

- **Model Training:** Training the model on the data.

- **Model Evaluation:** Assessing the model's performance.

- **Model Deployment:** Putting the model into production.



- Model Deployment: Putting the model into

High quality

features

new features

transform

format

selection



The Machine Learning Workflow


- **Data Collection:** Gathering the data.
- **Data Preprocessing:** Cleaning and preparing the data.
- **Feature Engineering:** Creating new features.

• **Model Selection:** Choosing the best model.

• **Model Training:** Training a model on the data.

• **Model Evaluation:** Assessing the model's performance.

• **Model Deployment:** Putting the model into production.



The Machine Learning Workflow

- **Data Collection:** Gathering the data.
- **Data Preprocessing:** Cleaning and preparing the data.
- **Feature Engineering:** Creating new features.
- **Model Selection:** Choosing the best model.

any ~ 50
108
202

- **Model Training:** Training the model on the data.
- **Model Evaluation:** Assessing the model's performance.
- **Model Deployment:** Putting the model into production.

Depends on data → Trees

cluster → unsupervised

AutoML



experience



The Machine Learning Workflow


- **Data Collection:** Gathering the data.
- **Data Preprocessing:** Cleaning and preparing the data.
- **Feature Engineering:** Creating new features.
- **Model Selection:** Choosing the best model.
- **Model Training:** Training the model on the data.

one-hot
encoding



3x3
2x3
1x3
1x3
1x3

1m
2m



Stewart

10 groups → groups
↓
synthetic
numeric

3x3=9
3x4=12

12x3=36
12x4=48



The Machine Learning Workflow

- **Data Collection:** Gathering the **data**.
- **Data Preprocessing:** Cleaning and preparing the **data**.
- **Feature Engineering:** Creating **new features**.
- **Model Selection:** Choosing the **best model**.
- **Model Training:** Training the **model** on the data.
- **Model Evaluation:** Assessing the **model's performance**

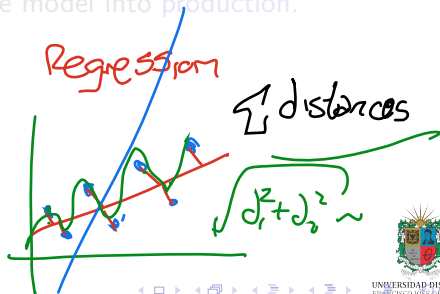

• **Model Deployment:** Putting the model into production.

Metrics \Rightarrow Error

Classification

Regression

• Matrix Confusion



The Machine Learning Workflow

- No SW* → *SW* → *SW* → *SW*
- **Data Collection:** Gathering the **data**.
 - **Data Preprocessing:** Cleaning and preparing the **data**.
 - **Feature Engineering:** Creating **new features**.
 - **Model Selection:** Choosing the **best model**.
 - **Model Training:** Training the **model** on the data.
 - **Model Evaluation:** Assessing the **model's** performance.
 - **Model Deployment:** Putting the model into **production**.



a = ""
b = ""
c = ""

Notebooks

Docker

Cloud (serverless)

classes
oop

API
Rest

direct
Pickle

Parameters

Features



Examining the Data

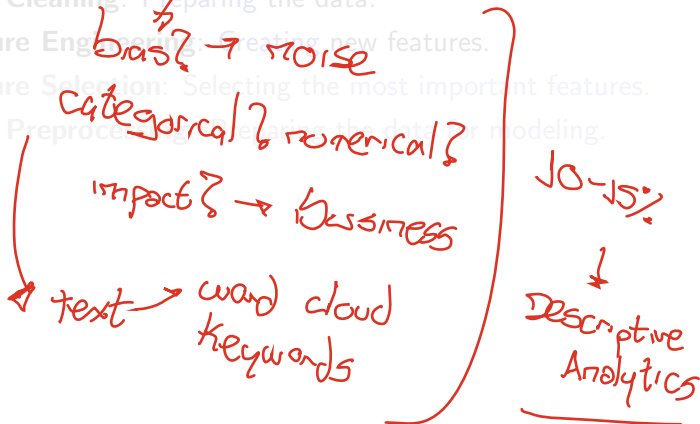
- **Data Exploration:** Understanding the data.

- Data Cleaning: Preparing the data.

- Feature Engineering: Creating new features.

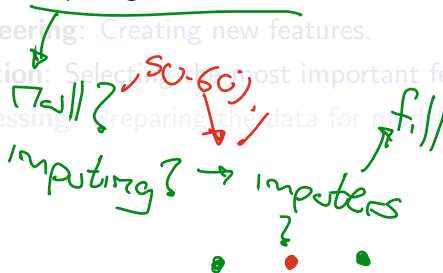
- Feature Selection: Selecting the most important features.

- Data Preprocessing: Preparing the data for modeling.



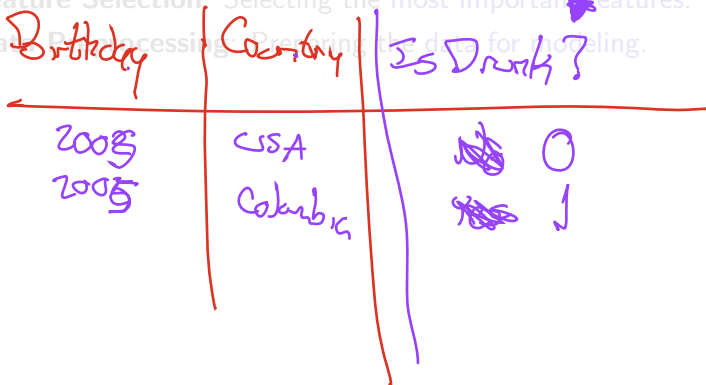
Examining the Data

- **Data Exploration:** Understanding the data.
- **Data Cleaning:** Preparing the data.
- **Feature Engineering:** Creating new features.
- **Feature Selection:** Selecting the most important features.
- **Data Preprocessing:** Preparing the data for modeling.



Examining the Data

- **Data Exploration:** Understanding the data.
- **Data Cleaning:** Preparing the data.
- **Feature Engineering:** Creating new features.
- **Feature Selection:** Selecting the most important features.
- **Data Preprocessing:** Preparing the data for modeling.



A hand-drawn table with three columns: 'Birthday', 'Country', and 'Is Drunk?'. The first column contains '2003' and '2005'. The second column contains 'USA' and 'Colombia'. The third column contains '0' and '1'. A purple arrow points to the 'Is Drunk?' column header.

Birthday	Country	Is Drunk?
2003	USA	0
2005	Colombia	1



Examining the Data

- **Data Exploration:** Understanding the data.
- **Data Cleaning:** Preparing the data.
- **Feature Engineering:** Creating new features.
- **Feature Selection:** Selecting the most important features.
- **Data Preprocessing:** Preparing the data for modeling.

features

A	B	C	D	E

Exp. 1. ABC

Exp. 2. ABCDE

Exp. 3. ACE

Exp. 4. BD



Examining the Data

- **Data Exploration:** Understanding the data.
- **Data Cleaning:** Preparing the data.
- **Feature Engineering:** Creating new features.
- **Feature Selection:** Selecting the most important features.
- **Data Preprocessing:** Preparing the data for modeling.

Categorical
Text } → Numeric → Vector

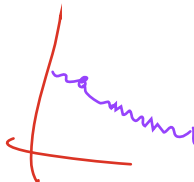


Algorithmic Bias

- **Algorithmic bias** is a **systematic error** in a model that results in **unfair outcomes**.
- It can be caused by **biased training data**, biased algorithms, or biased **decision-making**.

\$ Families →

Loj	car	codes	inc.
-	-	-	-
+	+	+	+
+	+	+	+
+	+	+	+
+	+	+	+



Outline

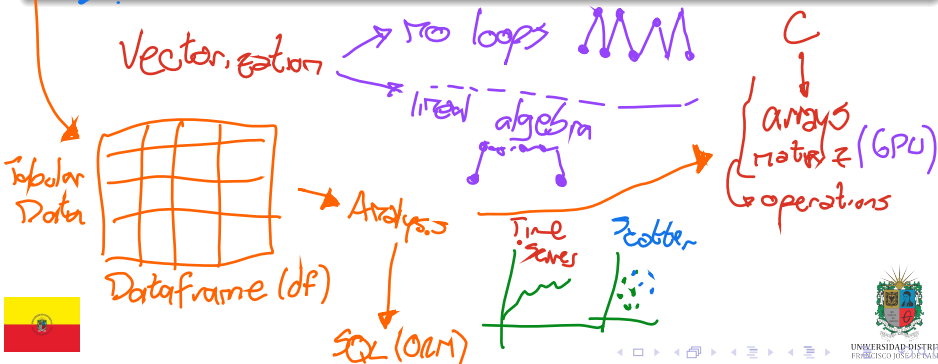
- 1 Fundamentals of Machine Learning
- 2 Python Tools for Machine Learning**
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation



Python Tools for Machine Learning

Python Tools

- NumPy: A library for numerical computing. → ~~Trans~~ ~~Ally~~ ~~dent~~
- Pandas: A library for data manipulation and analysis.
- Matplotlib: A library for data visualization.
- Scikit-learn: A library for machine learning.




Jupyter Notebooks

CC BY

- **Jupyter Notebooks** are a web-based interactive computing environment that allows you to create and share documents that contain live code, equations, visualizations, and narrative text.

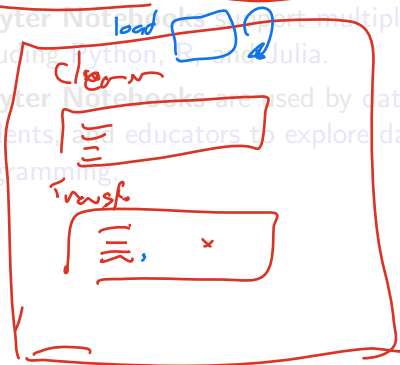
- Jupyter Notebooks are used for data cleaning, data transformation, statistical modeling, data visualization, machine learning, etc.
- Jupyter Notebooks support multiple programming languages, including Python, R, and Julia.
- Jupyter Notebooks are used by data scientists, researchers, students, and educators to explore data, create reports, and teach programming.



Jupyter Notebooks

- **Jupyter Notebooks** are a **web-based interactive computing environment** that allows you to create and share documents that contain **live code**, **equations**, **visualizations**, and **narrative text**.
- **Jupyter Notebooks** are used for data cleaning, data transformation, statistical modeling, data visualization, machine learning, and more.

- Jupyter Notebooks support multiple programming languages, including Python, R, and Julia.
- Jupyter Notebooks are used by data scientists, researchers, students, and educators to explore data, create reports, and teach programming.



Jupyter Notebooks

- **Jupyter Notebooks** are a **web-based interactive computing environment** that allows you to create and share documents that contain **live code**, **equations**, **visualizations**, and **narrative text**.
- **Jupyter Notebooks** are used for **data cleaning**, **data transformation**, **statistical modeling**, **data visualization**, **machine learning**, and **more**.
- **Jupyter Notebooks** support **multiple programming languages**, including **Python**, **R**, and **Julia**.
- **Jupyter Notebooks** are used by **scientists**, **researchers**, **students**, and **educators** to **explore data**, **create reports**, and **teach programming**.

Jupyter

Regular — Data Science



Jupyter Notebooks

- **Jupyter Notebooks** are a **web-based interactive computing environment** that allows you to create and share documents that contain **live code**, **equations**, **visualizations**, and **narrative text**.
- **Jupyter Notebooks** are used for **data cleaning**, **data transformation**, **statistical modeling**, **data visualization**, **machine learning**, and **more**.
- **Jupyter Notebooks** support **multiple programming languages**, including **Python**, **R**, and **Julia**.
- **Jupyter Notebooks** are used by **data scientists**, **researchers**, **students**, and **educators** to **explore data**, **create reports**, and **teach programming**.

Machine Learning Introduction



Lambda Functions

$$f \ x: x+2$$

Definition

A **lambda function** is a **small** **anonymous** function. A **lambda function** can take any number of **arguments**, but can only have one expression.

lambda arguments: expression

$x = \text{lambda } a: a+2$

$\text{print}(x(10)) \rightarrow 20$

↳ lambda (arguments) {
 return expression;



$\text{df}['adult'] = \text{lambda}$
 $\text{df}['age']: \text{if } x \geq 18 \text{ else } 0$



Numerical Python Library — Numpy

- **Numpy** is the core library for scientific computing in Python. It is the fundamental package for scientific computing with Python.
- **Numpy** is a general-purpose array-processing package. It provides a high-performance multidimensional array object, and tools for working with these arrays. *vector*
- **Numpy** was created by Travis Oliphant in 2005, and it is an open-source project. Coming soon, Numpy version 2.0 will be released. *2.1*



Lineal Algebra with Numpy

- **Numpy** provides a comprehensive set of linear algebra functions.
- **Numpy** provides the functionality to create and manipulate matrices.
- **Numpy** provides the functionality to solve linear systems of equations.
- **Numpy** provides the functionality to calculate the determinant of a matrix.
- **Numpy** provides the functionality to calculate the inverse of a matrix.

$$\begin{bmatrix} - & - & - \\ - & - & - \\ - & - & - \end{bmatrix}$$

Eigen
Solve Equations

$$\begin{bmatrix} 2 & 1 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$



Lineal Algebra with Numpy

- **Numpy** provides a **comprehensive set** of **linear algebra** functions.
- **Numpy** provides the **functionality** to **create** and **manipulate matrices**.
- **Numpy** provides the functionality to solve linear systems of equations.
- **Numpy** provides the functionality to calculate the determinant of a matrix.
- **Numpy** provides the functionality to calculate the inverse of a matrix.



Lineal Algebra with Numpy

- **Numpy** provides a **comprehensive set** of **linear algebra** functions.
- **Numpy** provides the **functionality** to **create** and **manipulate matrices**.
- **Numpy** provides the **functionality** to **solve linear systems** of equations.
- **Numpy** provides the **functionality** to **calculate the determinant** of a matrix.
- **Numpy** provides the **functionality** to **calculate the inverse** of a matrix.



Lineal Algebra with Numpy

- **Numpy** provides a **comprehensive set** of **linear algebra** functions.
- **Numpy** provides the **functionality** to **create** and **manipulate matrices**.
- **Numpy** provides the **functionality** to **solve linear systems** of equations.
- **Numpy** provides the **functionality** to **calculate** the **determinant** of a matrix.
- **Numpy** provides the **functionality** to **calculate** the **inverse** of a matrix.



Lineal Algebra with Numpy

- **Numpy** provides a **comprehensive set** of **linear algebra** functions.
- **Numpy** provides the **functionality** to **create** and **manipulate matrices**.
- **Numpy** provides the **functionality** to **solve linear systems** of equations.
- **Numpy** provides the **functionality** to **calculate** the **determinant** of a matrix.
- **Numpy** provides the **functionality** to **calculate** the **inverse** of a matrix.



Vectorization with Numpy

- **Vectorization** is the **process** of **converting** an **algorithm** from **operating** on a single value at a time to **operating** on a set of values at one time.
- **Vectorization** is the **process** of **replacing** explicit **loops** with **array expressions** or **matrix operations**.

- The advantages of vectorization are speed and clarity. The disadvantages are memory and complexity.
- Numpy provides the functionality to vectorize operations on arrays.



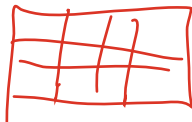
Vectorization with Numpy

- **Vectorization** is the process of converting an algorithm from operating on a single value at a time to operating on a set of values at one time.
- **Vectorization** is the process of replacing explicit loops with array expressions or matrix operations.
- The advantages of vectorization are speed and clarity. The disadvantages are memory and complexity.
- **Numpy** provides the functionality to vectorize operations on arrays.



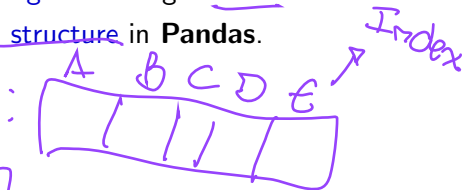
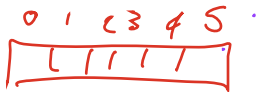
Introduction to Pandas

- **Pandas** is a fast, powerful, flexible, and easy-to-use open-source data manipulation and data analysis library built on top of the Python programming language.
- **Pandas** is a high-level data manipulation tool developed by Wes McKinney in 2008. → *Polars (Rust)*
- **Pandas** is a fast and efficient data manipulation tool that is built on top of NumPy. → *numerical*
- **Pandas** is one of the most popular and widely-used data manipulation libraries in the world.



The “Series” Data Structure

- A **Series** is a one-dimensional array-like object that contains a sequence of values and an associated array of data labels, called the index.
- The index of a **Series** is an array of labels that correspond to the values in the **Series**. The index of a **Series** is an optional parameter that defaults to a sequence of integers starting at zero.
- The **Series** object is a core data structure in **Pandas**.

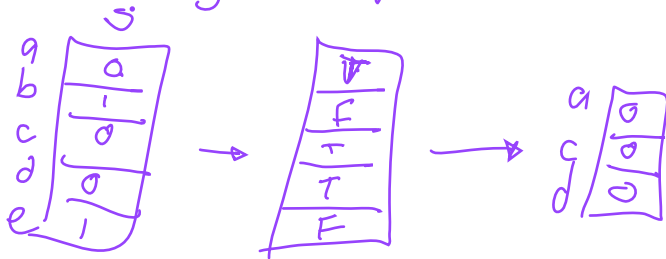


Querying a Series

- You could query a **Series** using indexing (boolean or fancy).

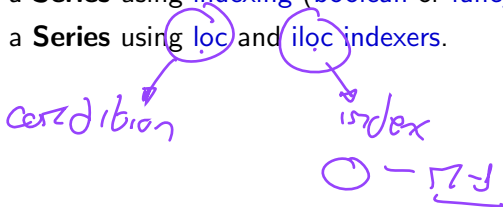
- You could query a Series using loc and ioc indexers.

$$s = s[cond] \rightarrow x=0$$



Querying a Series

- You could **query** a **Series** using **indexing** (boolean or fancy).
- You could **query** a **Series** using **loc** and **iloc** indexers.



 condition

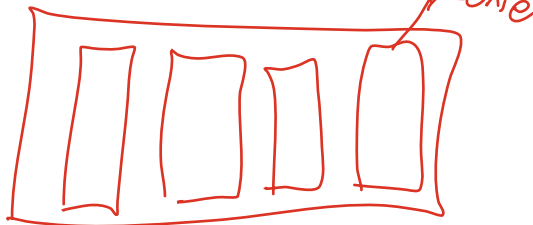
 index

 0 - 123



The "DataFrame" Data Structure

- A **DataFrame** is a two-dimensional labeled data structure with columns of potentially different types. *→ matrix*
- A **DataFrame** is a tabular data structure that is similar to a spreadsheet or a SQL table. *→ operations*
- A **DataFrame** is a core data structure in **Pandas**. It is a two-dimensional size-mutable data structure with labeled axes (rows and columns).
- A **DataFrame** is a container for Series objects. *→ serie*



DataFrame Indexing and Loading

- You could **index** a **DataFrame** using **column names**.
- You could **load** a **DataFrame** from a **CSV** file.
- You could **load** a **DataFrame** from a **JSON** file.
- You could **load** a **DataFrame** from a **SQL** database.



DataFrame Indexing and Loading

- You could **index** a **DataFrame** using **column names**.
- You could **load** a **DataFrame** from a **CSV file**.
- You could **load** a **DataFrame** from a **JSON file**.
- You could **load** a **DataFrame** from a **SQL database**.



DataFrame Indexing and Loading

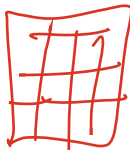
- You could **index** a **DataFrame** using **column names**.
- You could **load** a **DataFrame** from a **CSV file**.
- You could **load** a **DataFrame** from a **JSON file**.
- You could **load** a **DataFrame** from a **SQL database**.

different
SQL



DataFrame Indexing and Loading

- You could **index** a **DataFrame** using **column names**.
- You could **load** a **DataFrame** from a **CSV** file.
- You could **load** a **DataFrame** from a **JSON** file.
- You could **load** a **DataFrame** from a **SQL database**.



DateTime Handling in Pandas

- You could convert a string to a datetime object using the to_datetime() method.
- You could convert a datetime object to a string using the `strftime()` method.
Y-M-D H:M:S
- You could convert a datetime object to a timestamp using the `timestamp()` method.



DateTime Handling in Pandas

- You could **convert** a **string** to a **datetime** object using the **to_datetime()** method.
- You could **convert** a **datetime** object to a **string** using the **strftime()** method.
- You could **convert** a **datetime** object to a **timestamp** using the **timestamp()** method.



DateTime Handling in Pandas

- You could **convert** a **string** to a **datetime** object using the `to_datetime()` method.
- You could **convert** a **datetime** object to a **string** using the `strftime()` method.
- You could **convert** a **datetime** object to a **timestamp** using the `timestamp()` method.

Handwritten diagram illustrating the conversion of a datetime string to a timestamp:

Y-M-D H:M:S → Timestamp Y-M-D H:M:S.000 ms TZ

2025-06-18 10:00:00.000000 (7) → 0 70/10

The diagram shows a datetime string "2025-06-18 10:00:00.000000" being converted to a timestamp. The "7" in the original string is circled and labeled "7". An arrow points from the circled "7" to a "0" which is also labeled "70/10". Another arrow points from the "TZ" part of the timestamp format to the "0" in the timestamp.



Querying a DataFrame

- You could **query** a **DataFrame** using **indexing** (**boolean** or **fancy**).
- You could **query** a **DataFrame** using **loc** and **iloc** indexers.
- You could **query** a **DataFrame** using **query** method.



Querying a DataFrame

- You could **query** a **DataFrame** using **indexing** (**boolean** or **fancy**).
- You could **query** a **DataFrame** using **loc** and **iloc** indexers.
- You could **query** a **DataFrame** using **query** method.



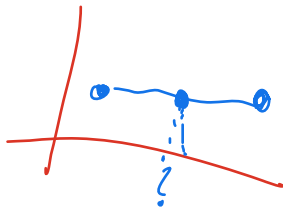
Querying a DataFrame

- You could **query** a **DataFrame** using **indexing** (**boolean** or **fancy**).
- You could **query** a **DataFrame** using **loc** and **iloc** indexers.
- You could **query** a **DataFrame** using **query method**.



Missing Values in a DataFrame

- You could **detect** missing values in a **DataFrame**. The **isnull()** method returns a **Boolean DataFrame** indicating the **presence** of missing values.
 - You could **fill** missing values in a **DataFrame**. The **fillna()** method returns a **DataFrame** with missing values filled.
 - You could **drop** missing values in a **DataFrame**. The **dropna()** method returns a **DataFrame** with missing values dropped.
- Handwritten notes:* "Clear" (above fillna), "Imputer" (above dropna), and "Fill" (with an arrow pointing from dropna to fillna).



Missing Values in a DataFrame

- You could **detect** missing values in a **DataFrame**. The **isnull()** method returns a **Boolean DataFrame** indicating the **presence** of missing values.
- You could **fill** missing values in a **DataFrame**. The **fillna()** method returns a **DataFrame** with missing values filled.
- You could **drop** missing values in a **DataFrame**. The **dropna()** method returns a **DataFrame** with missing values dropped.

Null → NaN



Missing Values in a DataFrame

- You could **detect** missing values in a **DataFrame**. The **isnull()** method returns a **Boolean DataFrame** indicating the **presence** of missing values.
- You could **fill** missing values in a **DataFrame**. The **fillna()** method returns a **DataFrame** with missing values filled.
- You could **drop** missing values in a **DataFrame**. The **dropna()** method returns a **DataFrame** with missing values dropped.

Null → NaN → row at least 1
 NaN → gone



Merging DataFrames

- You could **merge** two **DataFrames** using the **merge()** method.
- You could **concatenate** two **DataFrames** using the **concat()** method.
- You could **join** two **DataFrames** using the **join()** method.

Join
concatenations



Merging DataFrames

- You could **merge** two **DataFrames** using the **merge()** method.
- You could **concatenate** two **DataFrames** using the **concat()** method.
- You could ~~join~~ two **DataFrames** using the **join()** method.



Merging DataFrames

- You could **merge** two **DataFrames** using the **merge()** method.
- You could **concatenate** two **DataFrames** using the **concat()** method.
- You could **join** two **DataFrames** using the **join()** method.



GroupBy in Pandas

- You could **group** a **DataFrame** using the **groupby()** method.
- You could **aggregate** a **DataFrame** using the **agg()** method.
- You could **transform** a **DataFrame** using the **transform()** method.
- You could **filter** a **DataFrame** using the **filter()** method.



GroupBy in Pandas

- You could **group** a **DataFrame** using the **groupby()** method.
- You could **aggregate** a **DataFrame** using the **agg()** method.
- You could **transform** a **DataFrame** using the **transform()** method.
- You could **filter** a **DataFrame** using the **filter()** method.



GroupBy in Pandas

- You could **group** a **DataFrame** using the `groupby()` method.
- You could **aggregate** a **DataFrame** using the `agg()` method.
- You could **transform** a **DataFrame** using the `transform()` method.
- You could **filter** a **DataFrame** using the `filter()` method.



GroupBy in Pandas

- You could **group** a **DataFrame** using the `groupby()` method.
- You could **aggregate** a **DataFrame** using the `agg()` method.
- You could **transform** a **DataFrame** using the `transform()` method.
- You could **filter** a **DataFrame** using the `filter()` method.



Outline

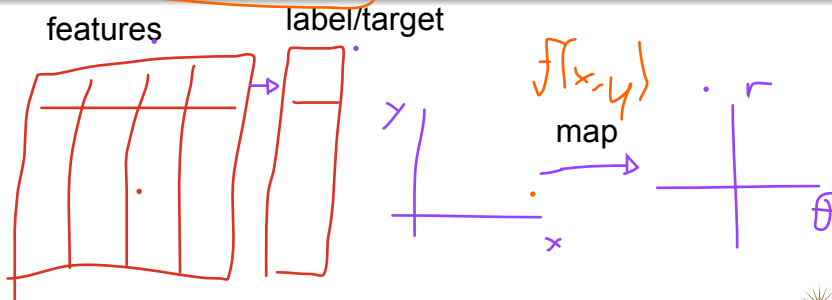
- 1 Fundamentals of Machine Learning
 -
- 2 Python Tools for Machine Learning
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation



Introduction to Supervised Machine Learning

Definition

- **Supervised learning** is a type of **machine learning** where the model is trained on **labeled data**.
- It involves training a model to **map** **input data** to **output data** based on example **input-output pairs**.



Overfitting and Underfitting

Overfitting

Overfitting occurs when a model learns the training data too well and performs poorly on new data.

Underfitting

Underfitting occurs when a model is too simple to capture the underlying structure of the data.



Overfitting and Underfitting

Overfitting

Overfitting occurs when a model learns the training data too well and performs poorly on new data.

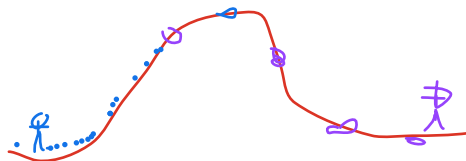
Underfitting

Underfitting occurs when a model is too simple to capture the underlying structure of the data.

Iterations

Hill Climbing

Learning Rate



Supervised Learning Datasets

- **Training Dataset:** The data used to train the model.

- **Validation Dataset:** The data used to tune the model hyperparameters.

- **Test Dataset:** The data used to evaluate the model performance.
Learning a pattern to map
features -> label

80-70
9



Supervised Learning Datasets

- **Training Dataset:** The data used to **train the model**.
- **Validation Dataset:** The data used to **tune** the model **hyperparameters**.
- **Test Dataset:** The data used to **evaluate** the model **performance**.

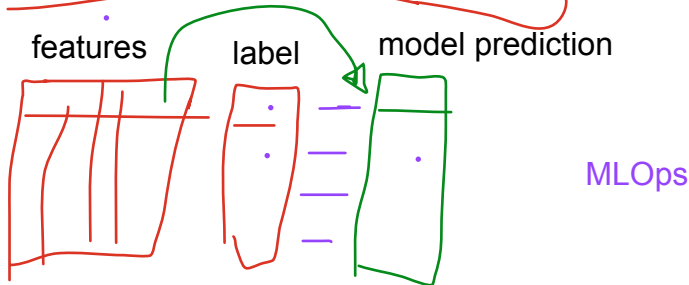
model setup

- * iterations
- * learning rate
- * optimization metric (error)



Supervised Learning Datasets

- **Training Dataset:** The data used to **train the model**.
- **Validation Dataset:** The data used to **tune** the model **hyperparameters**.
- **Test Dataset:** The data used to **evaluate** the model **performance**.



Cross-Validation

- **Cross-validation** is a technique for assessing the performance of a model.
- It involves splitting the data into multiple subsets, training the model on some subsets, and evaluating it on others.
- Common cross-validation **techniques** include k-fold cross-validation and leave-one-out cross-validation. 13 4
- Cross-validation helps to reduce overfitting and provides a more accurate estimate of the model's performance.



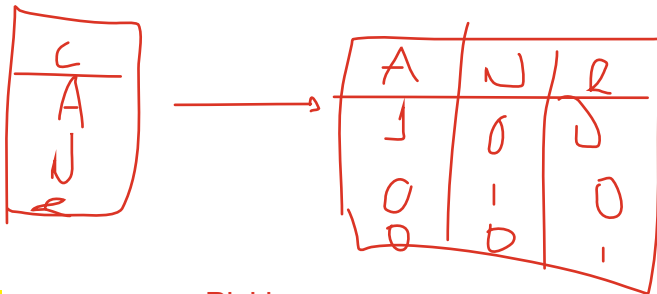
Training 1: 1234 Testing: 5
 • Training 2: 1235 Testing: 4



One-Hot Encoding

One-Hot Encoding

- One-hot encoding is a technique for converting categorical variables into numerical variables.
- It creates a binary vector for each category, with a 1 for the *category* and 0s for all other categories.




Pickle •



Data Leakage

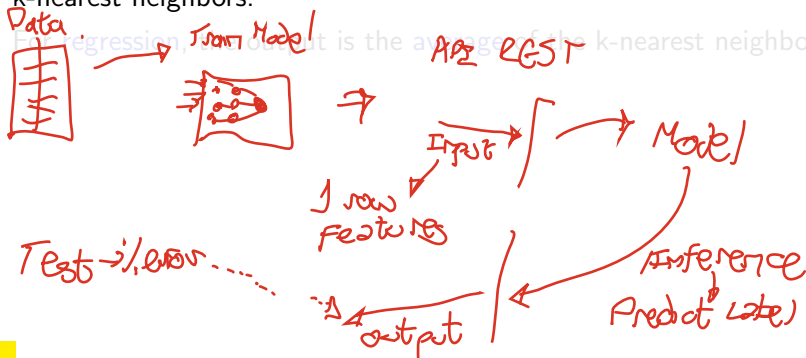
- **Data leakage** occurs when information from the test set is inadvertently used to train the model.
- It can lead to overfitting and inflated performance metrics.
- Common sources of **data leakage** include target leakage, train-test contamination, and information leakage.
- To prevent **data leakage**, it is important to carefully separate the training and test data and avoid using information from the test set during training.



K-Nearest Neighbors: Classification and Regression


- **K-Nearest Neighbors (KNN)** is a simple algorithm that stores all available cases and classifies new cases based on a **similarity measure**.
- It can be used for both **classification** and **regression** tasks.
- For **classification**, the output is the **class label** of the majority of the k-nearest neighbors.

- For **regression**, the output is the **average** of the k-nearest neighbors.



K-Nearest Neighbors: Classification and Regression

- **K-Nearest Neighbors (KNN)** is a simple algorithm that stores all available cases and classifies new cases based on a similarity measure.
- It can be used for both **classification** and **regression** tasks.
- For **classification**, the output is the **class label** of the majority of the k-nearest neighbors.
- For **regression**, the output is the **average** of the **k**-nearest neighbors.



Linear Regression with Least Squares

Linear Regression

- **Linear regression** is a type of regression analysis used for predicting the value of a continuous dependent variable.
- It works by finding the line that best fits the data.

Least Squares

Least squares is a method for finding the best-fitting line by minimizing the sum of the squared differences between the predicted and actual values.



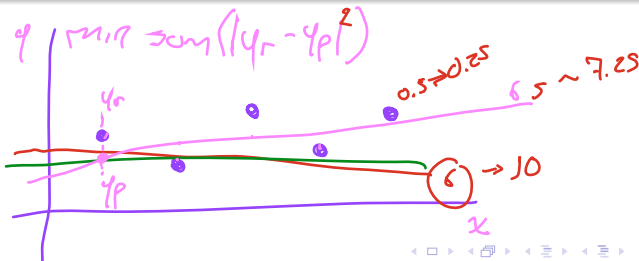
Linear Regression with Least Squares

Linear Regression

- **Linear regression** is a type of **regression analysis** used for predicting the value of a **continuous dependent variable**.
- It works by finding the **line that best fits the data**.

Least Squares

Least squares is a method for finding the **best-fitting line** by **minimizing the sum** of the squared differences between the **predicted** and **actual** values.

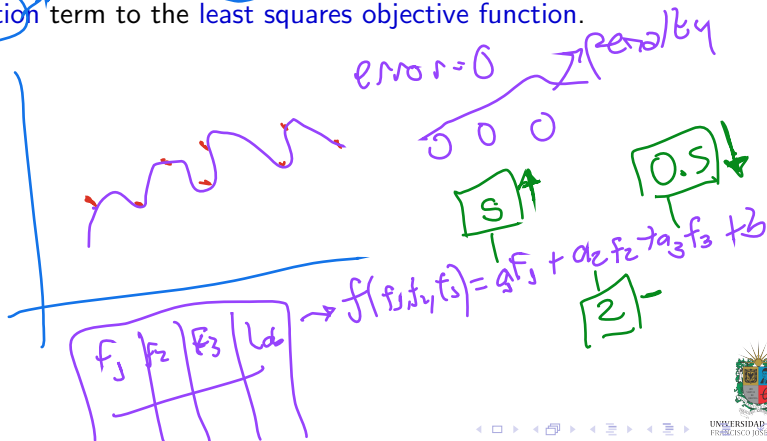


Ridge & Lasso

L1 → overfitting

L2 → feature importance

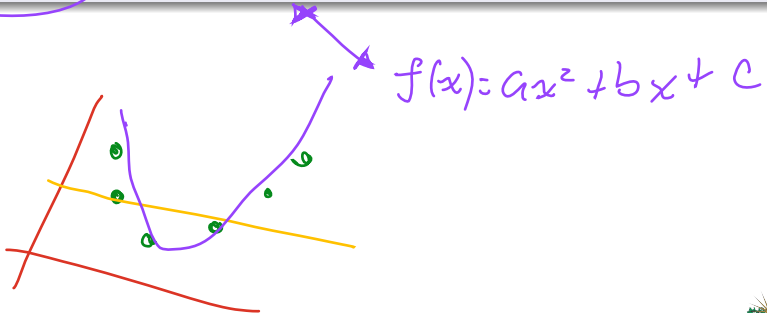
Ridge regression & **Lasso regression** are a type of linear regression that includes a penalty term to prevent overfitting. It works by adding a regularization term to the least squares objective function.



Polynomial Regression

Polynomial Regression

- **Polynomial regression** is a type of regression analysis that models the relationship between the independent and dependent variables as an n th-degree polynomial.
- It can capture non-linear relationships between the variables.



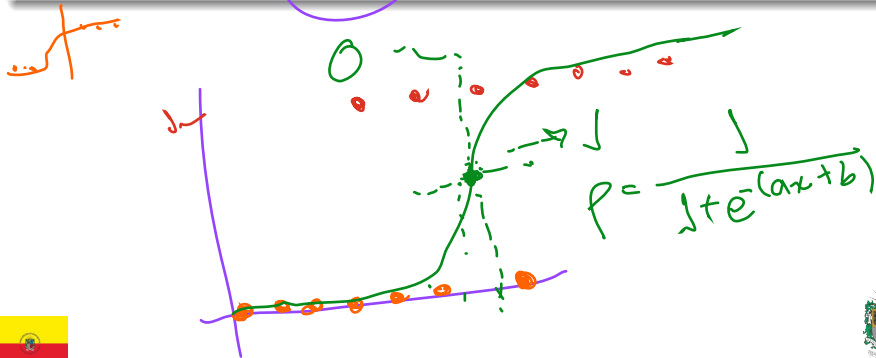
$$f(x) = ax^2 + bx + c$$



Logistic Regression

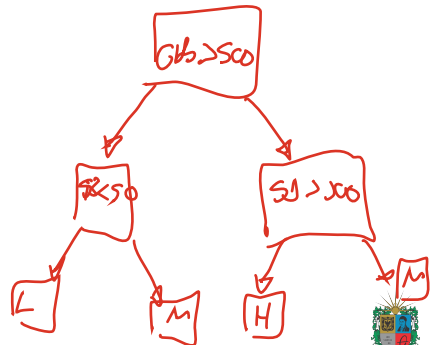
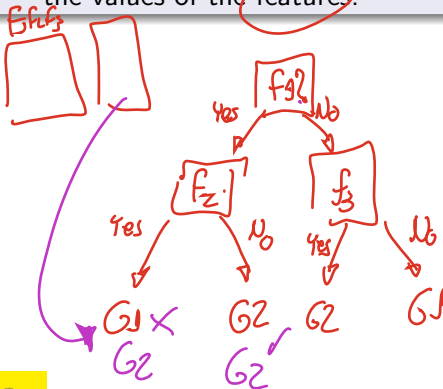
Logistic Regression

- **Logistic regression** is a type of regression analysis used for predicting the outcome of a categorical dependent variable.
- It is used for binary classification tasks, where the output is a probability between 0 and 1.



Decision Trees

- **Decision trees** are a type of **machine learning model** that can be used for both **classification** and **regression** tasks.
- They work by recursively **partitioning** the data into **subsets** based on the values of the features.



Naive Bayes Classifier

- The **naive Bayes classifier** is a simple probabilistic **classifier** based on Bayes' theorem.
- It assumes that the features are **conditionally independent** given the class label.

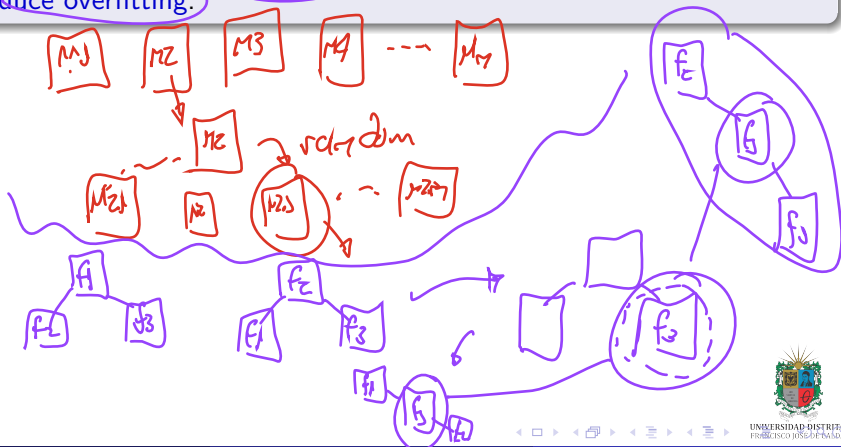
$P_{\text{rain}} \rightarrow P_{\text{wind}} \rightarrow P_{\text{sun}} \rightarrow P_{\text{dependent}}$

Baseline / performance



Random Forest

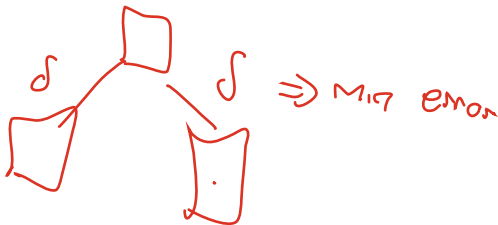
- **Random forest** is an **ensemble learning** method that combines **multiple decision trees** to create a strong predictive model.
- It works by building **multiple trees** and averaging their predictions to reduce overfitting.



Gradient Boosted Decision Trees

- **Gradient boosted decision trees** are an ensemble learning method that combines multiple decision trees and gradient descent optimization to create a strong predictive model.
- They work by building trees sequentially, with each tree correcting the errors of the previous trees.

Bagging



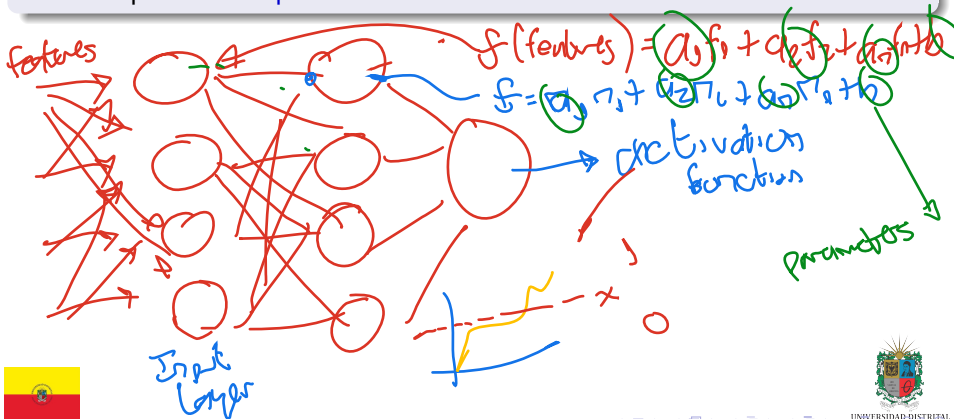
Neural Networks

ReLU

Softmax



- **Neural networks** are a type of **machine learning model** inspired by the **human brain**.
- They consist of **layers** of interconnected nodes that process **input data** and produce **output data**.



Outline

- 1 Fundamentals of Machine Learning
- 2 Python Tools for Machine Learning
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation**



Model Evaluation & Selection

- **Model Evaluation**: Assessing the **performance** of a model.
- **Model Selection**: Choosing the **best model** for the task.

	M1		M2	
	error1	error2	error1	error2
Dataset 1	x_1	y_1	x_2	y_2
Dataset 2	x_3	y_3	x_4	y_4
Dataset 3	x_5	y_5	x_6	y_6

error
predict vs expected

S-JO \Rightarrow average



Confusion Matrices

↳ Classification

Definition

- A **confusion matrix** is a **table** that summarizes the **performance** of a **classification model**.
- It shows the number of **true positives**, **true negatives**, **false positives**, and **false negatives**.


		real / (expected)	
		1	0
Predicted	1	True Positive → 1.0	False Positive → 0.0
	0	False Negative → 0.0	True Negative → 1.0



Basic Evaluation Metrics

- **Accuracy**: The proportion of correct predictions.

- **Precision**: The proportion of true positives among all positive predictions.
- **Recall**: The proportion of true positives among all actual positives.
- **F1 Score**: The harmonic mean of precision and recall.



$$\leadsto \text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN} \quad \text{Total rows test} \quad \text{e.g. 1.0}$$




Basic Evaluation Metrics

- **Accuracy:** The proportion of **correct predictions**.
- **Precision:** The proportion of true positives among all positive predictions.

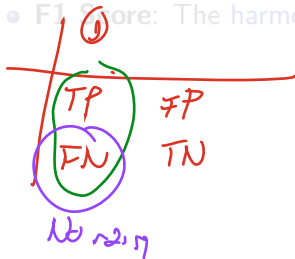
• **Recall:** The proportion of true positives among all actual positives.

• **F1 Score:** The harmonic mean of precision and recall.



Basic Evaluation Metrics

- **Accuracy:** The proportion of **correct predictions**.
- **Precision:** The proportion of **true positives** among **all positive predictions**.
- **Recall:** The proportion of **true positives** among **all actual positives**.
- **F1 Score:** The harmonic mean of precision and recall.



$$\text{recall} = \frac{TP}{TP + FN}$$

↓

J.O



Basic Evaluation Metrics

- **Accuracy:** The proportion of correct predictions.
- **Precision:** The proportion of true positives among all positive predictions.
- **Recall:** The proportion of true positives among all actual positives.
- **F1 Score:** The harmonic mean of precision and recall.

$$F_1 \text{ Score} = \frac{\text{Precision} + \text{Recall}}{2 \times \text{Precision} \times \text{Recall}}$$

$\Rightarrow \frac{2}{\frac{1}{p} + \frac{1}{r}}$




Classifier Metrics

- **ROC Curve:** A plot of the true positive rate against the false positive rate.

- Precision-Recall Curve: A plot of precision against recall.


- AUC-ROC: The area under the ROC curve.

- AUC-PR: The area under the precision-recall curve.




Classifier Metrics

- **ROC Curve:** A plot of the true positive rate against the false positive rate.
- **Precision-Recall Curve:** A plot of precision against recall.
- **AUC-ROC:** The area under the ROC curve.
- **AUC-PR:** The area under the precision-recall curve.




Classifier Metrics

- **ROC Curve:** A plot of the true positive rate against the false positive rate.
- **Precision-Recall Curve:** A plot of precision against recall.
- **AUC-ROC:** The area under the ROC curve.



Classifier Metrics

- **ROC Curve:** A plot of the true positive rate against the false positive rate.
- **Precision-Recall Curve:** A plot of precision against recall.
- **AUC-ROC:** The area under the ROC curve.
- **AUC-PR:** The area under the precision-recall curve.



Regression Metrics

- **Mean Squared Error:** The average of the squared differences between the predicted and actual values.

- **Mean Absolute Error:** The average of the absolute differences between the predicted and actual values.

- **R-Squared:** The proportion of the variance in the dependent variable that is predictable from the independent variables.

- **Adjusted R-Squared:** A modified version of R-squared that adjusts for the number of predictors in the model.

$$mse = \frac{(y_{i, \text{pred}} - y_{i, \text{best}})^2}{n}$$

- **Root Mean Squared Error:** The square root of the mean squared error.

$$\sqrt{mse}$$



Regression Metrics

- **Mean Squared Error:** The **average** of the **squared differences** between the predicted and actual values.
- **Mean Absolute Error:** The **average** of the **absolute differences** between the predicted and actual values.
- **R-Squared:** The **proportion** of the **variance** in the dependent variable that is **predictable** from the independent variables.
- **Adjusted R-Squared:** A **modified version** of **R-squared** that adjusts for the **number of predictors** in the model.
- **Root Mean Squared Error:** The **square root** of the **mean squared error**.

$$mae = \frac{|y_{pred} - y_{best}|}{len(y)}$$



Regression Metrics

- **Mean Squared Error:** The **average** of the **squared differences** between the predicted and actual values.
- **Mean Absolute Error:** The **average** of the **absolute differences** between the predicted and actual values.
- **R-Squared:** The **proportion** of the **variance** in the **dependent variable** that is predictable from the independent variables.

Handwritten notes for R-Squared:
 - A purple line connects "proportion" to "variance".
 - A purple line connects "variance" to "features".
 - A purple line connects "dependent variable" to "y pred".
 - A purple line connects "features" to "var (features)".
 - A purple line connects "var (features)" to "var (y pred)".
 - A purple line connects "var (y pred)" to "corr".
- **Adjusted R-Squared:** A modified version of R-squared that adjusts for the number of predictors in the model.
- **Root Mean Squared Error:** The square root of the mean squared error.



Regression Metrics

- **Mean Squared Error:** The **average** of the **squared differences** between the predicted and actual values.
 - **Mean Absolute Error:** The **average** of the **absolute differences** between the predicted and actual values.
 - **R-Squared:** The **proportion** of the **variance** in the dependent variable that is predictable from the independent variables.
 - **Adjusted R-Squared:** A modified version of **R-squared** that adjusts for the **number of predictors** in the model.
 - **Root Mean Squared Error:** The square root of the mean squared error.
- features most significance*



Regression Metrics

- **Mean Squared Error:** The **average** of the **squared differences** between the predicted and actual values.
- **Mean Absolute Error:** The **average** of the **absolute differences** between the predicted and actual values.
- **R-Squared:** The **proportion** of the **variance** in the dependent variable that is predictable from the independent variables.
- **Adjusted R-Squared:** A modified version of R-squared that adjusts for the **number of predictors** in the model.
- **Root Mean Squared Error:** The **square root** of the mean squared error.

$$rmse = \sqrt{\frac{\sum_{i=0}^n (y_i^{pred} - y_i^{test})^2}{n}}$$

↓

○



Outline

$$f() = a_0 r_1 + a_1 r_2 + \dots + a_{64} r_{64} + b$$

↓ Trillion

- 1 Fundamentals of Machine Learning
- 2 Python Tools for Machine Learning
- 3 Supervised Machine Learning
- 4 Machine Learning Models Evaluation

$$f = a_0 f_1 + a_1 f_2 + a_2 f_3 + a_3 f_4 + a_4 f_5 + c$$



Thanks!

Questions?



Repo: <https://github.com/EngAndres/ud-public/tree/main/courses/systems-sciences-foundations>



UNIVERSIDAD DISTRITAL
FRANCISCO JOSÉ DE CALDAS